

Learn to Build Automated Software Analysis Tools with Graph Paradigm and Interactive Visual Framework

GIAN, September 12-16, 2016

Suresh C. Kothari
Richardson Professor
Department of Electrical and Computer Engineering

Ben Holland, Iowa State University

Module VI: Software Analysis Lab Exercises and Projects

Acknowledgement: Team members at Iowa State University and EnSoft, DARPA contracts FA8750-12-2-0126 & FA8750-15-2-0080

Software Lab Exercises Projects

Hopefully, you can use this material in your research and teaching, build on it, and share what you have added. THANK YOU.

- Goals:

- Learn to solve complex software problems by applying critical thinking.
 - Practice the graph paradigm with Atlas to apply critical thinking to large software.

- Sample Lab Exercises and Projects:

- **Lab exercises:** (1) Experiments using Atlas to understand and apply program analysis with graph paradigm, (2) Implement software metrics and analyze results using Java code, (3) Detect a class of software safety and security vulnerabilities
 - **Project** – Use graph models to verify the Linux kernel for a safety property

- Planned activities:

- Understand the importance of the posed problem and its challenges.
 - Identify the relevant fundamental principles and develop a problem solving strategy.
 - Learn by performing experiments.
 - Learn to communicate the solution precisely and succinctly.

Lab Exercise 1: Understand and apply program analysis

- Goal

- Hands-on experience to reinforce the course material.
 - Prepare participants to make effective use of the graph paradigm in their research, education, and practice.

- Planned Activities

- Use the analysis and visualization capabilities interactively.
 - Write and execute short program analysis queries.
 - Introduce the eXtensible Common Software Graph schema for C and Java.

Lab Exercise 2: Understand and apply program analysis

- Goal: Learn how to compute software metrics using Atlas and analyze the pros and cons of a particular metric
- The lab will have experiments to study:
 - What indicators does well designed software have?
 - What indicators does poorly designed software have?

Software Architecture (SA)

- A good SA is essential for easy software maintenance (modification of the software product after delivery to correct faults, to improve performance or other attributes) and adding new features.
- Our focus: the aspect of SA that deals with architecture of the modules (packages and classes) and their interconnections.

Motivating Questions....

- How difficult can (even a simple) a code change be?
 - Does it cause a cascade of subsequent changes in dependent modules? (Rigidity)
 - Does it break software in multiple places – even in areas not conceptually related to the original code change? (Fragility)
- How easy is it to reuse a software module?
 - What is the overhead in separating the desirable parts from undesirable parts? Is the overhead so much that rewriting the software would be easier? (Immobility)
- Is the right way the hard way?
 - Do the design constraints make it hard to follow the right way, and instead an improperly designed “hack” is easy to implement. (Viscosity)

References

- Software Package Metrics

- Design Principles and Design Patterns, Robert C. Martin¹, 2000. Article available at: http://www.objectmentor.com/resources/articles/Principles_and_Patterns.pdf
- Wikipedia: http://en.wikipedia.org/wiki/Software_package_metrics
- Book “Agile Software Development: Principles, Patterns, and Practices”: http://books.google.com/books/about/Agile_Software_Development.html?id=OHYhAQAAIAAJ

¹Martin is a software professional since 1970 and an international software consultant since 1990. In 2001, he initiated the meeting of the group that created agile software development from extreme programming techniques.

See http://en.wikipedia.org/wiki/Robert_Cecil_Martin

Package Coupling Principles

- Applications are networks of interrelated packages.
- The interrelationship should obey certain principles for efficient software maintenance.
 - “Package Coupling Principles”
- We capture packages interrelations by package dependency metrics.
- Package X is dependent on package Y if X cannot be compiled without Y.

Package Coupling Principles

- Package X is dependent on package Y if
 - Some code in X read/writes into a field declared in some class in Y
 - A class in X extends/implements a class in Y
 - Some code in X calls a method in Y

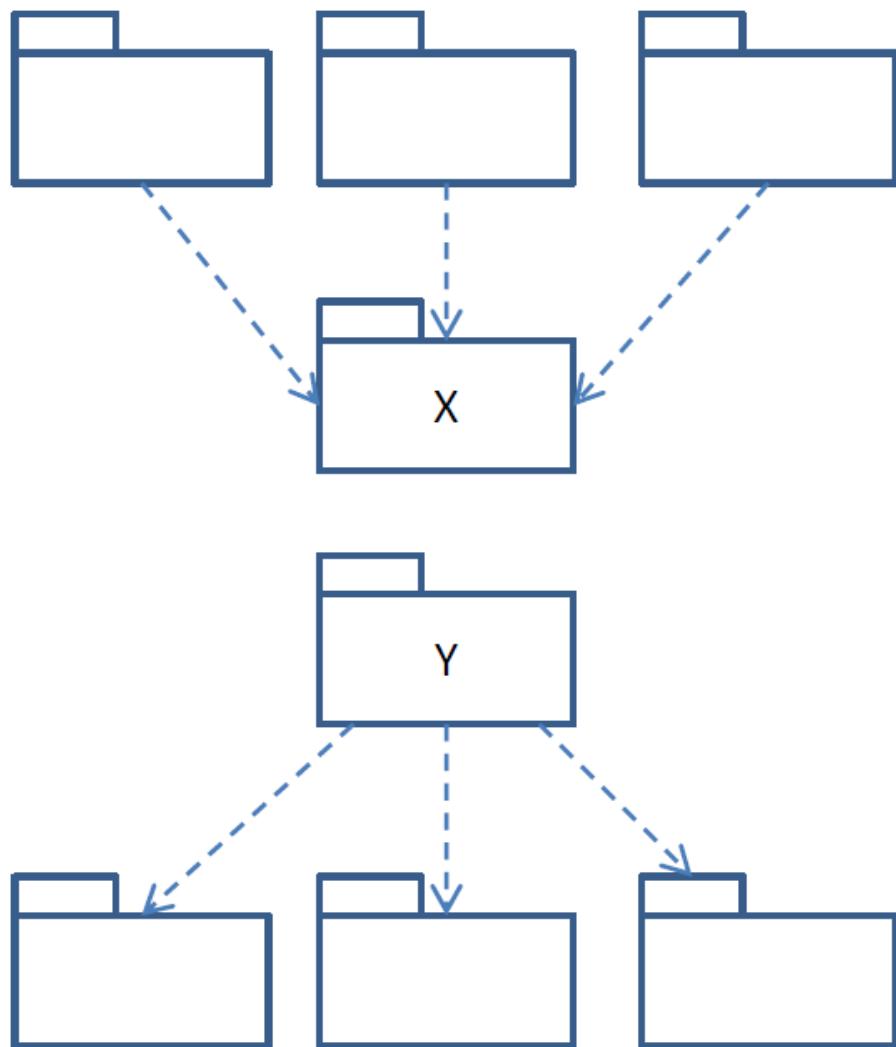
The Acyclic Dependencies Principle

- o *The dependencies between packages must not form cycles.*
 - Code-changes are usually localized to a package because a good design groups related classes together in a package. The changed package should be compiled and tested with the packages it depends on. If there is a cyclic dependency, a change in any one package in the cycle requires the test suite to be built with all the packages in the cycle – clearly not a desirable thing.

The Stable Dependencies Principle (SDP)

- A package should be stable if it has many packages dependent on it.
 - Because it requires a great deal of work to reconcile any changes in the stable package with all the dependent packages.
 - For such a package should be stable, it should not be heavily dependent on other packages.

Dependent vs. Independent Packages

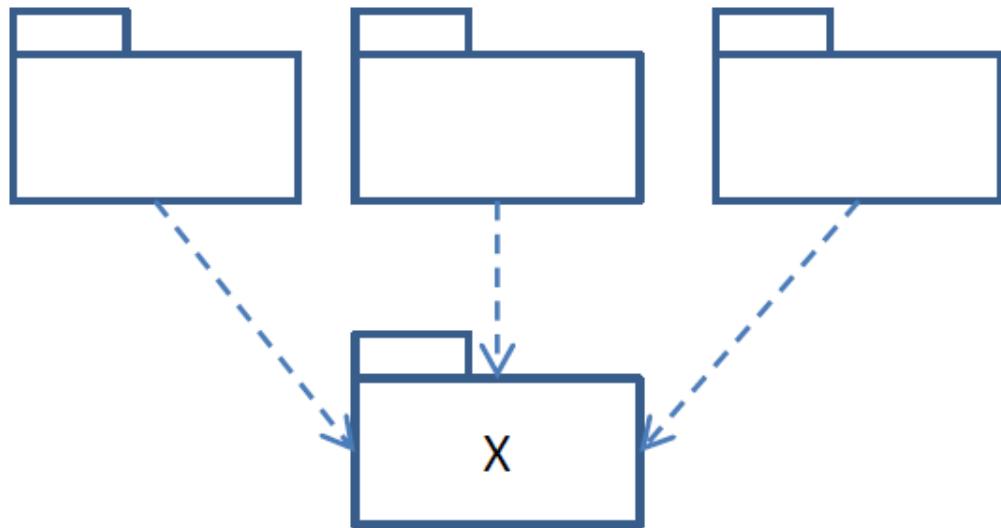


- X is independent
 - It has three packages dependent on it.
 - X depends upon nothing, so it has no external influence to make it change. We say it is *independent*.
- Y is dependent
 - It has no packages depending on it.
 - Y depends on three packages, so changes may come from three external sources. We say that Y is *dependent*.

Stability Metrics

- Efferent Coupling (C_e): The number of classes outside the package that the classes inside the package depend upon (outgoing dependencies).
- Afferent Coupling (C_a): The number of classes outside the package that depend upon the classes inside the package (incoming dependencies).
- Instability (I): $C_e/(C_a+C_e)$
 - Range [0,1]; 0: very stable, 1: very unstable.
- SDP can be rephrased as: “Depend upon packages whose instability metric is lower than yours.”
- We can envision the packages structure of our application as a set of interconnected packages with unstable packages at the top, and stable packages on the bottom. In this view, all dependencies point downwards.

The Stable Abstractions Principle (SAP)

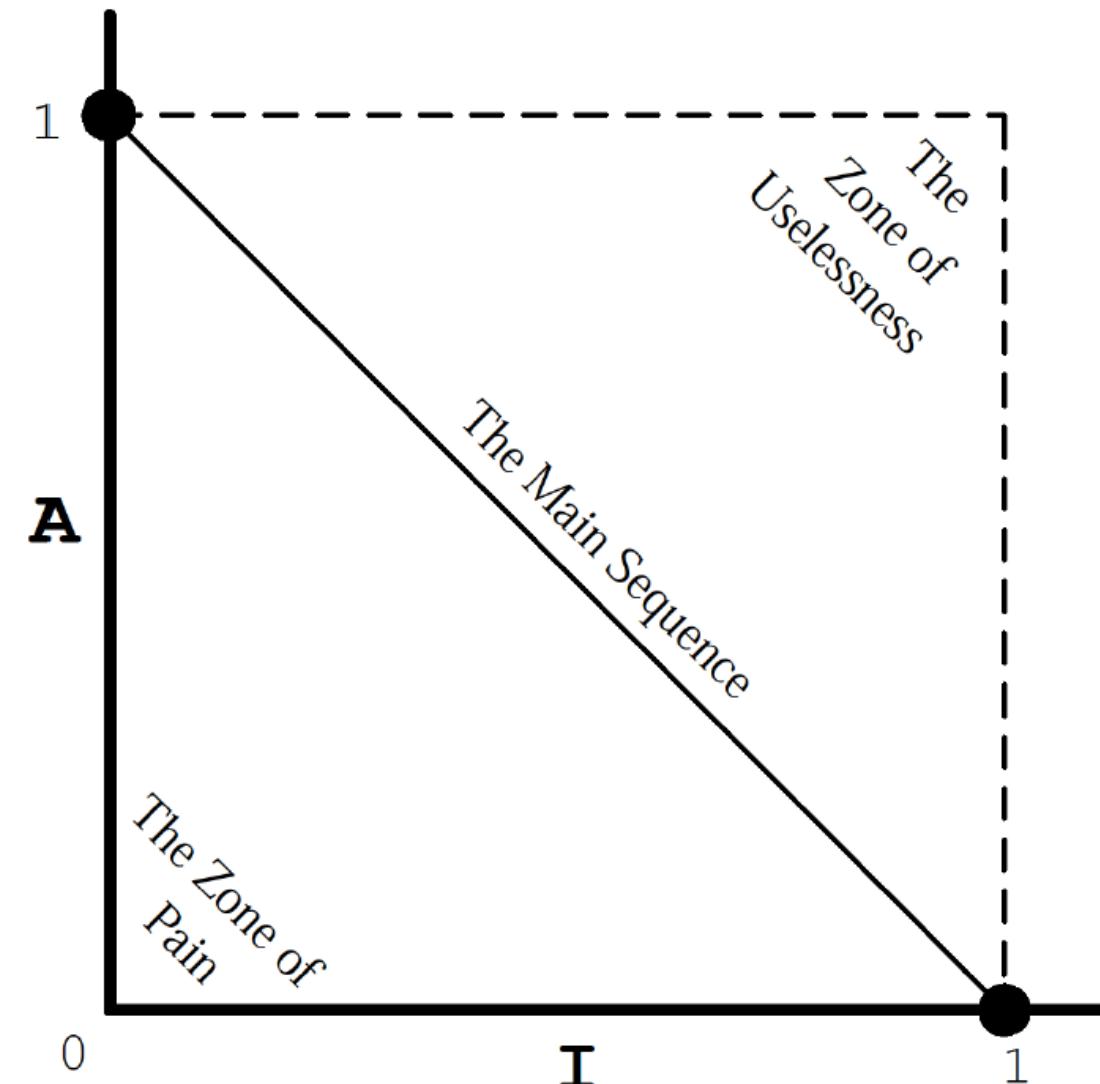


- **SAP:** Heavily used packages should have a high-degree of abstraction.
- A high-degree of abstraction implies that packages are easy to change or extend.

Abstractness Metrics

- N_c : Number of classes in the package.
- N_a : Number of abstract classes in the package.
- Abstractness (A) = N_a/N_c
 - Range [0,1]; 0: no abstract classes, 1: all classes are abstract
- SAP can be restated as: I should increase as abstractness decreases. That is, concrete packages should be unstable (flexible) while abstract packages should be stable.

Abstractness vs Instability



- Ideally a package should lie on “The Main Sequence” line
- Distance (D): $|A+I-1|$
 - Range $[0,1]$; 0: package directly on the main sequence, 1: farthest away from the main sequence.

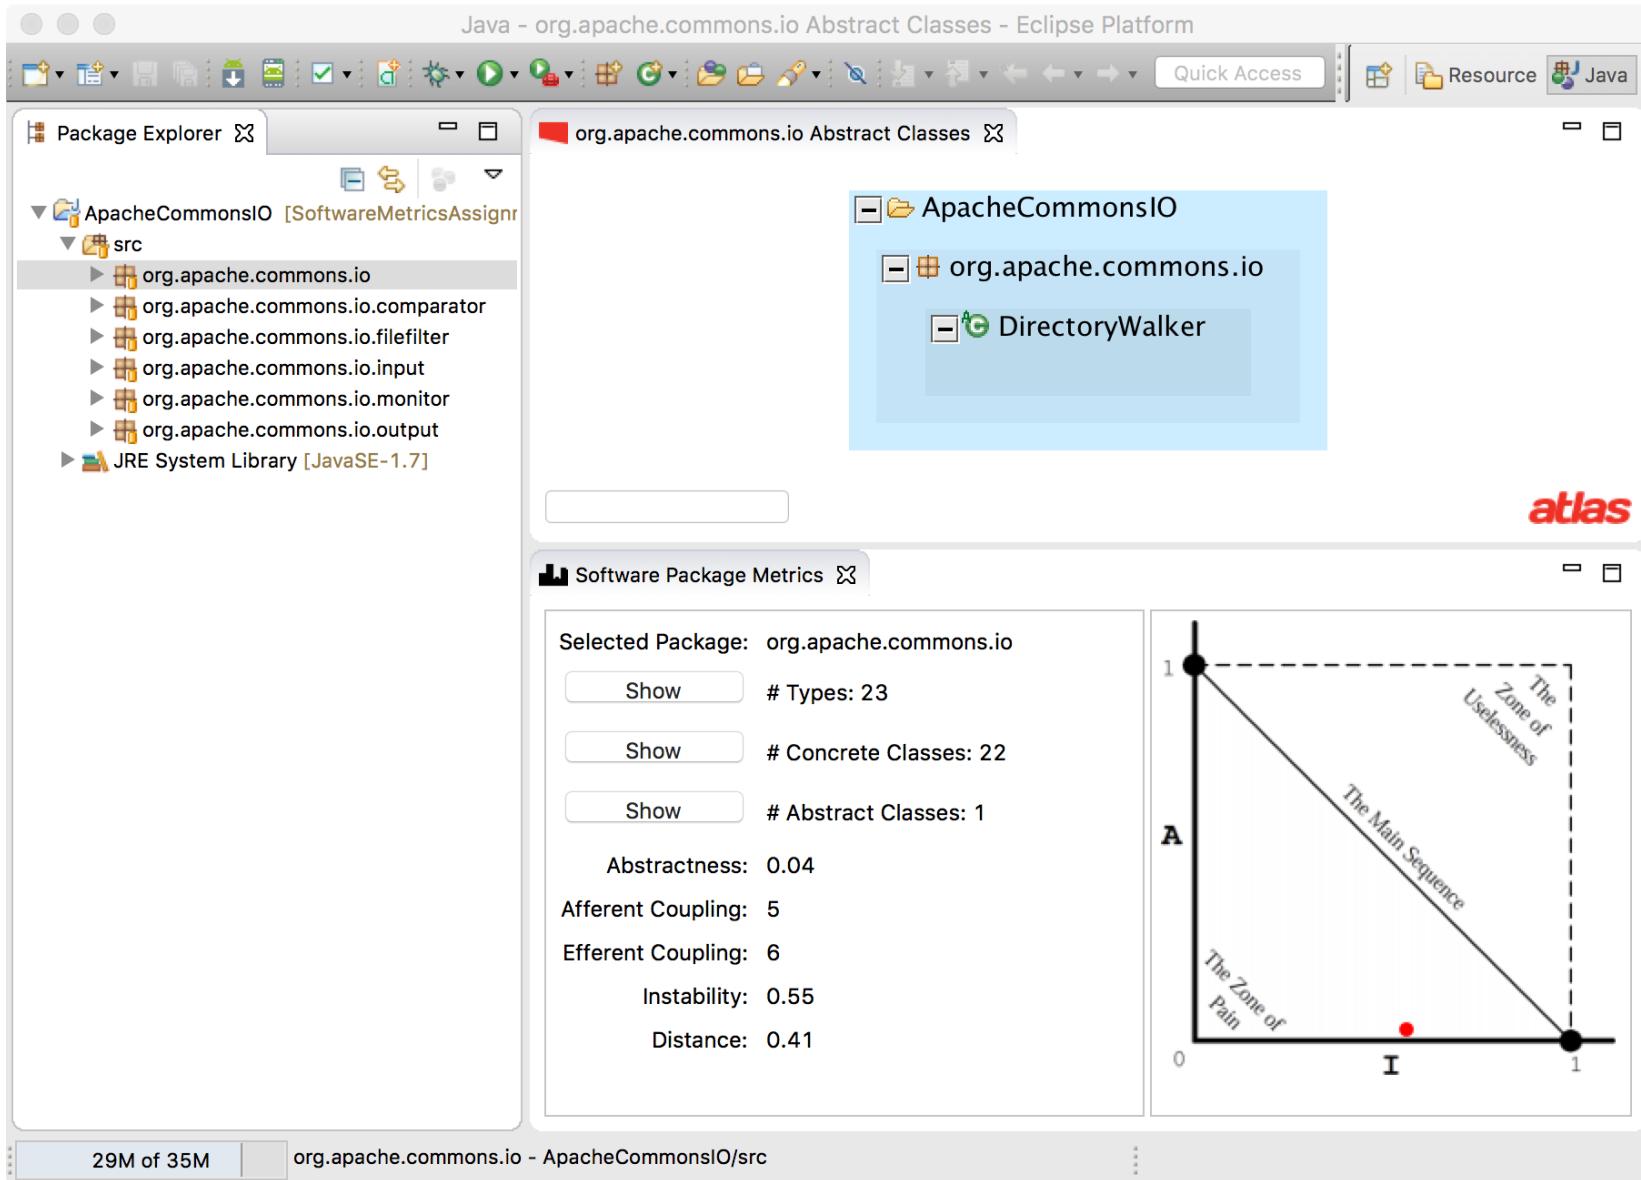
Reflections on Software Metrics

- These metrics are meant to provide a measure of how good the software architecture is.
- But, they are imperfect, and reliance upon them as the sole indicator of a sturdy architecture would be foolhardy.

Classroom Activity

- Develop a program analysis tool to compute the previously described software metrics
 - Implement TODOs in base project:
<https://github.com/benjholla/SoftwareMetricsAssignment>
- Evaluate the Apache Commons IO library

Classroom Activity



Classroom Activity

- With regard to the Distance (D) from the Main Sequence is the org.apache.commons.io package better or worse compared to the org.apache.commons.io.filefilter package?
- Do you feel these metrics accurately capture the quality of this code?

Classroom Activity: Solution

```
public static Q getClasses(Q graph){  
    return graph.nodesTaggedWithAny(XCSG.Java.AbstractClass, XCSG.Java.Class);  
}  
  
public static Q getInterfaces(Q graph){  
    return graph.nodesTaggedWithAll(XCSG.Java.Interface);  
}  
  
public static Q getAbstractClasses(Q graph){  
    return graph.nodesTaggedWithAny(XCSG.Java.AbstractClass, XCSG.Java.Interface);  
}  
  
public static Q getConcreteClasses(Q graph) {  
    return getClasses(graph).difference(getAbstractClasses(graph));  
}
```

Classroom Activity: Solution

```
public static double getAbstractness(Node pkg) throws ArithmeticException {  
    // Step 1) Get the types declared under the package  
    Q packageTypes = getPackageTypes(pkg);  
    // Step 2) From the discovered class nodes select the abstract classes  
    Q abstractClasses = getAbstractClasses(packageTypes);  
    // Step 4) Return the ratio of the number of abstract classes (and interfaces) in the package to the total number of types in the package  
    return (double) countNodes(abstractClasses) / (double) countNodes(packageTypes);  
}
```

Classroom Activity: Solution

```
public static Q getAfferentCouplings (Node pkg){  
    // Step 1) Get the methods declared under the package  
    Q packageMethods = getPackageMethods(pkg);  
    // Step 2) Create a subgraph of CALL edges from the universe  
    Q callEdges = Common.universe().edgesTaggedWithAny(XCSG.Call).retainEdges();  
    // Step 3) Within the calls subgraph get the predecessors of the package methods (calling methods)  
    Q callingMethods = callEdges.predecessors(packageMethods);  
    // Step 4) Return the packages of the calling methods  
    return getPackagesOfProgramElements(callingMethods);  
}
```

Classroom Activity: Solution

```
public static Q getEfferentCouplings (Node pkg){  
    // Step 1) Get the methods declared under the package  
    Q packageMethods = getPackageMethods(pkg);  
    // Step 2) Create a subgraph of CALL edges from the universe  
    Q callEdges = Common.universe().edgesTaggedWithAny(XCSG.Call).retainEdges();  
    // Within the calls subgraph get the successors of the package methods (called methods)  
    Q calledMethods = callEdges.successors(packageMethods);  
    // Step 4) Return the packages of the called methods  
    return getPackagesOfProgramElements(calledMethods);  
}
```

Classroom Activity: Solution

```
public static double getInstability(Node pkg) {  
    double ce = (double) countNodes(getEfferentCouplings(pkg));  
    double ca = (double) countNodes(getAfferentCouplings(pkg));  
    return ce / (ca + ce);  
}  
  
public static double getDistance(Node pkg) {  
    double a = getAbstractness(pkg);  
    double i = getInstability(pkg);  
    return Math.abs(a + i - 1);  
}
```

Classroom Activity: Solution

| Package | Types | Concrete Classes | Abstract Classes | Abstractness | Afferent Coupling | Efferent Coupling | Instability | Distance | Properties |
|----------------------------------|-------|------------------|------------------|--------------|-------------------|-------------------|-------------|----------|----------------|
| org.apache.commons.io | 23 | 22 | 1 | 0.04 | 5 | 6 | 0.55 | 0.41 | Least Abstract |
| org.apache.commons.io.comparator | 10 | 9 | 1 | 0.1 | 4 | 3 | 0.43 | 0.47 | |
| org.apache.commons.io.filefilter | 25 | 22 | 3 | 0.12 | 3 | 6 | 0.67 | 0.21 | Least Stable |
| org.apache.commons.io.input | 27 | 24 | 3 | 0.11 | 6 | 3 | 0.33 | 0.56 | Most Stable |
| org.apache.commons.io.monitor | 5 | 4 | 1 | 0.2 | 3 | 4 | 0.57 | 0.23 | Most Abstract |
| org.apache.commons.io.output | 19 | 18 | 1 | 0.05 | 4 | 3 | 0.43 | 0.52 | |

- With regard to the Distance (D) from the Main Sequence is the org.apache.commons.io package better or worse compared to the org.apache.commons.io.filefilter package?
 - According to the metric, the distance D from the main sequence indicates org.apache.commons.io is worse designed than org.apache.commons.io.filefilter, since the org.apache.commons.io.filefilter package is significantly closer to 0 than the org.apache.commons.io package.

Classroom Activity: Solution

| Package | Types | Concrete Classes | Abstract Classes | Abstractness | Afferent Coupling | Efferent Coupling | Instability | Distance | Properties |
|----------------------------------|-------|------------------|------------------|--------------|-------------------|-------------------|-------------|----------|----------------|
| org.apache.commons.io | 23 | 22 | 1 | 0.04 | 5 | 6 | 0.55 | 0.41 | Least Abstract |
| org.apache.commons.io.comparator | 10 | 9 | 1 | 0.1 | 4 | 3 | 0.43 | 0.47 | |
| org.apache.commons.io.filefilter | 25 | 22 | 3 | 0.12 | 3 | 6 | 0.67 | 0.21 | Least Stable |
| org.apache.commons.io.input | 27 | 24 | 3 | 0.11 | 6 | 3 | 0.33 | 0.56 | Most Stable |
| org.apache.commons.io.monitor | 5 | 4 | 1 | 0.2 | 3 | 4 | 0.57 | 0.23 | Most Abstract |
| org.apache.commons.io.output | 19 | 18 | 1 | 0.05 | 4 | 3 | 0.43 | 0.52 | |

- Do you feel these metrics accurately capture the quality of this code?
 - Depends. Code metrics are an approximation that capture some properties of the software. Keep in mind that Apache Commons IO is a library. What would happen if we consider the libraries interactions with the JDK? What would happen if we consider the libraries interactions with an application using the library?

Lab Exercise 3: Detect a class of software vulnerabilities

- Goal
 - Bring out the challenges of intricate data and control-flow dependencies.
 - Demonstrate how graph models can address the challenges.
- Planned Activities
 - Perform experiments to understand the challenges of multiplicative growth of execution paths and path feasibility.
 - Identify the fundamentals necessary to address the challenges.
 - Apply the principles to design graph models for effective detection and verification.
 - Practice critical thinking through experiments to design the graph models.

A software vulnerability problem

- Problem: Check whether the memory allocation in the XINU disk driver code leads to memory leak.
- Given: relevant disk driver routines – about 200 lines of C code.

Problem Details

Problem Overview: The given code includes some disk driver routines. `getbuf()` and `freebuf()` are respectively the calls to *allocate* and *deallocate* memory dynamically. The goal is to verify that an *allocation* is followed by *deallocation* on all feasible execution paths. The code presents two allocation instances: (a) a simple instance in `dsread()` where allocation is followed by deallocation in the same function, (b) a complex instance in `dswrite()` where allocation is *not* followed by deallocation in the same function. How to verify that allocation in `dswrite()` is followed by deallocation on all feasible execution paths? How many paths are there? Where could be the matching deallocation, if not in `dswrite()`?

Side Questions: Why is the deallocation not in `dswrite()` itself? Is there a good justification to not have the deallocation in `dswrite()`? what is that justification?

Discussion Points: Why is the problem so hard? What are the prevalent practices to solve the problem? Why do we need a fundamentally different approach to solve the problem?

Memory leak verification

- o Problem: Verify that each instance $A(x)$ of memory allocation (event E1) is followed by deallocation $D(X)$ (event E2) on every *feasible* execution path.
 - This is a 2-event problem.
 - Non-occurrence of an event is fault!

Lock/Unlock Pairing verification

- o Problem: Verify that each instance $L(x)$ of Lock (event E1) is followed by Unlock $U(X)$ (event E2) on every *feasible* execution path.
 - This is a 2-event problem.
 - Non-occurrence of an event is fault!

Anti-Matching Pair (AMP) verification

- Problem: Verify that event $E1(X)$ is *not* followed by event $E2(X)$ on every *feasible* execution path.
 - This is a 2-event problem.
 - Occurrence of an event is fault!

A concrete example of AMP

- **Sensitive Source Event** : Occurrence of an API with the permission to *read* phone state. Let us call this $E1(X)$ where X is the pointer to the phone state.
- **Malicious Sink**: Occurrence of an API with the permission to *access Internet*. Let us call this $E2(X)$ where X is the pointer to the phone state.
- **Problem**: Verify that event $E1(X)$ is *not* followed by event $E2(X)$ on every *feasible* execution path.

MP: False positives and false negatives

| REALITY | VERIFICATION RESULT | ACCURACY |
|------------------------------------------|---------------------|----------|
| E1(x) followed by E2(x) on all paths | Violation YES | FP |
| E1(x) followed by E2(x) on all paths | Violation NO | Correct |
| E1(x) not followed by E2(x) on some path | Violation YES | Correct |
| E1(x) not followed by E2(x) on some path | Violation NO | FN |

AMP: False positives and false negatives

| REALITY | VERIFICATION RESULT | ACCURACY |
|------------------------------------------|---------------------|----------|
| E1(x) not followed by E2(x) on all paths | Violation YES | FP |
| E1(x) not followed by E2(x) on all paths | Violation NO | Correct |
| E1(x) followed by E2(x) on some path | Violation YES | Correct |
| E1(x) followed by E2(x) on some path | Violation NO | FN |

Sophisticated vulnerabilities

- The vulnerability has a catastrophic unaccepted behavior.
- A reactive fix is not a good solution.
- The trigger for enacting the vulnerability is obscure and rare.
- Testing by itself is not a viable verification option.
- Discovering the vulnerability is like looking for the needle in the haystack, not knowing what the needle looks like.
- The vulnerabilities are often hidden due to invisible control and/or data flows in the program.

The commonly used reliability testing, software reliability measures, analysis of factors causing defects, and the cybersecurity precautions are generally not applicable or effective in the case of critical vulnerabilities.

Defining problem hardness

- The hardness of verifying one instance can be very different than the hardness of verifying another instance.
- Without knowing the hardness, we are working in the dark to advance the research.
- What defines the hardness?

Memory Leak Example 1

```
dsread(devptr, buff, block)
    struct devsw *devptr;
    char *buff;
    DBADDR block;
{
    struct dreq *drptr;
    int stat;
    char ps;

    disable(ps);
    drptr = (struct dreq *) getbuf(dsksrbp);
    drptr->drdba = block;
    drptr->drpid = currpid;
    drptr->drbuff = buff;
    drptr->drop = DREAD;
    if ((stat=dskenq(drptr, devptr->dvioblk)) == DONQ) {
        stat = drptr->drstat;
    }
    freebuf(drptr);
    restore(ps);
    return(stat);
}
```

Verification Easy!

Memory Leak Example 2

```
dswrite(devptr, buff, block)
    struct devsw      *devptr;
    char *buff;
    DBADDR   block;
{
    struct dreq *drptr;
    char ps;

    disable(ps);
    drptr = (struct dreq *) getbuf(dsksrbp);
    drptr->drbuff = buff;
    drptr->drdba = block;
    drptr->drpid = currpid;
    drptr->drop = DWRITE;
    dskenq(drptr, devptr->dvioblk);
    restore(ps);
    return(OK);
}
```

Verification Hard!

Let us get to the fundamentals of what makes it hard.

Let's try to Characterize *Hardness*

```
dswrite(devptr, buff, block)
    struct devsw *devptr;
    char *buff;
    DBADDR block;
{
    struct dreq *drptr;
    char ps;

    disable(ps);
    drptr = (struct dreq *) getbuf(dskrbp);
    drptr->drbuff = buff;
    drptr->drdba = block;
    drptr->drpid = currpid;
    drptr->drop = DWRITE;
    dskenq(drptr, devptr->dvioblk);
    restore(ps);
    return(OK);
}
```

Interprocedural – not a good enough characterization of hardness, Why?

Let's Look a Little Deeper - 1

```
dskenq(drptr, dsprt)          4 paths
    struct dreq *drptr;
    struct dsblk *dsprt;
{ struct dreq *p, *q;
if ( (q=dsprt->dreqlst) == DRNULL ) {
1   |   dsprt->dreqlst = drptr;
1   |   dskstrt(dsprt);
1   |   return();
for () {
2   |   if ((st=dskqopt(p, q, drptr)!=SYSERR))
2   |       return();
2   |   if ( )
3   |       q->drnext = drptr;
3   |       return(); }
4   |   q->drnext = drptr;
4   |   return(); }
```

Path 1: drptr passed to dskstrt(), which does not deallocate –a memory leak?

Let's Look a Little Deeper - 2

```
dskenq(drptr, dsprt)
    struct dreq *drptr;
    struct dsblk *dsprt;
{ struct dreq *p, *q;
if ( (q=dsprt->dreqlst) == DRNULL ) {
1 |     dsprt->dreqlst = drptr;
    dskstrt(dsprt);
    return();
for () {
2 |     if ((st=dskqopt(p, q, drptr)!=SYSERR))
        return();
    if ( )
        3 |         q->drnext = drptr;
        return(); }
    q->drnext = drptr;
4 | return(); }
```

More paths ➔ Harder

Hardness should account for # paths.

Paths 3 & 4: drptr not passed to callee – memory leak?

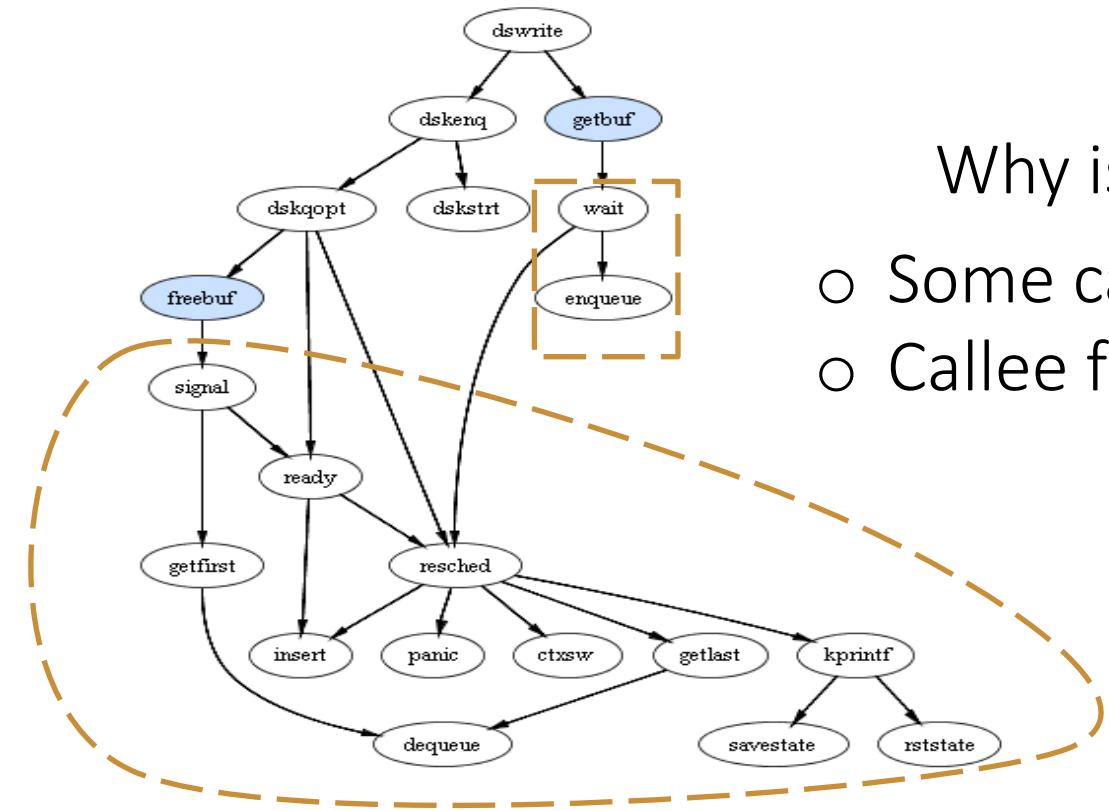
Let's Look a Little Deeper - 3

```
dskenq(drptr, dsprt)
    struct dreq *drptr;
    struct dsblk *dsprt;
{ struct dreq *p, *q;
if ( (q=dsprt->dreqlst) == DRNULL ) {
1   |   dsprt->dreqlst = drptr;
1   |   dskstrt(dsprt);
1   |   return();
for () {
2   |   if ((st=dskqopt(p, q, drptr)!=SYSERR))
2   |       return();
2   |   if ( )
3   |       q->drnext = drptr;
3   |       return(); }
4   |   q->drnext = drptr;
4   |   return(); }
```

drptr gets passed back to the caller(s) on paths 3 and 4.

Let's Look a Little Deeper - 4

Counting the number of *callee* functions not a good metric to quantify hardness.



Why is it not a good metric?

- Some callee functions *not* relevant.
- Callee functions are not the only relevant functions.

Call graph is neither *necessary* nor *sufficient* to identify *relevant functions*.

Let's Look a Little Deeper - 5

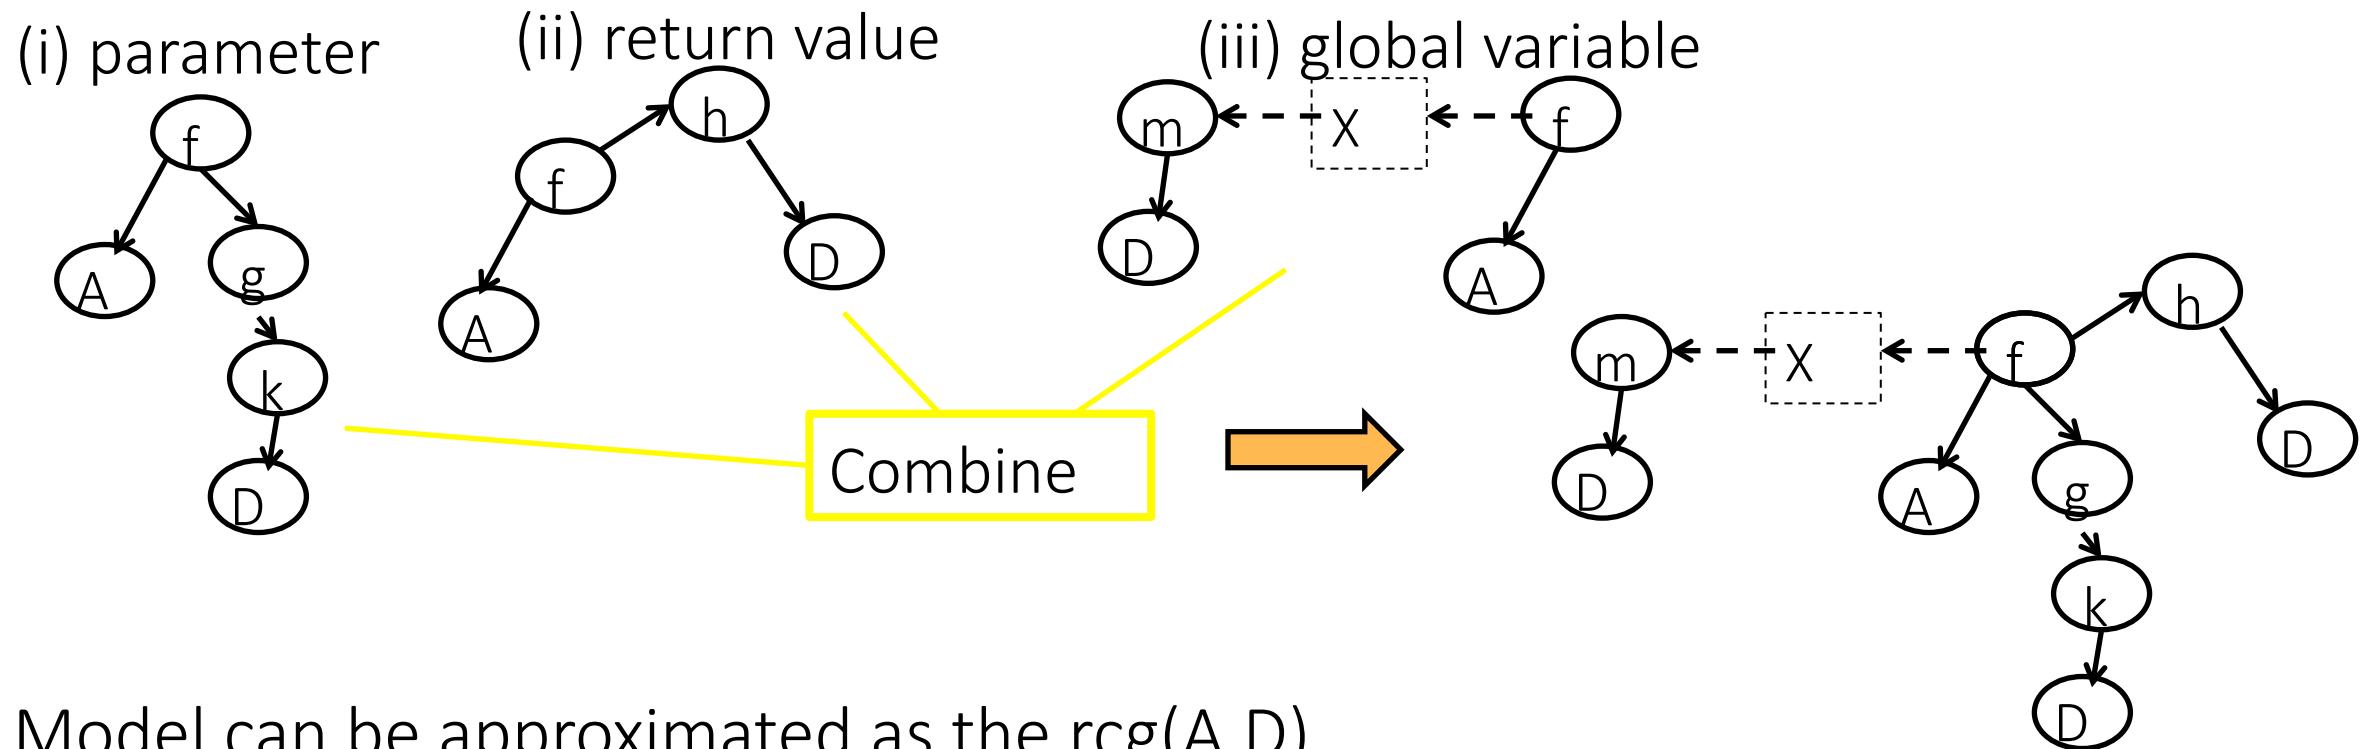
Secondly, it is not just about how many functions interact but also about how functions interact.

Fundamental Interaction Mechanisms:

1. Function f passes a pointer as a parameter to a callee function g .
2. Function g passes a pointer as a parameter or return value back to a caller function g .
3. Function f assigns a pointer to a global variable D .

Does one mechanism make verification harder than the other? Why?

A comprehensive model of communication among functions



Model can be approximated as the $\text{rcg}(A, D)$

A: memory allocation function, D: deallocation function

$\text{rcg}(A, D)$: the *reverse call graph* with A and D as leaves

Exploration experiments using Atlas

At this point, let us explore using Atlas.

Explorations will be about:

1. Identifying the relevant functions and their interactions.
2. Addressing the challenge of multitude of paths and the need to check their feasibility.

Some queries for the exploration

Background: Memory is allocated for a structure of type dreq and drptr points to a structure of type dreq. getbuf and freebuf respectively *allocate* and *deallocate* memory.

Queries:

1. Find functions which pass drptr to callee as a parameter, along with call relationships.
2. Find functions which pass drptr back to caller as a parameter, along with call relationships.
3. Find functions that write drptr to a global variable.
4. Find functions that read drptr from a global variable.
5. Find functions that have a call chain reaching to either getbuf or freebuf.
6. Find the union of the above, along with call relationships.
7. Find functions that have: (i) a call chain reaching to either getbuf or freebuf, and (ii) reference (read or write) drptr.
 - A. Will #6 and #7 give the same answer? Why or why not?
 - B. Is either of the answer necessary and/or sufficient for verifying the memory leak?

Interactive use of Atlas

- We explore visual models interactively.
- We will show an example of how interactive experiments can be useful for discovering new ideas and invent new visual models.
- The Lab covers the Atlas features:
 - Querying through Interpreter
 - Querying by referencing selecting elements
 - Smart Views
- The lab covers the visual models: Call Graph (CG), Reverse Call Graph (RCG), Control Flow Graph (CFG), Data Flow Graph (DFG), Program Slices, and Matching Pair Graph (MPG).

Project Explorer

```

arpfind.c
arpinit.c
blkcopy.c
blkequ.c
bpdump.c
chprioc.c
ckmode.c
cksum.s
clkinit.c
clkint.c
clkint.s
close.c
conf.c
control.c
create.c
csv.s
ctxsw.c
ctxsw.s
devdump.c
dfalloc.c
dfdsrch.c
dgalloc.c
dgclose.c
dgcntl.c
dgdump.c
dginit.c
dgmcntl.c
dgmopen.c
dgparse.c
dgread.c
dgwrite.c
dot2ip.c
dscntl.c
dsinit.c
dsinter.c
dskbcpy.c
dskenq.c
dskqopt.c
dskstrt.c
dsksync.c
dsopen.c
dsread.c
dsseek.c
dswrite.c
dudir.c

```

dswrite.c

```

/* dswrite.c - dswrite */

#include <conf.h>
#include <kernel.h>
#include <proc.h>
#include <disk.h>

* dswrite -- write a block (system buffer) onto a disk device..
dswrite(devptr, buff, block)
    struct devsw *devptr;
    char *buff;
    DBADDR block;
{
    struct dreq *drptr;
    char ps;

    disable(ps);
    drptr = (struct dreq *) getbuf(dskrbp);
    drptr->drbuff = buff;
    drptr->drdba = block;
    drptr->drpid = currpid;
    drptr->drop = DWRITE;
    dskenq(drptr, devptr->dvioblk);
    restore(ps);
    return(OK);
}

```

// Method for creating a call graph

```

public Q cg(Q function) {
    return edges(XCSG.Call).forward(function);
}

```

// Invoke CG Script through Query Interpreter

```

show(cg(functions("dswrite")))

```

Script: Data Flow

Atlas Shell (Project: AtlasToolbox)

```

var y = cg(x)
y: com.ensoftcorp.atlas.core.query.Q = <Atlas query expression>
show(y)

```

Evaluate: show(cg(functions("dswrite")))

Element Detail View

- dswrite
 - @ XCSG.ModelElement.id = [Xinu/Xinu/sys/dswrite.c#dswrite(struct devsw*, char*, DBADDR)]
 - @ XCSG.ModelElement.name = dswrite
 - @ XCSG.ModelElement.sourceCorrespondence = [322+7+/Xinu/sys/dswrite.c]
 - ##index
 - external
 - isDef
 - XCSG.Function
 - XCSG.Language.C

Project Explorer

Graph 3 Graph 4 Graph 5 Graph 6 x_net.c Graph 7 28

Atlas Connection View Atlas Smart View

arpfind.c arpinit.c blkcopy.c blkequ.c bpdump.c chprio.c ckmode.c cksum.s clkinit.c clkint.c clkint.s close.c conf.c control.c create.c csv.s ctxsw.c ctxsw.s devdump.c dfalloc.c dfdsrch.c dgalloc.c dgclose.c dgctl.c dgdump.c dginit.c dgmcntl.c dgopen.c dgparse.c dgread.c dgwrite.c dot2ip.c dscntl.c dsinit.c dsinter.c dskbcpy.c dskkenq.c dskqopt.c dskstrt.c dsksync.c dsopen.c dsread.c dsseek.c dswrite.c dudir.c

Resulting Call Graph (CG)

Script: Data Flow

Problems Tasks Console Properties Atlas Shell (AtlasToolbox)

```

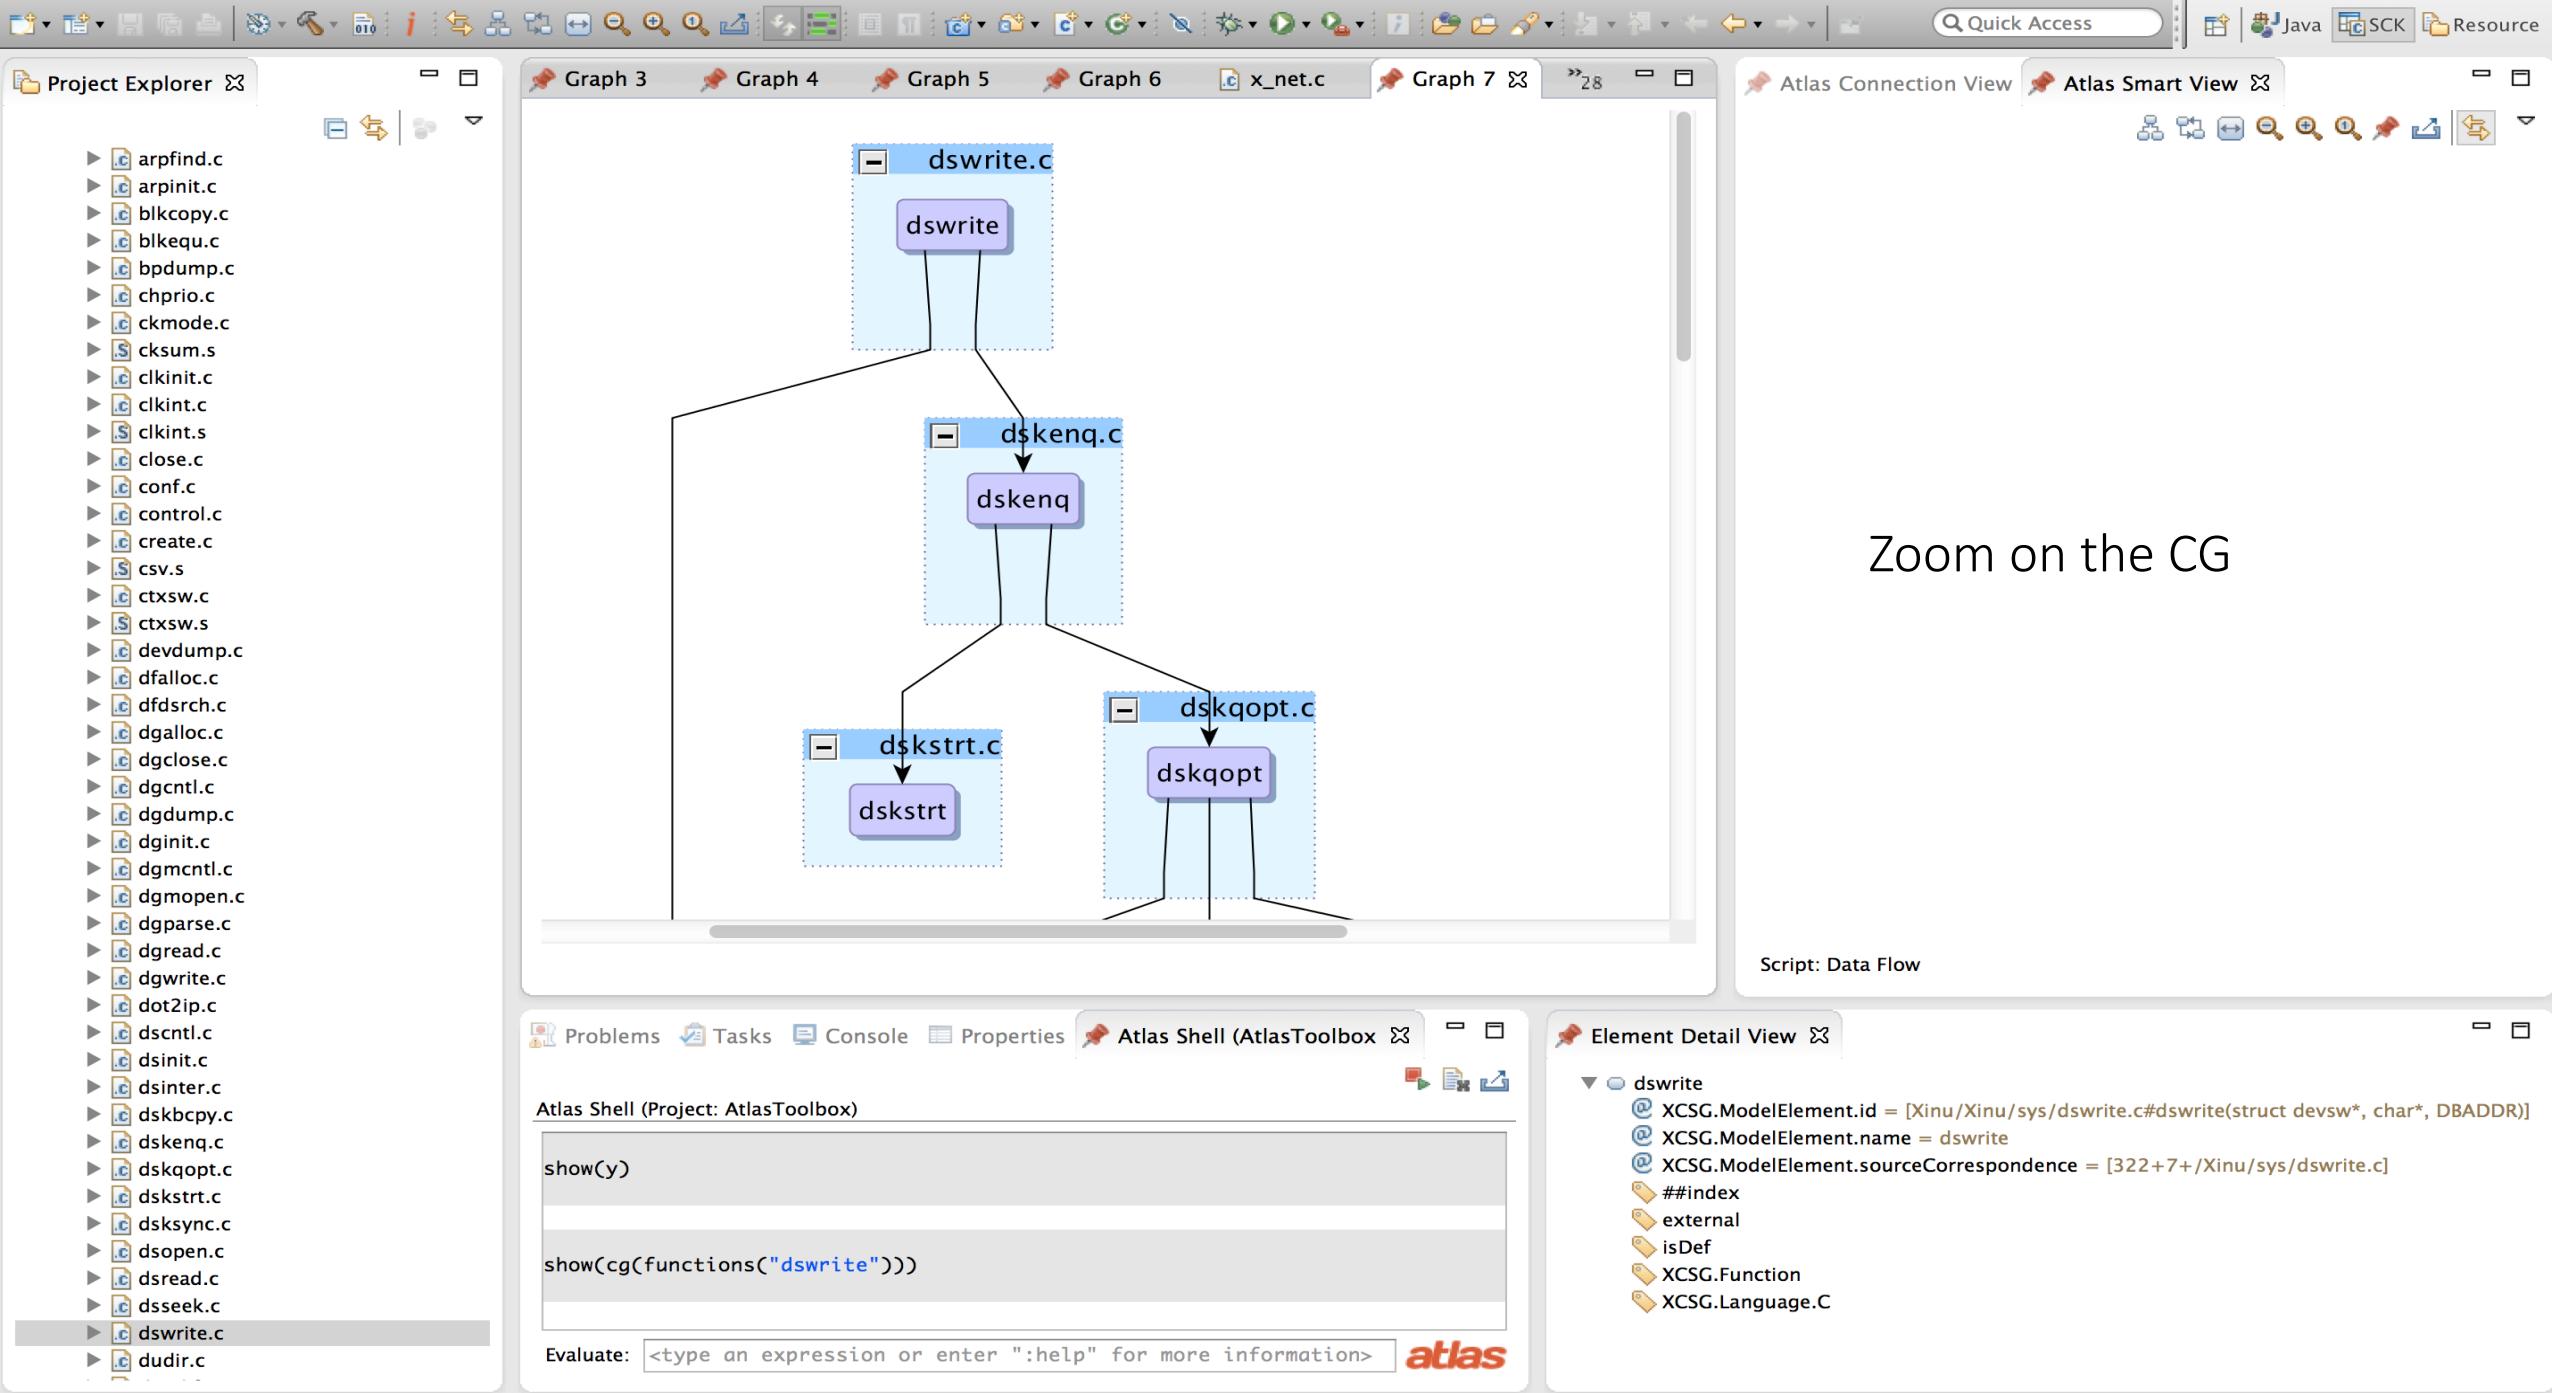
show(y)
show(CG(functions("dswrite")))

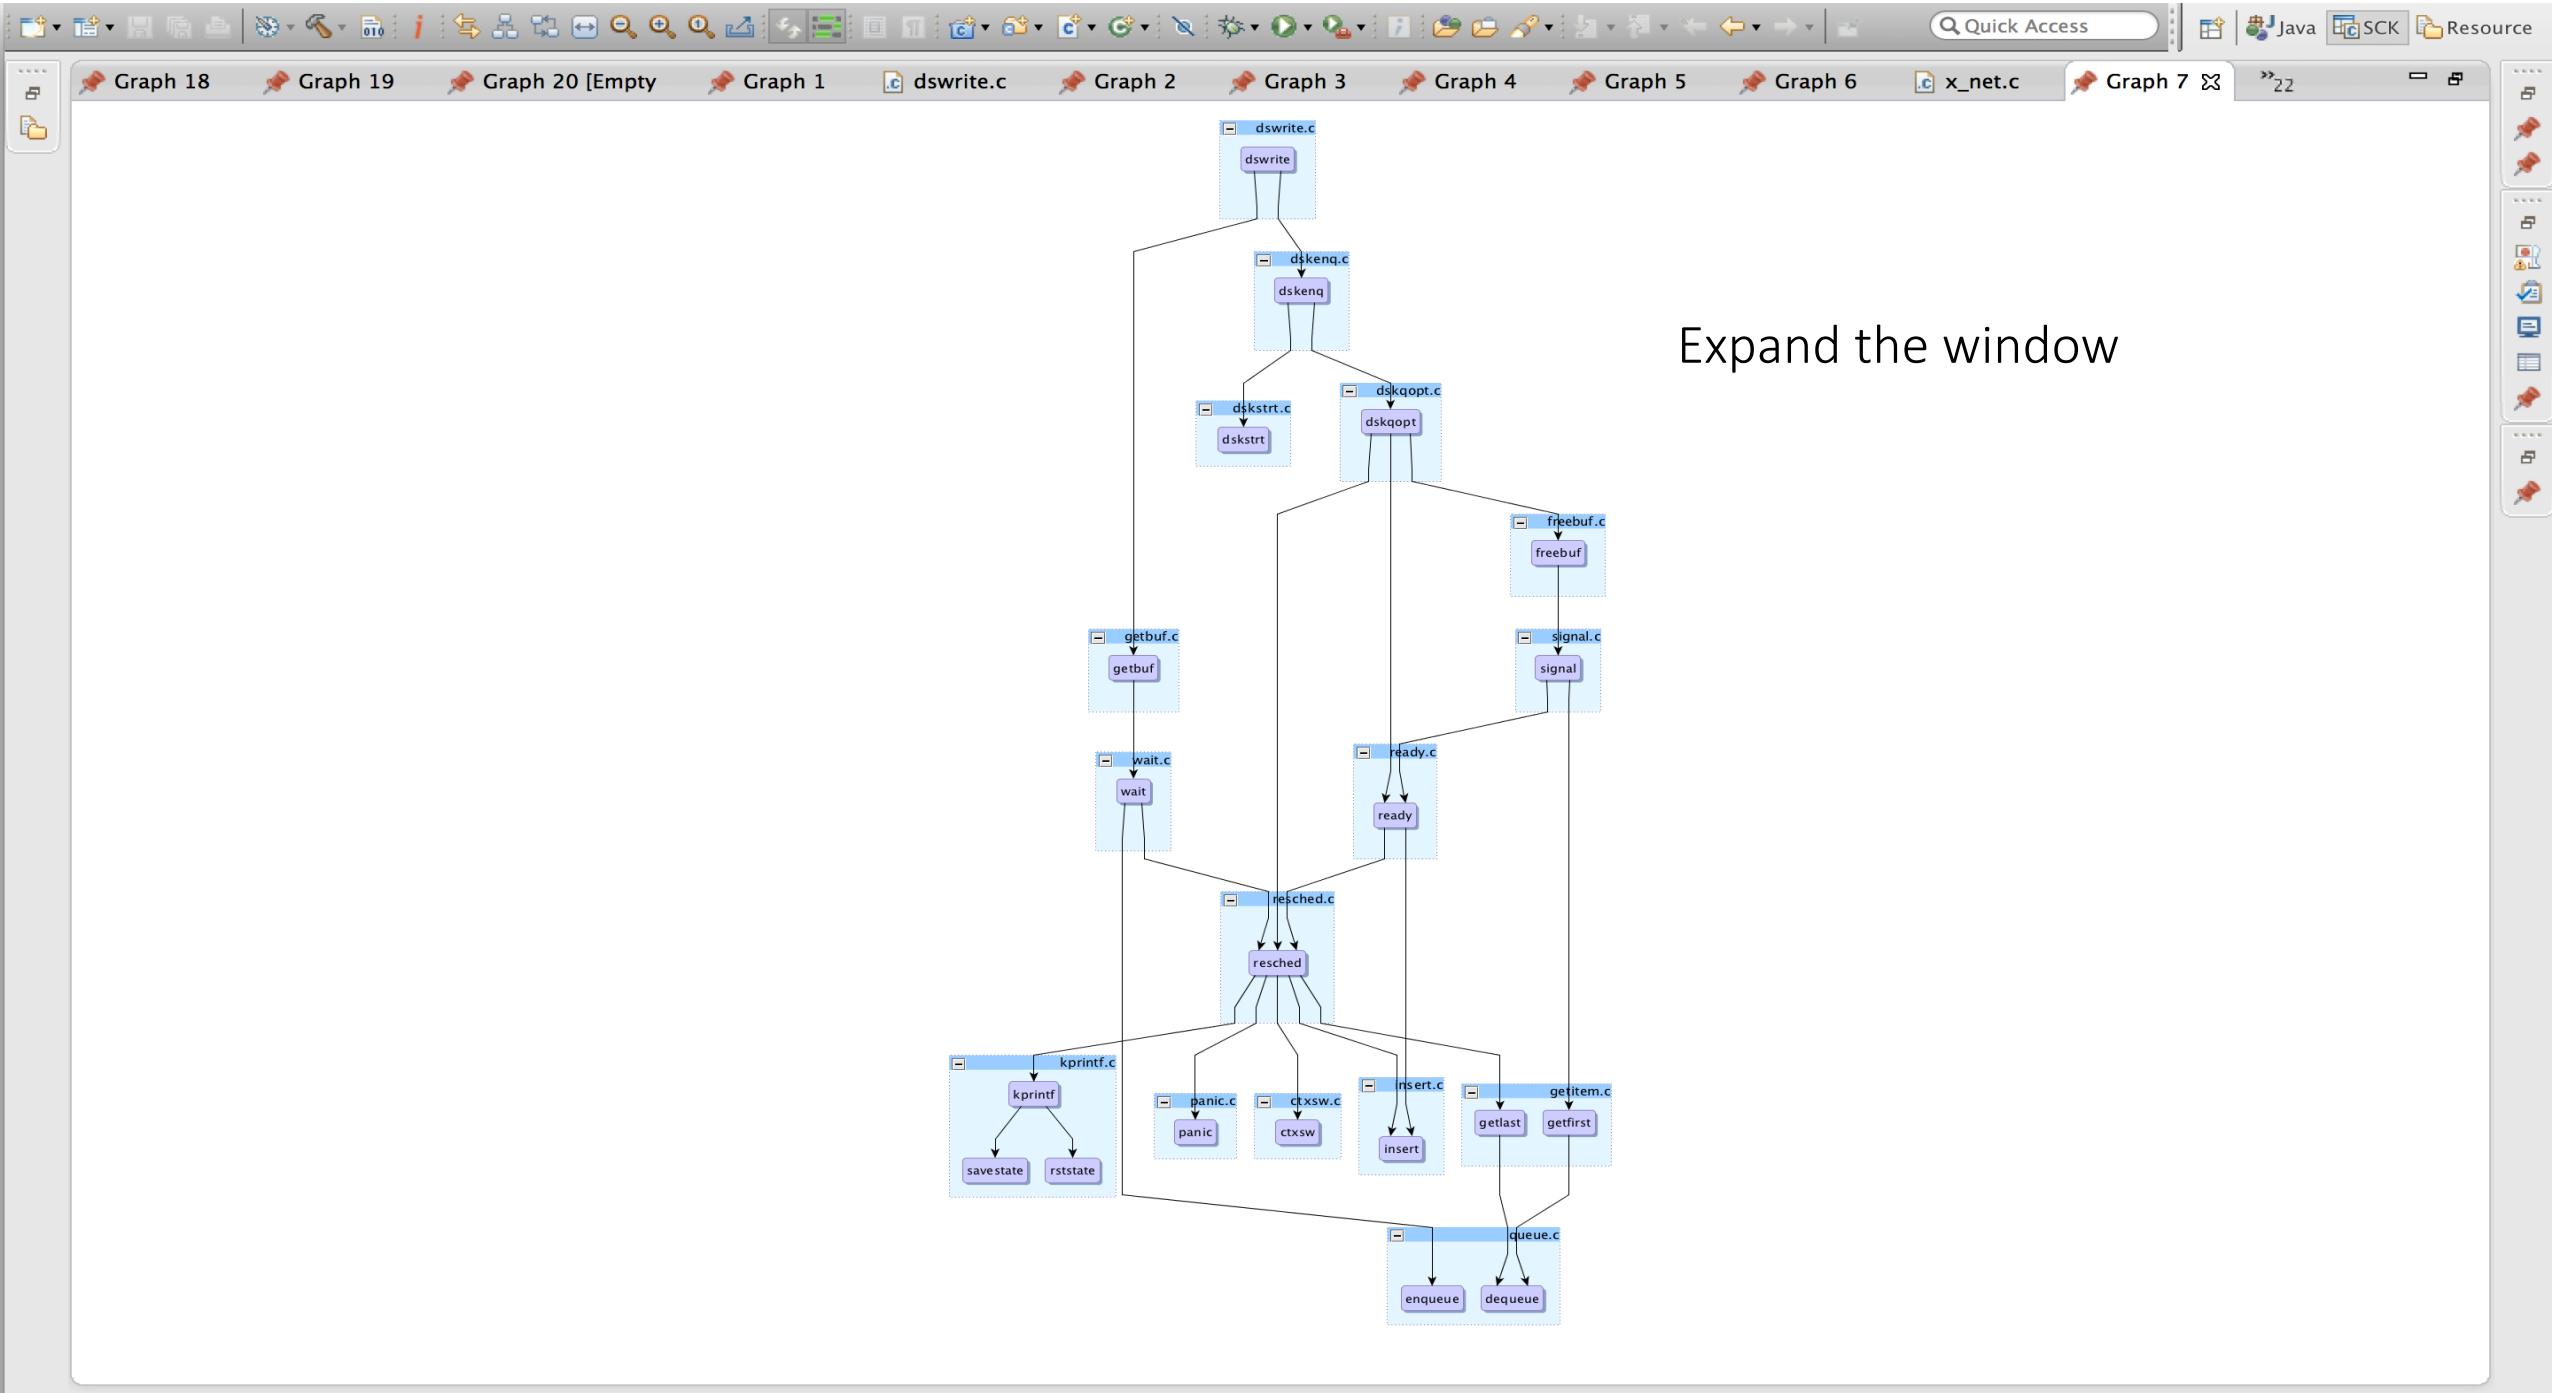
```

Evaluate: <type an expression or enter ":help" for more information>

Element Detail View

- dswrite
 - XCSG.ModelElement.id = [Xinu/Xinu/sys/dswrite.c#dswrite(struct devsw*, char*, DBADDR)]
 - XCSG.ModelElement.name = dswrite
 - XCSG.ModelElement.sourceCorrespondence = [322+7+/Xinu/sys/dswrite.c]
 - ##index
 - external
 - isDef
 - XCSG.Function
 - XCSG.Language.C





Project Explorer

- arpfind.c
- arpinit.c
- blkcopy.c
- blkequ.c
- bpdump.c
- chprioc.c
- ckmode.c
- cksum.s
- clkinit.c
- clkint.c
- clkint.s
- close.c
- conf.c
- control.c
- create.c
- csv.s
- ctxsw.c
- ctxsw.s
- devdump.c
- dfalloc.c
- dfdsrch.c
- dgalloc.c
- dgclose.c
- dgctl.c
- ddgdump.c
- dginit.c
- dgmctrl.c
- dgmopen.c
- dgparse.c
- dgread.c
- dgwrite.c
- dot2ip.c
- dscntl.c
- dsinit.c
- dsinter.c
- dskbcpy.c
- dskkenq.c
- dskqopt.c
- dskstrt.c
- dsksync.c
- dsopen.c
- dsread.c
- dsseek.c
- dswrite.c
- dudir.c

Graph 7 Graph 8 Graph 9 Graph 10 Graph 11

Queries by selecting model elements

Select a node from CG

show(CG(selected))

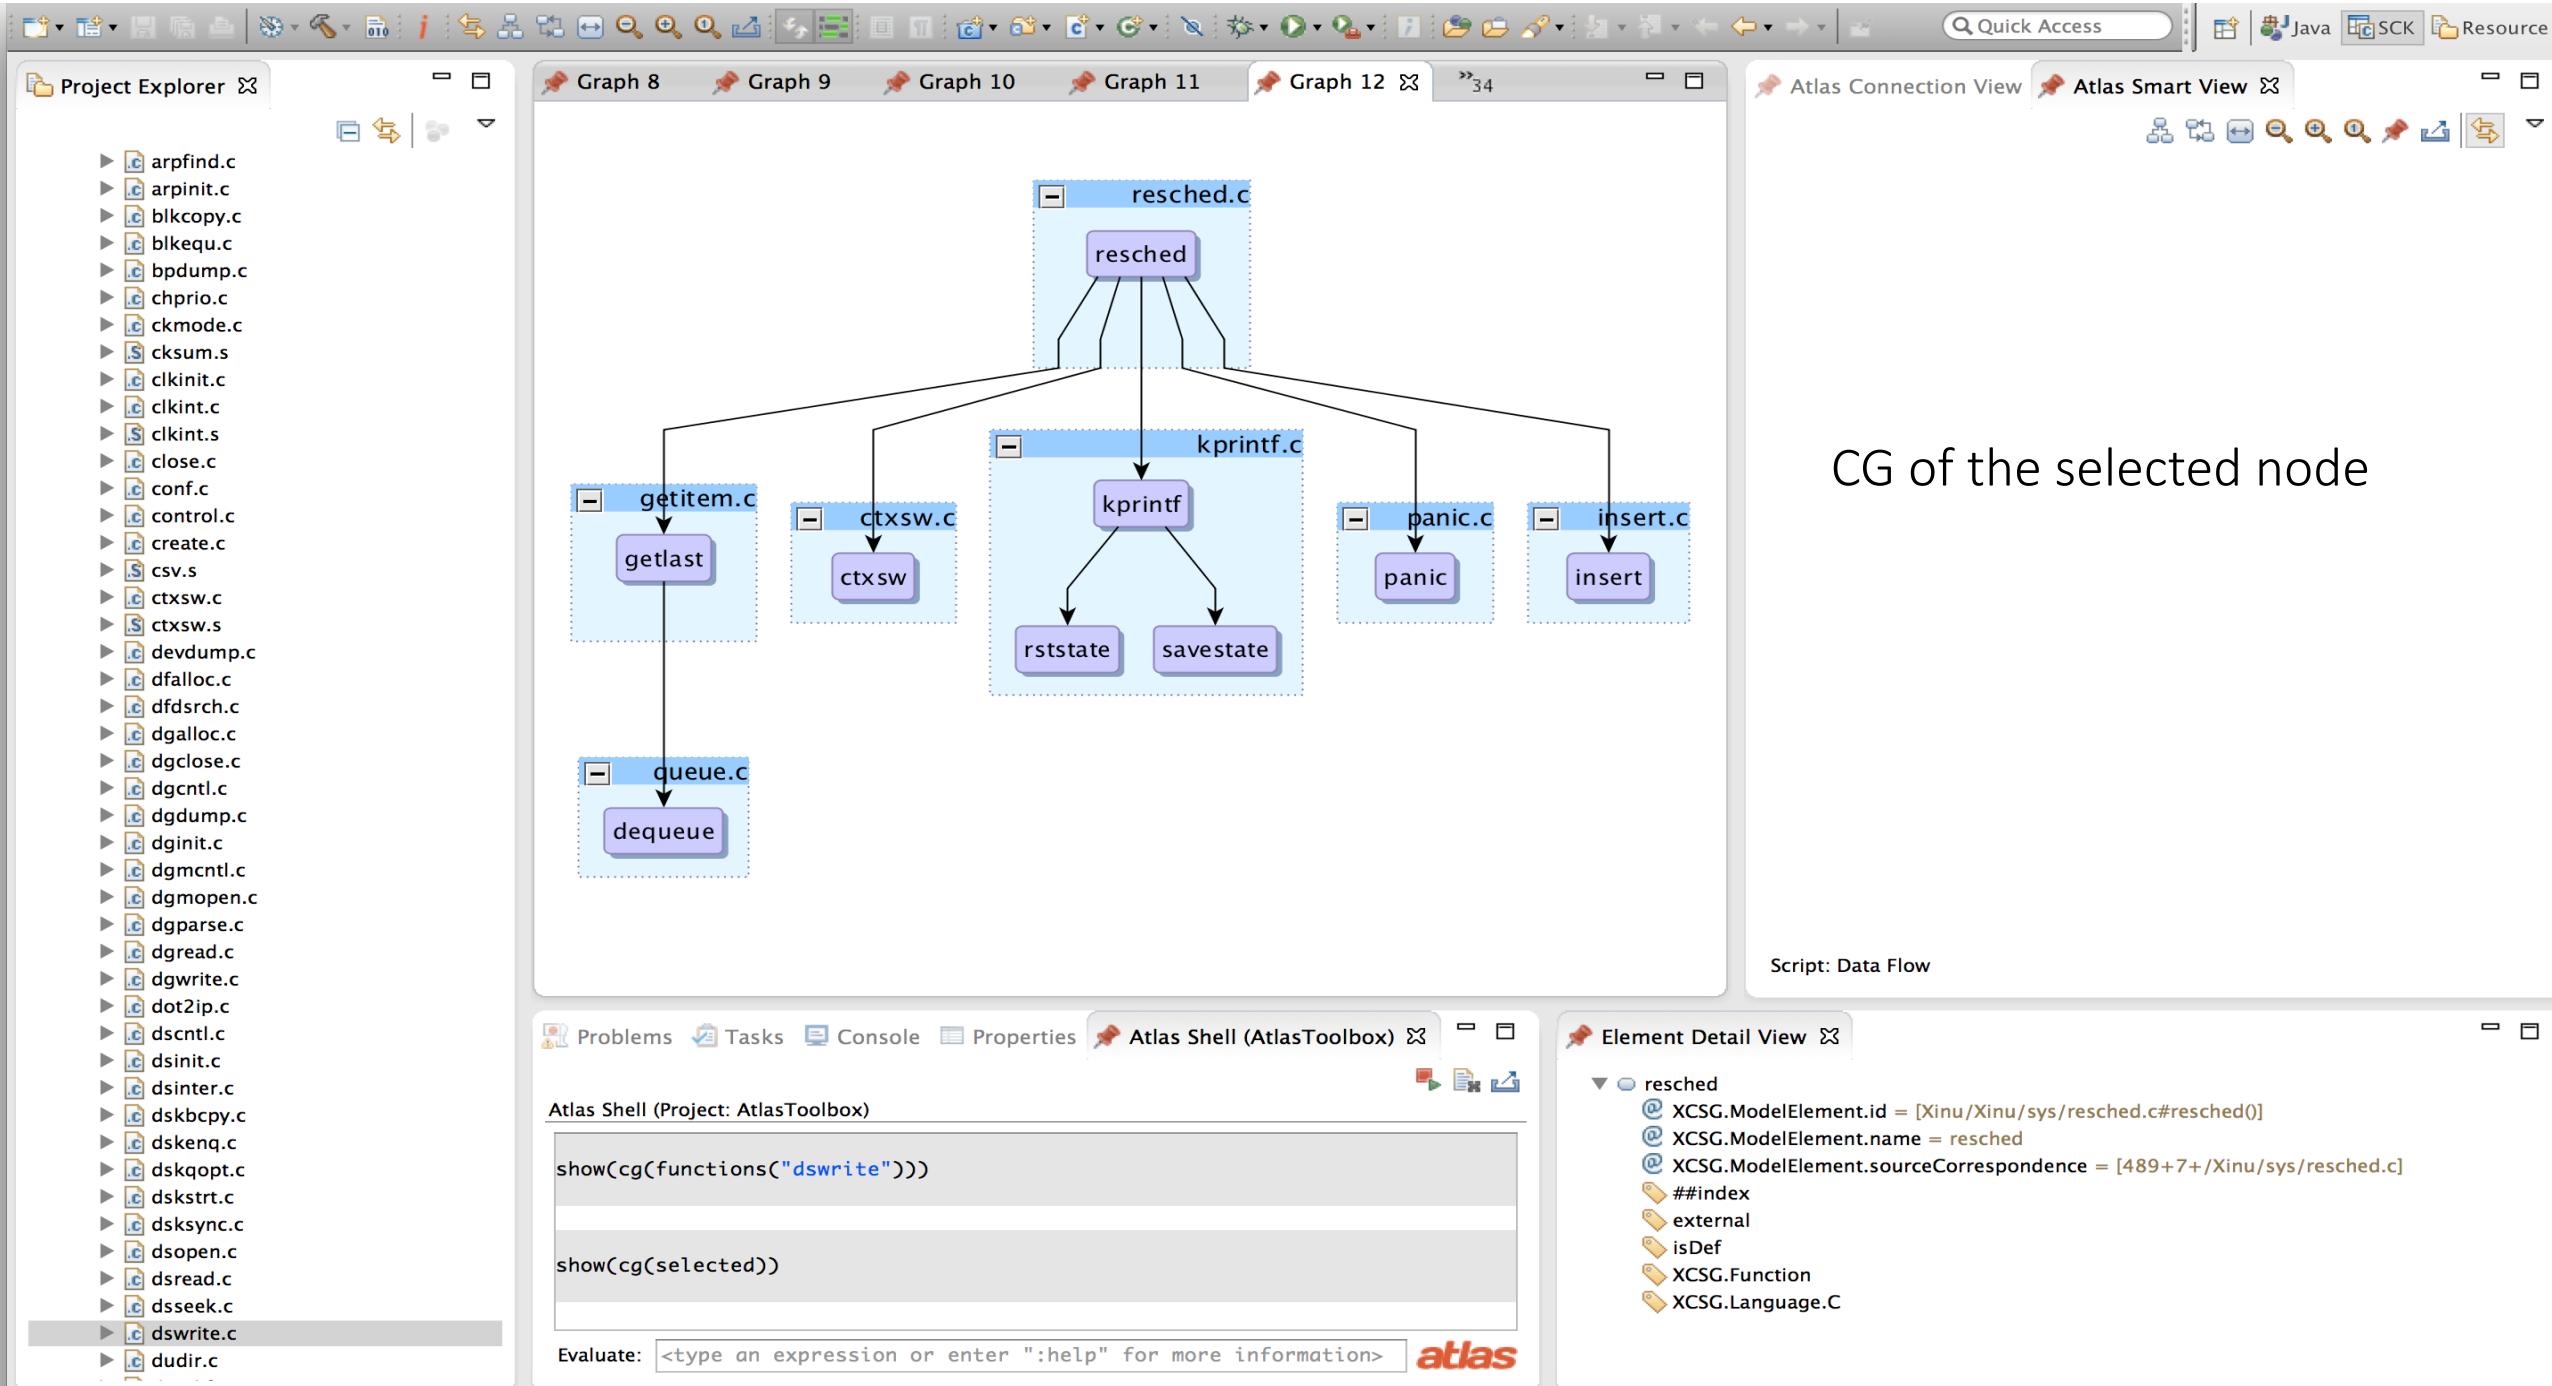
Script: Data Flow

Atlas Shell (Project: AtlasToolbox)

```
show(CG(selected))
show(CG(functions("dswrite")))
Evaluate: <type an expression or enter ":help" for more information>
```

Element Detail View

- resched
 - XCSG.ModelElement.id = [Xinu/Xinu/sys/resched.c#resched()]
 - XCSG.ModelElement.name = resched
 - XCSG.ModelElement.sourceCorrespondence = [489+7+/Xinu/sys/resched.c]
 - ##index
 - external
 - isDef
 - XCSG.Function
 - XCSG.Language.C



Smart View

Select a model for smart view (e.g. RCG)

Click a model node

Selected model appears here

Smart View Pane

Script: Reverse Call

Element Detail View

```

resched
  @ XCSCG.ModelElement.id = [Xinu/Xinu/sys/resched.c#resched()]
  @ XCSCG.ModelElement.name = resched
  @ XCSCG.ModelElement.sourceCorrespondence = [489+7+/Xinu/sys/resched.c]
    ##index
    external
    isDef
    XCSCG.Function
    XCSCG.Language.C
  
```

Project Explorer

Graph 9 Graph 10 Graph 11 Graph 12 Graph 13 >35

Atlas Connection View Atlas Smart View

Call
Call Step
Reverse Call
Forward Call
Data Flow
Data Flow Step
Reverse Data Flow
Forward Data Flow
Instantiation
References
Declarations
Type Hierarchy

Problems Tasks Console Properties Atlas Shell (AtlasToolbox)

```

show(CG(selected))
show(CG(functions("dswrite")))
  
```

Evaluate: <type an expression or enter ":help" for more information>

Smart view generates RCGs on the fly as we select different nodes

The screenshot shows the Atlas IDE interface with several windows:

- Project Explorer:** Lists numerous C source files (e.g., arpfind.c, arpinit.c, blkcopy.c, etc.) and header files.
- Graph 13:** Displays a complex call graph with many nodes and edges, representing the full system's control flow.
- Atlas Connection View:** Shows a simplified view of selected nodes: `dsseek`, `dswrite`, and `dsread`. Each node has a corresponding function name below it: `dskenq`.
- Atlas Shell (AtlasToolbox):** Contains a script window with the following content:


```
show(CG(selected))
show(CG(functions("dswrite")))
```
- Element Detail View:** Provides detailed information about the selected node `dskenq`, including its ID, name, source correspondence, and various properties like index, external, isDef, XCSG.Function, and XCSG.Language.C.

Project Explorer

Graph 9 Graph 10 Graph 11 Graph 12 Graph 13

Atlas Connection View Atlas Smart View

Smart view generates RCGs on the fly as we select different nodes

dsread.c

dsseek.c

dswrite.c

Script: Reverse Call

Problems Tasks Console Properties

Atlas Shell (AtlasToolbox)

```
show(CG(selected))
show(CG(functions("dswrite")))
```

Evaluate: <type an expression or enter ":help" for more information>

Element Detail View

dskqopt

- XCSG.ModelElement.id = [Xinu/Xinu/sys/dskqopt.c#dskqopt(struct dreq*, struct dreq*, struct freebuf*)]
- XCSG.ModelElement.name = dskqopt
- XCSG.ModelElement.sourceCorrespondence = [310+7+/Xinu/sys/dskqopt.c]
- ##index
- external
- isDef
- XCSG.Function
- XCSG.Language.C

Smart view generates RCGs on the fly as we select different nodes

Experiment and invent new visual models:

1. As we are viewing RCGs for different nodes, we discover a node with a large RCG.
2. What does a large RCG tell about the node?
3. How can we use that knowledge to refine visual models?

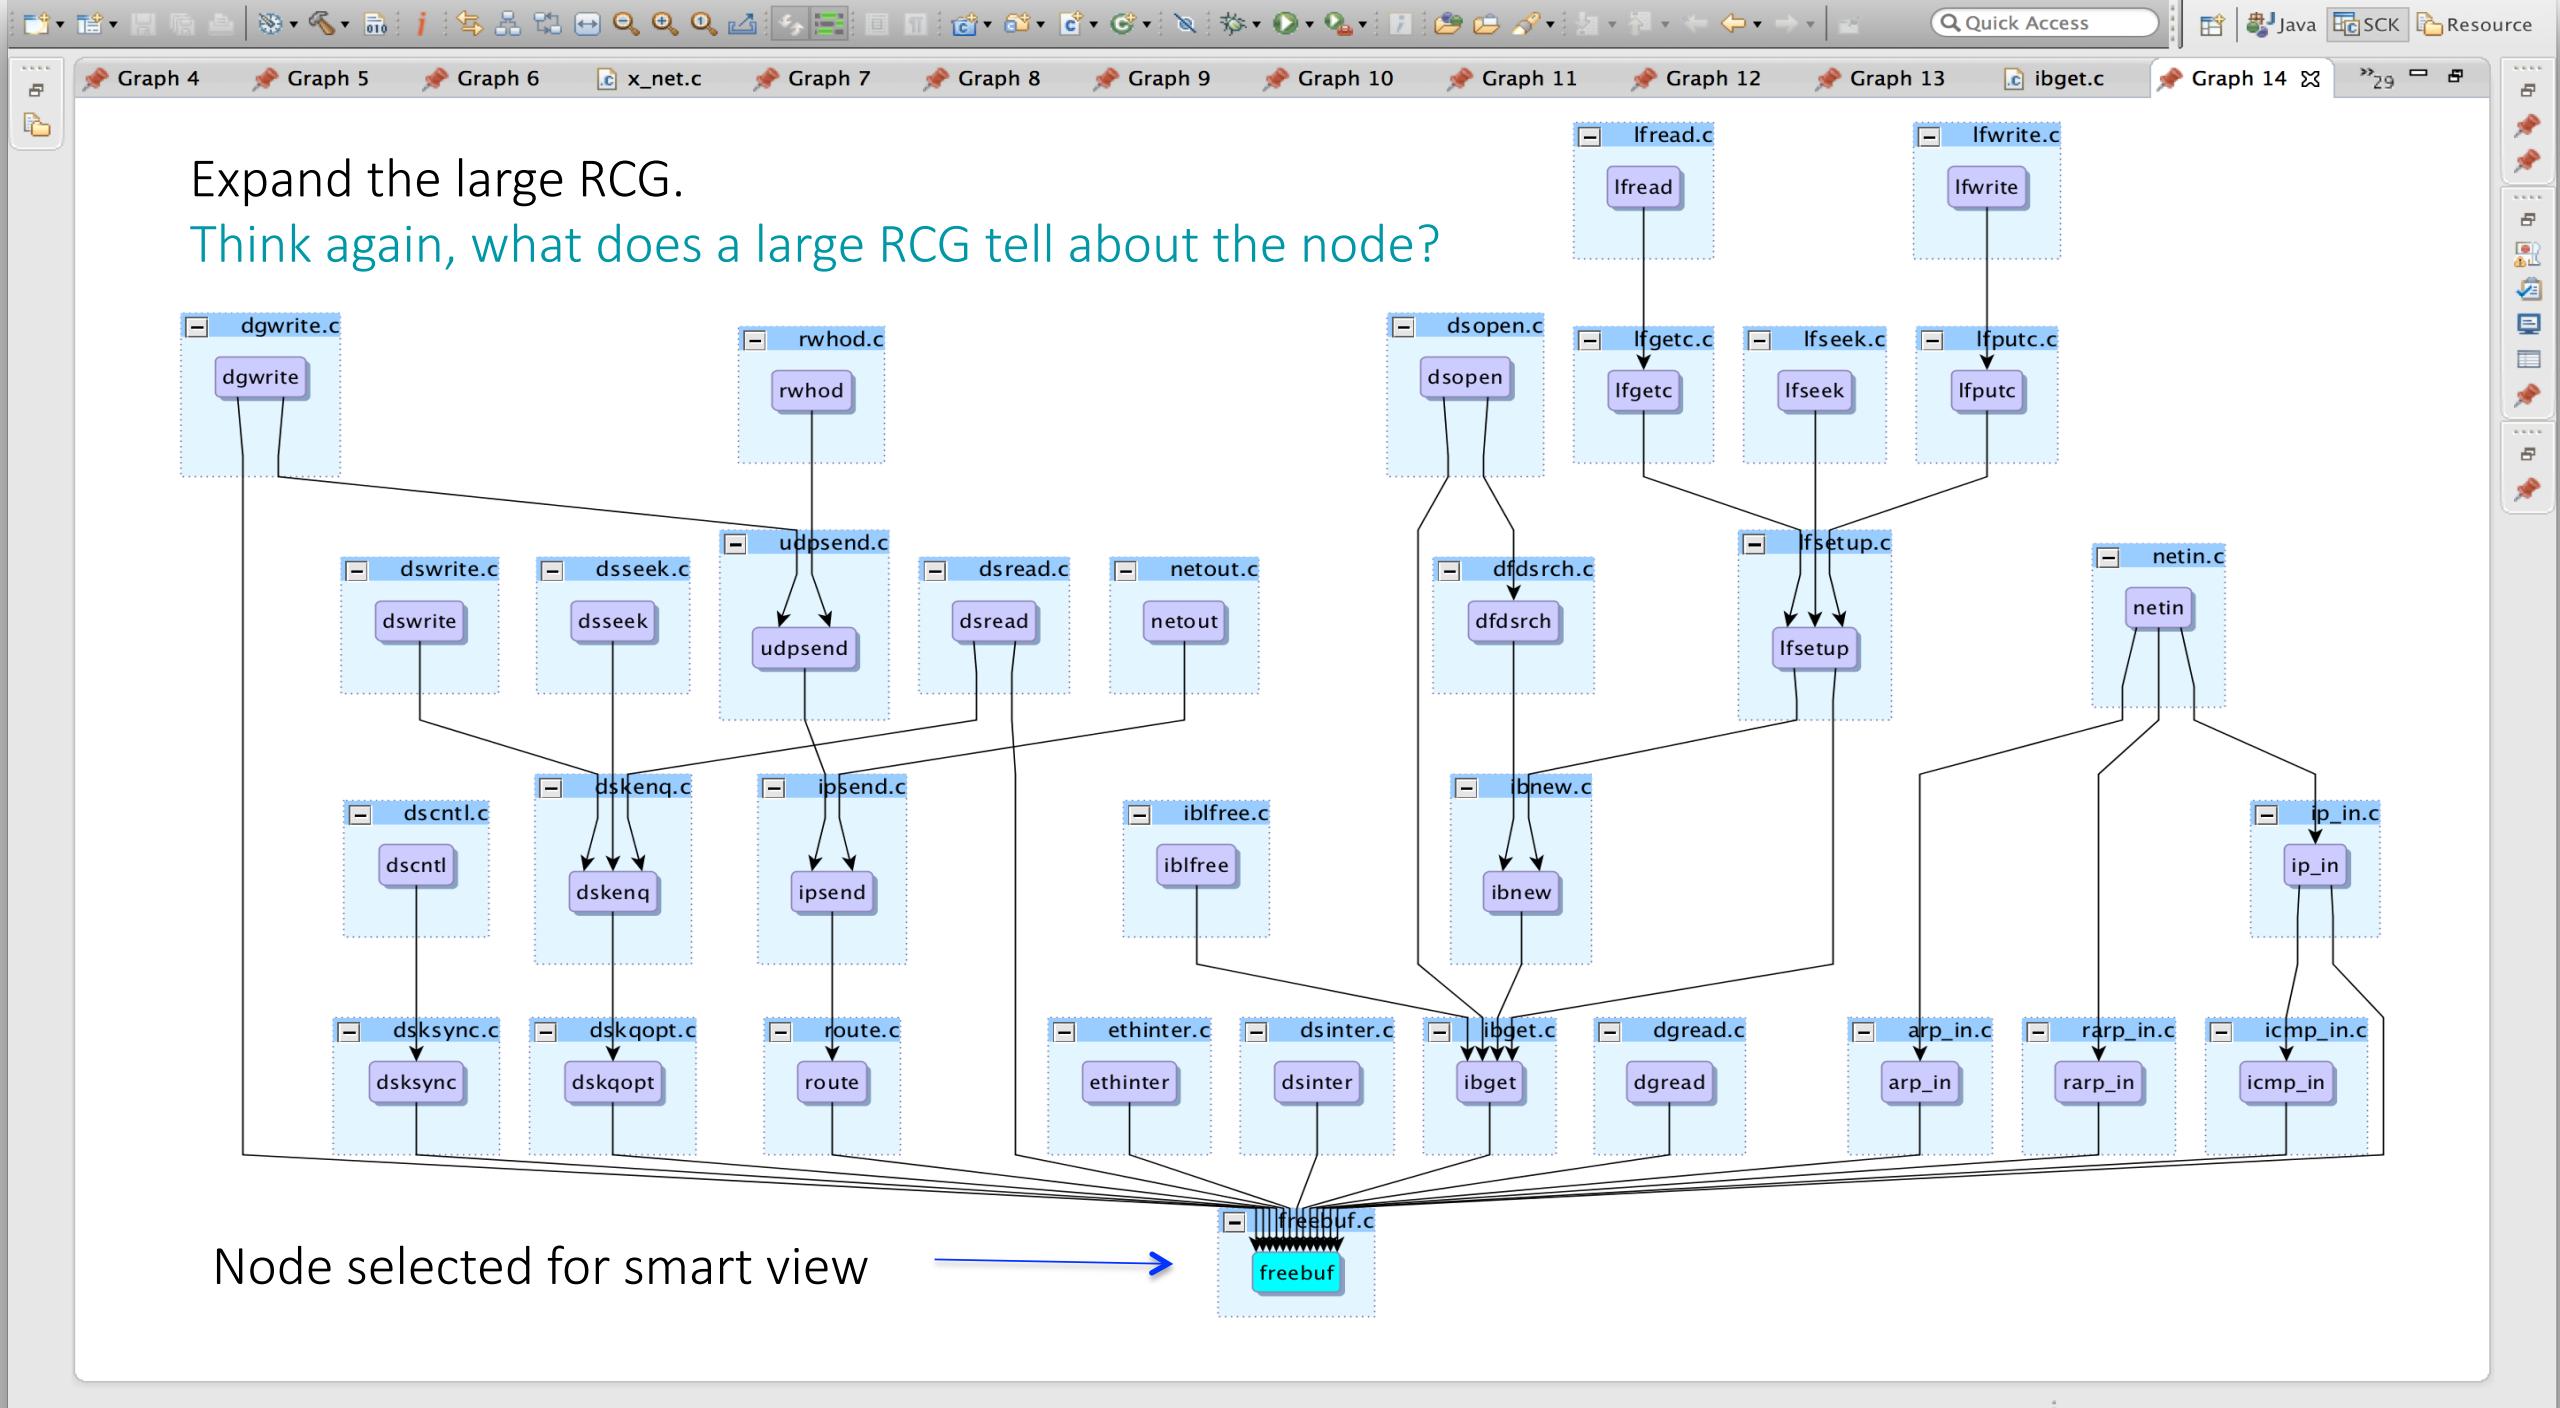
Atlas Shell (Project: AtlasToolbox)

```
show(CG(selected))
show(CG(functions("dswrite")))
```

Evaluate: <type an expression or enter ":help" for more information>

Element Detail View

- freebuf
 - @ XCSG.ModelElement.id = [Xinu/Xinu/sys/freebuf.c#freebuf(int*)]
 - @ XCSG.ModelElement.name = freebuf
 - @ XCSG.ModelElement.sourceCorrespondence = [339+7+/Xinu/sys/freebuf.c]
 - ##index
 - external
 - isDef
 - XCSG.Function
 - XCSG.Language.C



Project Explorer

Graph 10 Graph 11 Graph 12 Graph 13 ibget.c 36

```
/* ibget.c - ibget */
#include <conf.h>
#include <kernel.h>
#include <disk.h>
#include <lfile.h>
#include <dir.h>

* ibget -- get an iblock from disk given its number.
ibget(diskdev, inum, loc)
int diskdev;
IBADDR inum;
struct iblk *loc;
{
    char *from, *to;
    int i;
    char *buff;

    buff = getbuf(dskdbp);
    read(diskdev, buff, ibtodb(inum));
    from = buff + ibdisp(inum);
    to = (char *)loc;
    for (i=0 ; i<sizeof(struct iblk) ; i++)
        *to++ = *from++;
    freebuf(buff);
}
```

Atlas Connection View Atlas Smart View

See the corresponding code by clicking on nodes

Script: Reverse Call

Problems Tasks Console Properties Atlas Shell (AtlasToolbox)

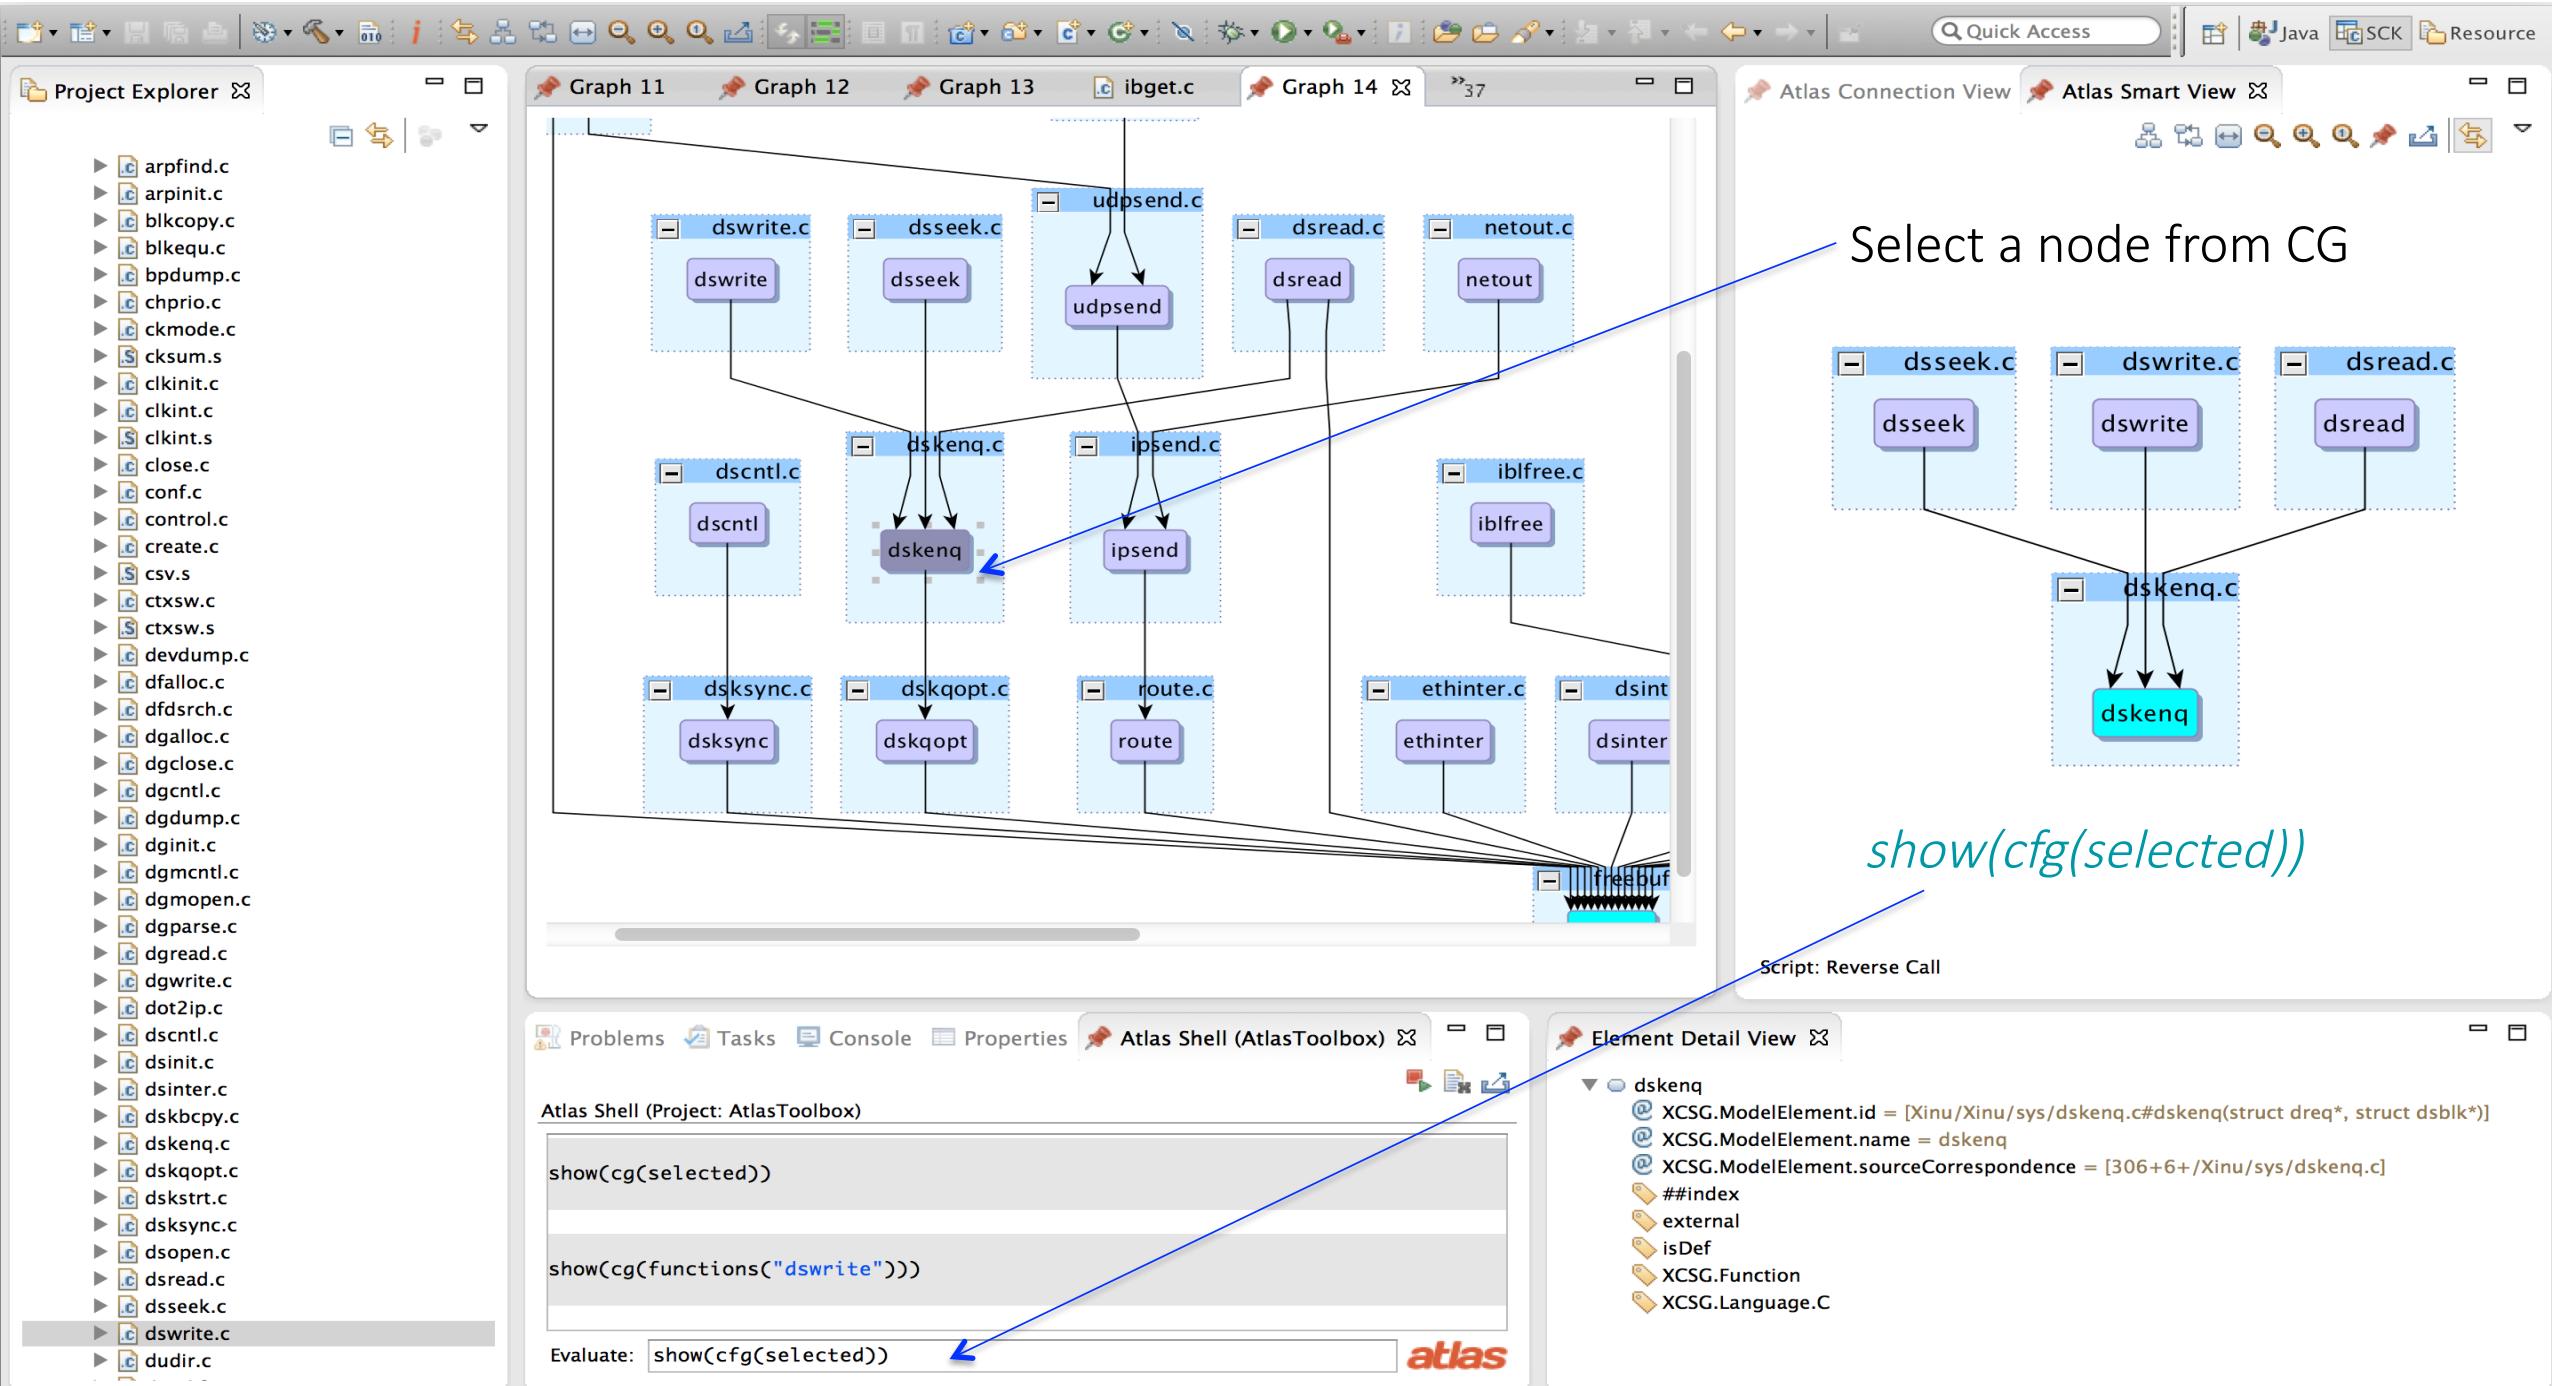
```
show(CG(selected))
show(CG(functions("dswrite")))
```

Evaluate: <type an expression or enter ":help" for more information> atlas

Element Detail View

- ibget.c
 - @ XCSG.ModelElement.id = [Xinu/Xinu/sys/ibget.c]
 - @ XCSG.ModelElement.name = ibget.c
 - @ XCSG.ModelElement.sourceCorrespondence = [0+624+/Xinu/sys/ibget.c]
 - ##Index
 - XCSG.C.TranslationUnit
 - XCSG.Language.C

Writable Smart Insert 1 : 1



Project Explorer

- arpfind.c
- arpinit.c
- blkcopy.c
- blkreq.c
- bpdump.c
- chprio.c
- ckmode.c
- cksum.s
- clkinit.c
- clkint.c
- clkint.s
- close.c
- conf.c
- control.c
- create.c
- csv.s
- ctxsw.c
- ctxsw.s
- devdump.c
- dalloc.c
- dfdsrch.c
- dgalloc.c
- dgclose.c
- dgcntl.c
- gdump.c
- dginit.c
- dgmctl.c
- dgmopen.c
- dgparse.c
- dgread.c
- dgwrite.c
- dot2ip.c
- dscntl.c
- dsinit.c
- dsinter.c
- dskbcpv.c
- dskenq.c
- dskqopt.c
- dskstrt.c
- dsksync.c
- dsopen.c
- dsread.c
- dsseek.c
- dswrite.c
- dudir.c

Graph 12 Graph 13 ibget.c Graph 14 Graph 15 »38

Atlas Connection View Atlas Smart View

Control Flow Graph (CFG) of selected node

```

graph TD
    Start(( )) --> DREQStruct[struct dreq *p, *q]
    DREQStruct --> DBADDR[DBADDR block]
    DBADDR --> INT[int st]
    INT --> Cond{if ((q == drptr->dreqst) == (drq->q))}
    Cond --> DRPTRSet[drptr->dreqst = drptr;]
    DRPTRSet --> DRNNULL[drptr->drnext = DRNULL;]
    DRNNULL --> DSSTRIDSPTR[dskstridsptr;]
    DSSTRIDSPTR --> Return2[return (2) ->]
    DSSTRIDSPTR --> BlockSet[block = drptr->drdb;]
    BlockSet --> ForInit[for init]
    ForInit --> ForLoop{for (p = (dreq->q))}
    ForLoop --> Cond2{if (p->drdb == block && (st == dsqopt(p, q) || drptr != -1))}
    Cond2 --> ReturnSt[return (st) ->]
    Cond2 --> IfBlock{if ((q->drdb <= block && block < p->drdb) || (q->drdb > block && block > p->drdb))}
    IfBlock --> DRPRTUpdate[drptr->drnext = p;]
    IfBlock --> DRPRTUpdate[drptr->drnext = drptr;]
    DRPRTUpdate --> DRNNULLSet[q->drnext = DRNULL;]
    DRNNULLSet --> Return2[return (2) ->]
    IfBlock --> Update[for update]
    Update --> DRPRTUpdate[drptr->drnext = drptr;]
    DRPRTUpdate --> DRNNULLSet[q->drnext = DRNULL;]
    DRNNULLSet --> Return2[return (2) ->]

```

dsseek.c dswrite.c dsread.c

dsseek dswrite dsread

dskenq

Script: Reverse Call

Problems Tasks Console Properties Atlas Shell (AtlasToolbox)

```

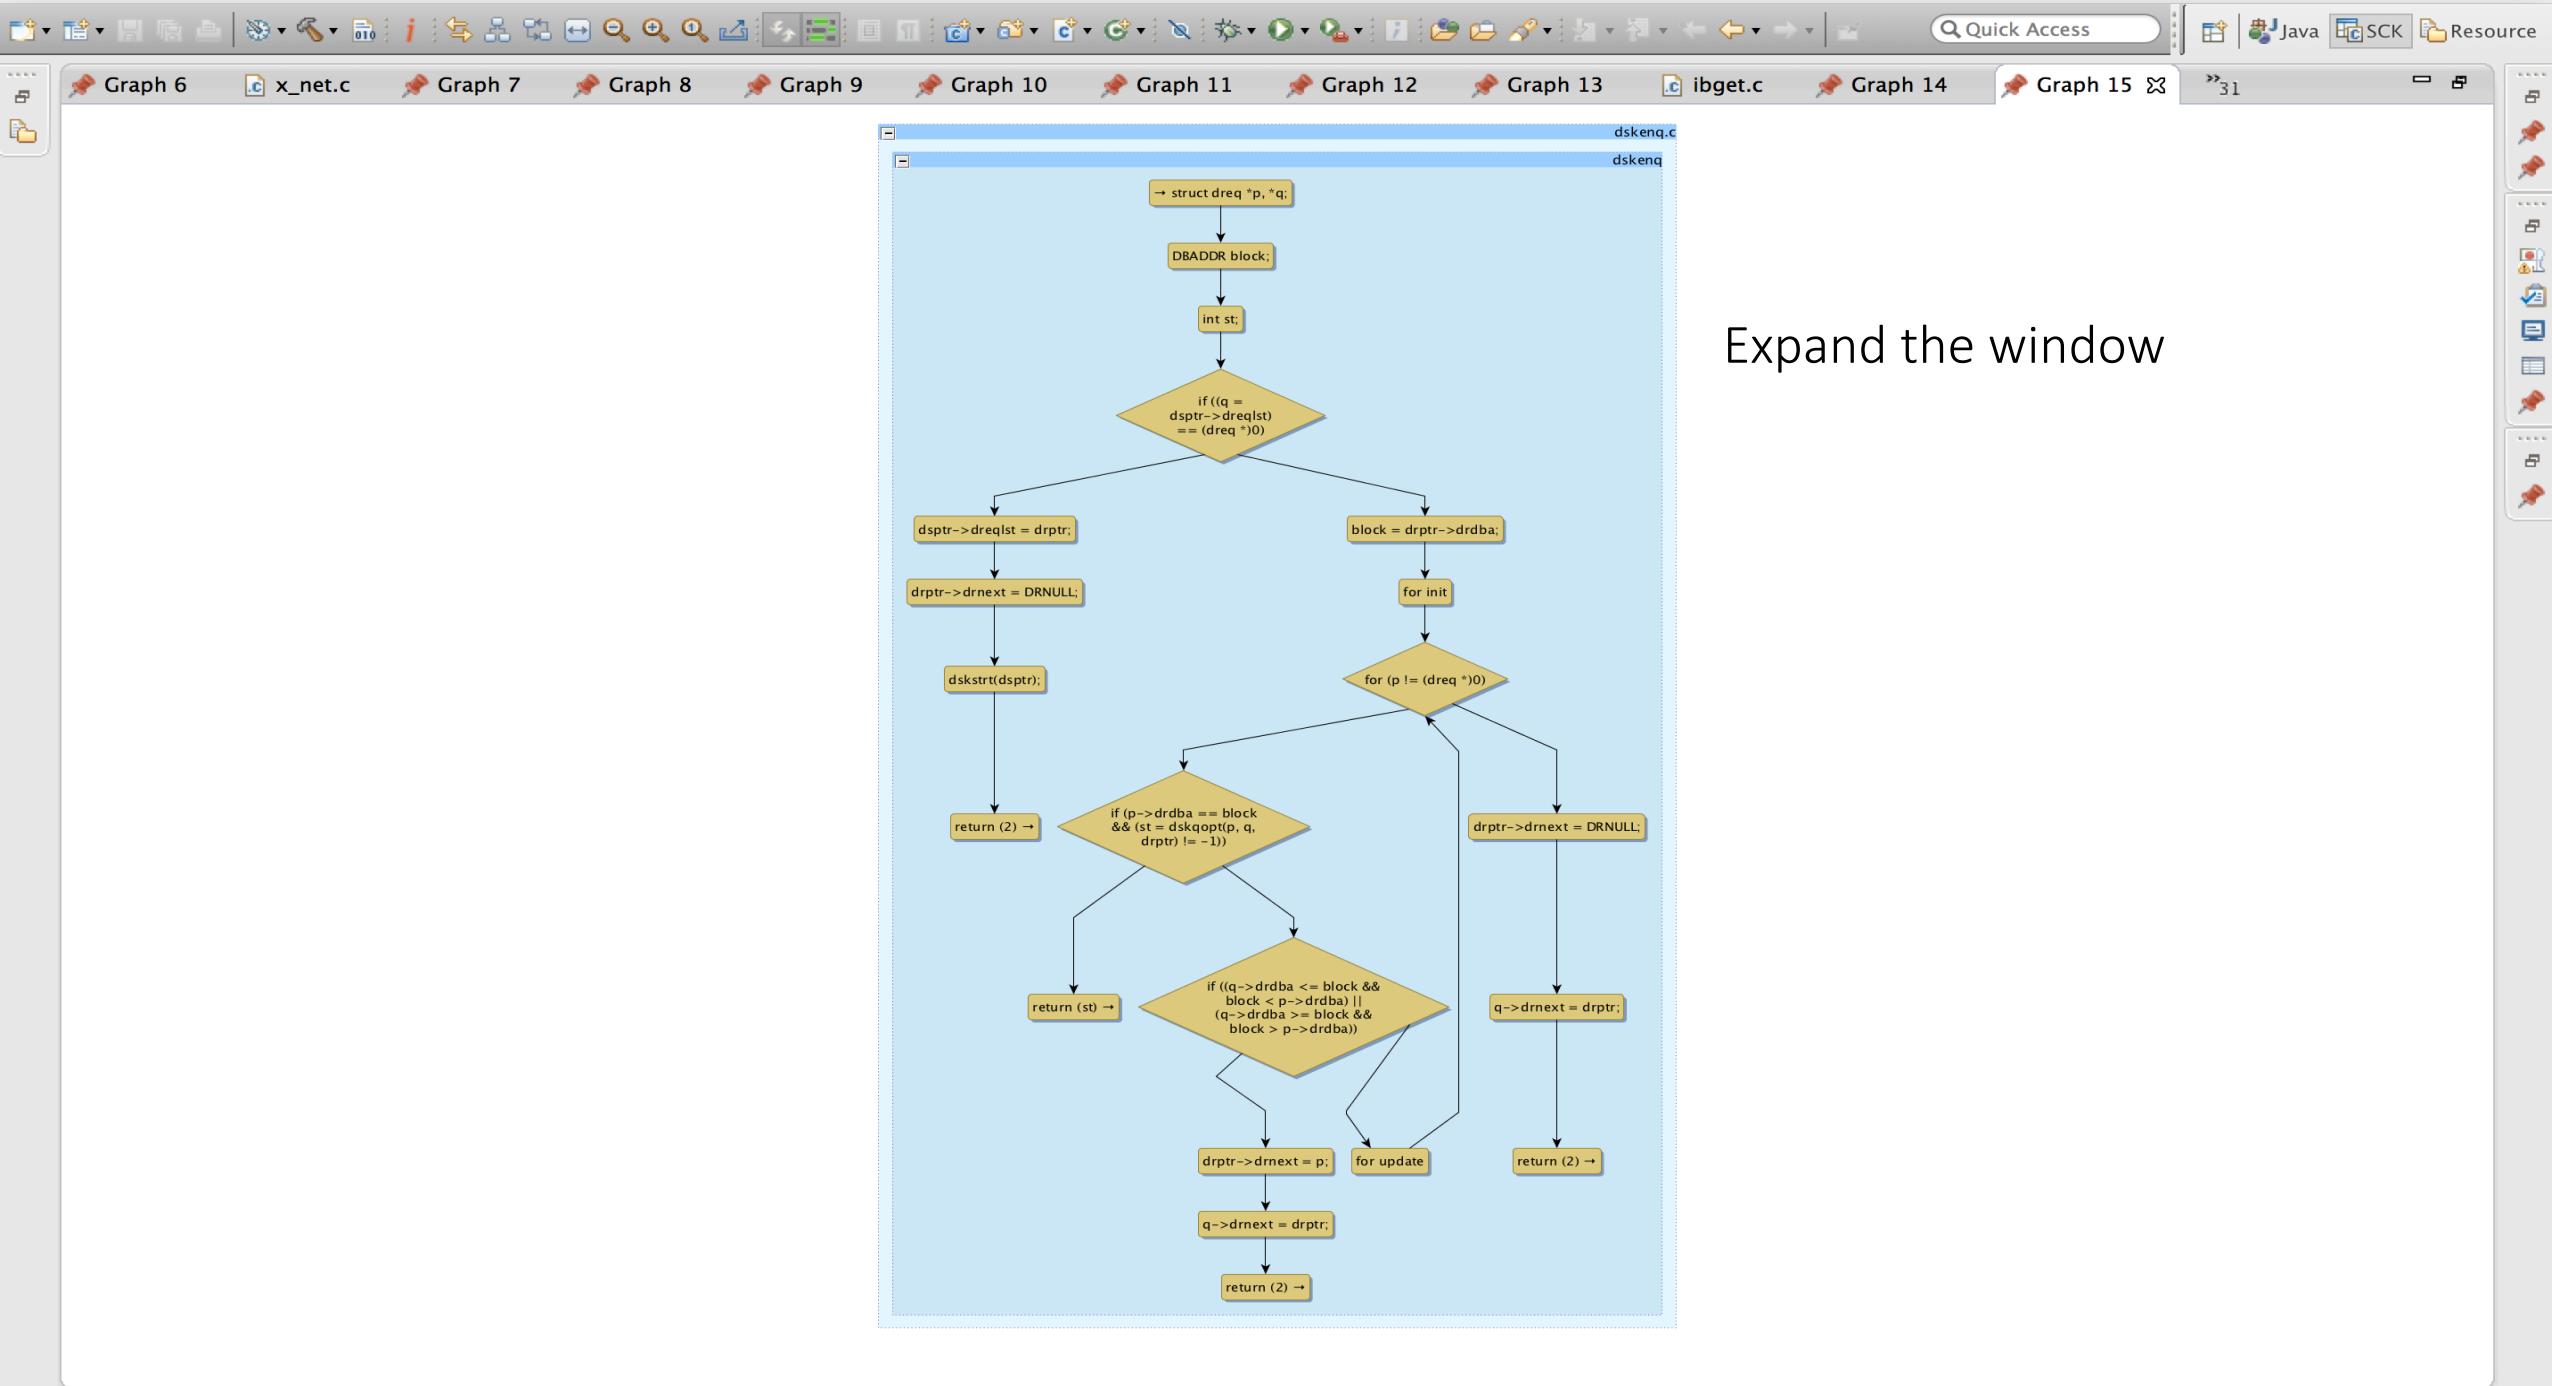
show(cfg(functions("dswrite")))
show(cfg(selected))

```

Evaluate: <type an expression or enter ":help" for more information>

Element Detail View

- dskenq
 - XCSG.ModelElement.id = [Xinu/Xinu/sys/dskenq.c#dskenq(struct dreq*, struct dsblk*)]
 - XCSG.ModelElement.name = dskenq
 - XCSG.ModelElement.sourceCorrespondence = [306+6+/Xinu/sys/dskenq.c]
 - ##index
 - external
 - isDef
 - XCSG.Function
 - XCSG.Language.C



Project Explorer

- bpdump.c
- chprio.c
- ckmode.c
- cksum.s
- clkinit.c
- clkint.c
- clkint.s
- close.c
- conf.c
- control.c
- create.c
- csv.s
- ctxsw.c
- ctxsw.s
- devdump.c
- dfalloc.c
- dfdsrc.h
- dgalloc.c
- dgclose.c
- dgctrl.c
- dgdump.c
- dginit.c
- dgmctrl.c
- dgmopen.c
- dgparse.c
- dgread.c
- dgwrite.c
- dot2ip.c
- dscnt.c
- dsinit.c
- dsinter.c
- dskbcpy.c
- dskenq.c
- dskqopt.c
- dskstrt.c
- dsksync.c
- dsopen.c
- dsread.c
- dsseek.c
- dswrite.c**
- dudir.c
- dumkfs.c
- ethinit.c
- ethinter.c
- ethread.c
- ethtrans.c

dswrite.c

```
/* dswrite.c - dswrite */
#include <conf.h>
#include <kernel.h>
#include <proc.h>
#include <disk.h>

* dswrite -- write a block (system buffer) onto a disk device.
dswrite(devptr, buff, block)
    struct devsw *devptr;
    char *buff;
    DBADDR block;
{
    struct dreq *drptr;
    char ps;

    disable(ps);
    drptr = (struct dreq *) getbuf(dskrbp);
    drptr->drbuff = buff;
    drptr->drdba = block;
    drptr->drpid = currpid;
    drptr->drop = DWRITE;
    dskenq(drptr, devptr->dvioblk);
    restore(ps);
    return(OK);
}
```

drptr = (struct dreq *) getbuf(dskrbp);

dskenq(drptr, devptr->dvioblk);

var dswritedfg = dfg("dswrite", "getbuf")

Atlas Connection View

Atlas Smart View

dswrite.c

dswrite

Script: Reverse Call

Element Detail View

- dswrite
 - XCSG.ModelElement.id = [Xinu/Xinu/sys/dswrite.c#dswrite(struct devsw*, char*, DBADDR)]
 - XCSG.ModelElement.name = dswrite
 - XCSG.ModelElement.sourceCorrespondence = [322+7+/Xinu/sys/dswrite.c]
 - ##index
 - external
 - isDef
 - XCSG.Function
 - XCSG.Language.C

atlas

Problems Tasks Console Properties Atlas Shell (AtlasToolbox)

show(projection)

show(forwardSlice(projection, selected))

Evaluate: var dswritedfg = dfg("dswrite", "getbuf")

Writable Smart Insert 12 : 1

Project Explorer

- bpdump.c
- chpri.o.c
- ckmode.c
- cksum.s
- clkinit.c
- clkint.c
- clkint.s
- close.c
- conf.c
- control.c
- create.c
- csv.s
- ctxsw.c
- ctxsw.s
- devdump.c
- dfalloc.c
- dfdsrch.c
- dalloc.c
- dgclose.c
- dgcntl.c
- ddgdump.c
- dginit.c
- dgmctrl.c
- dgmopen.c
- dgparse.c
- dgread.c
- dgwrite.c
- dot2ip.c
- dscntl.c
- dsinit.c
- dsinter.c
- dskbcpy.c
- dskenq.c
- dskqopt.c
- dskstrt.c
- dsksync.c
- dsopen.c
- dsread.c
- dsseek.c
- dswrite.c
- dudir.c
- dumkfs.c
- ethinit.c
- ethinter.c
- ethread.c
- ethtrans.c

Graph 20 Graph 21 Graph 22 Graph 23 Graph 25 >47

Atlas Connection View Atlas Smart View

show(dswritedfg)

Script: Reverse Call

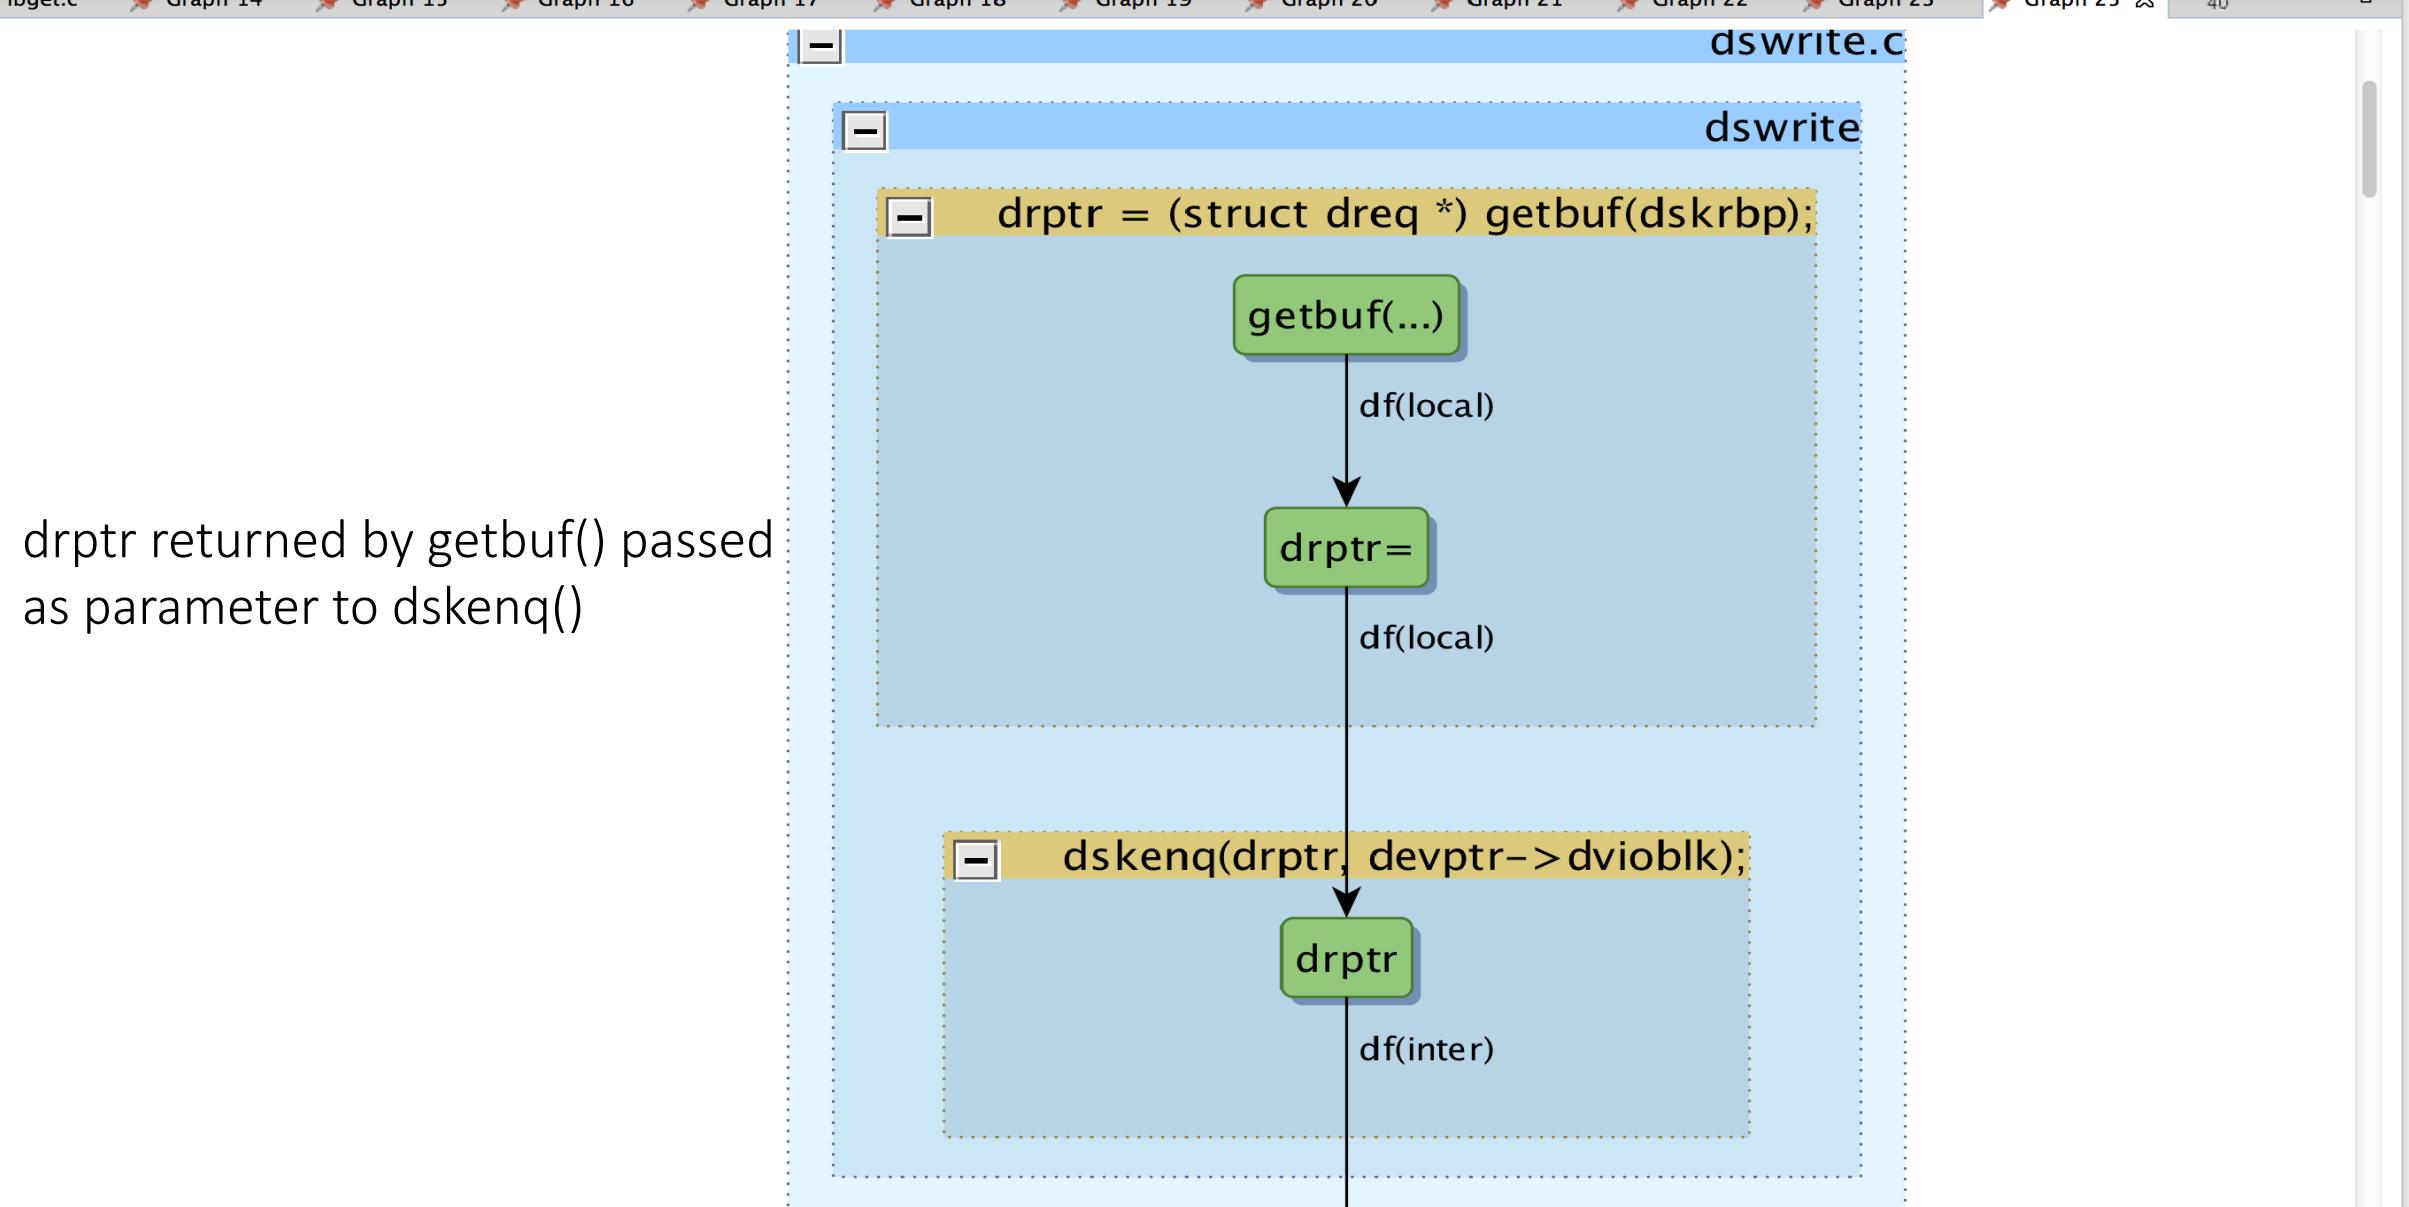
Atlas Shell (Project: AtlasToolbox)

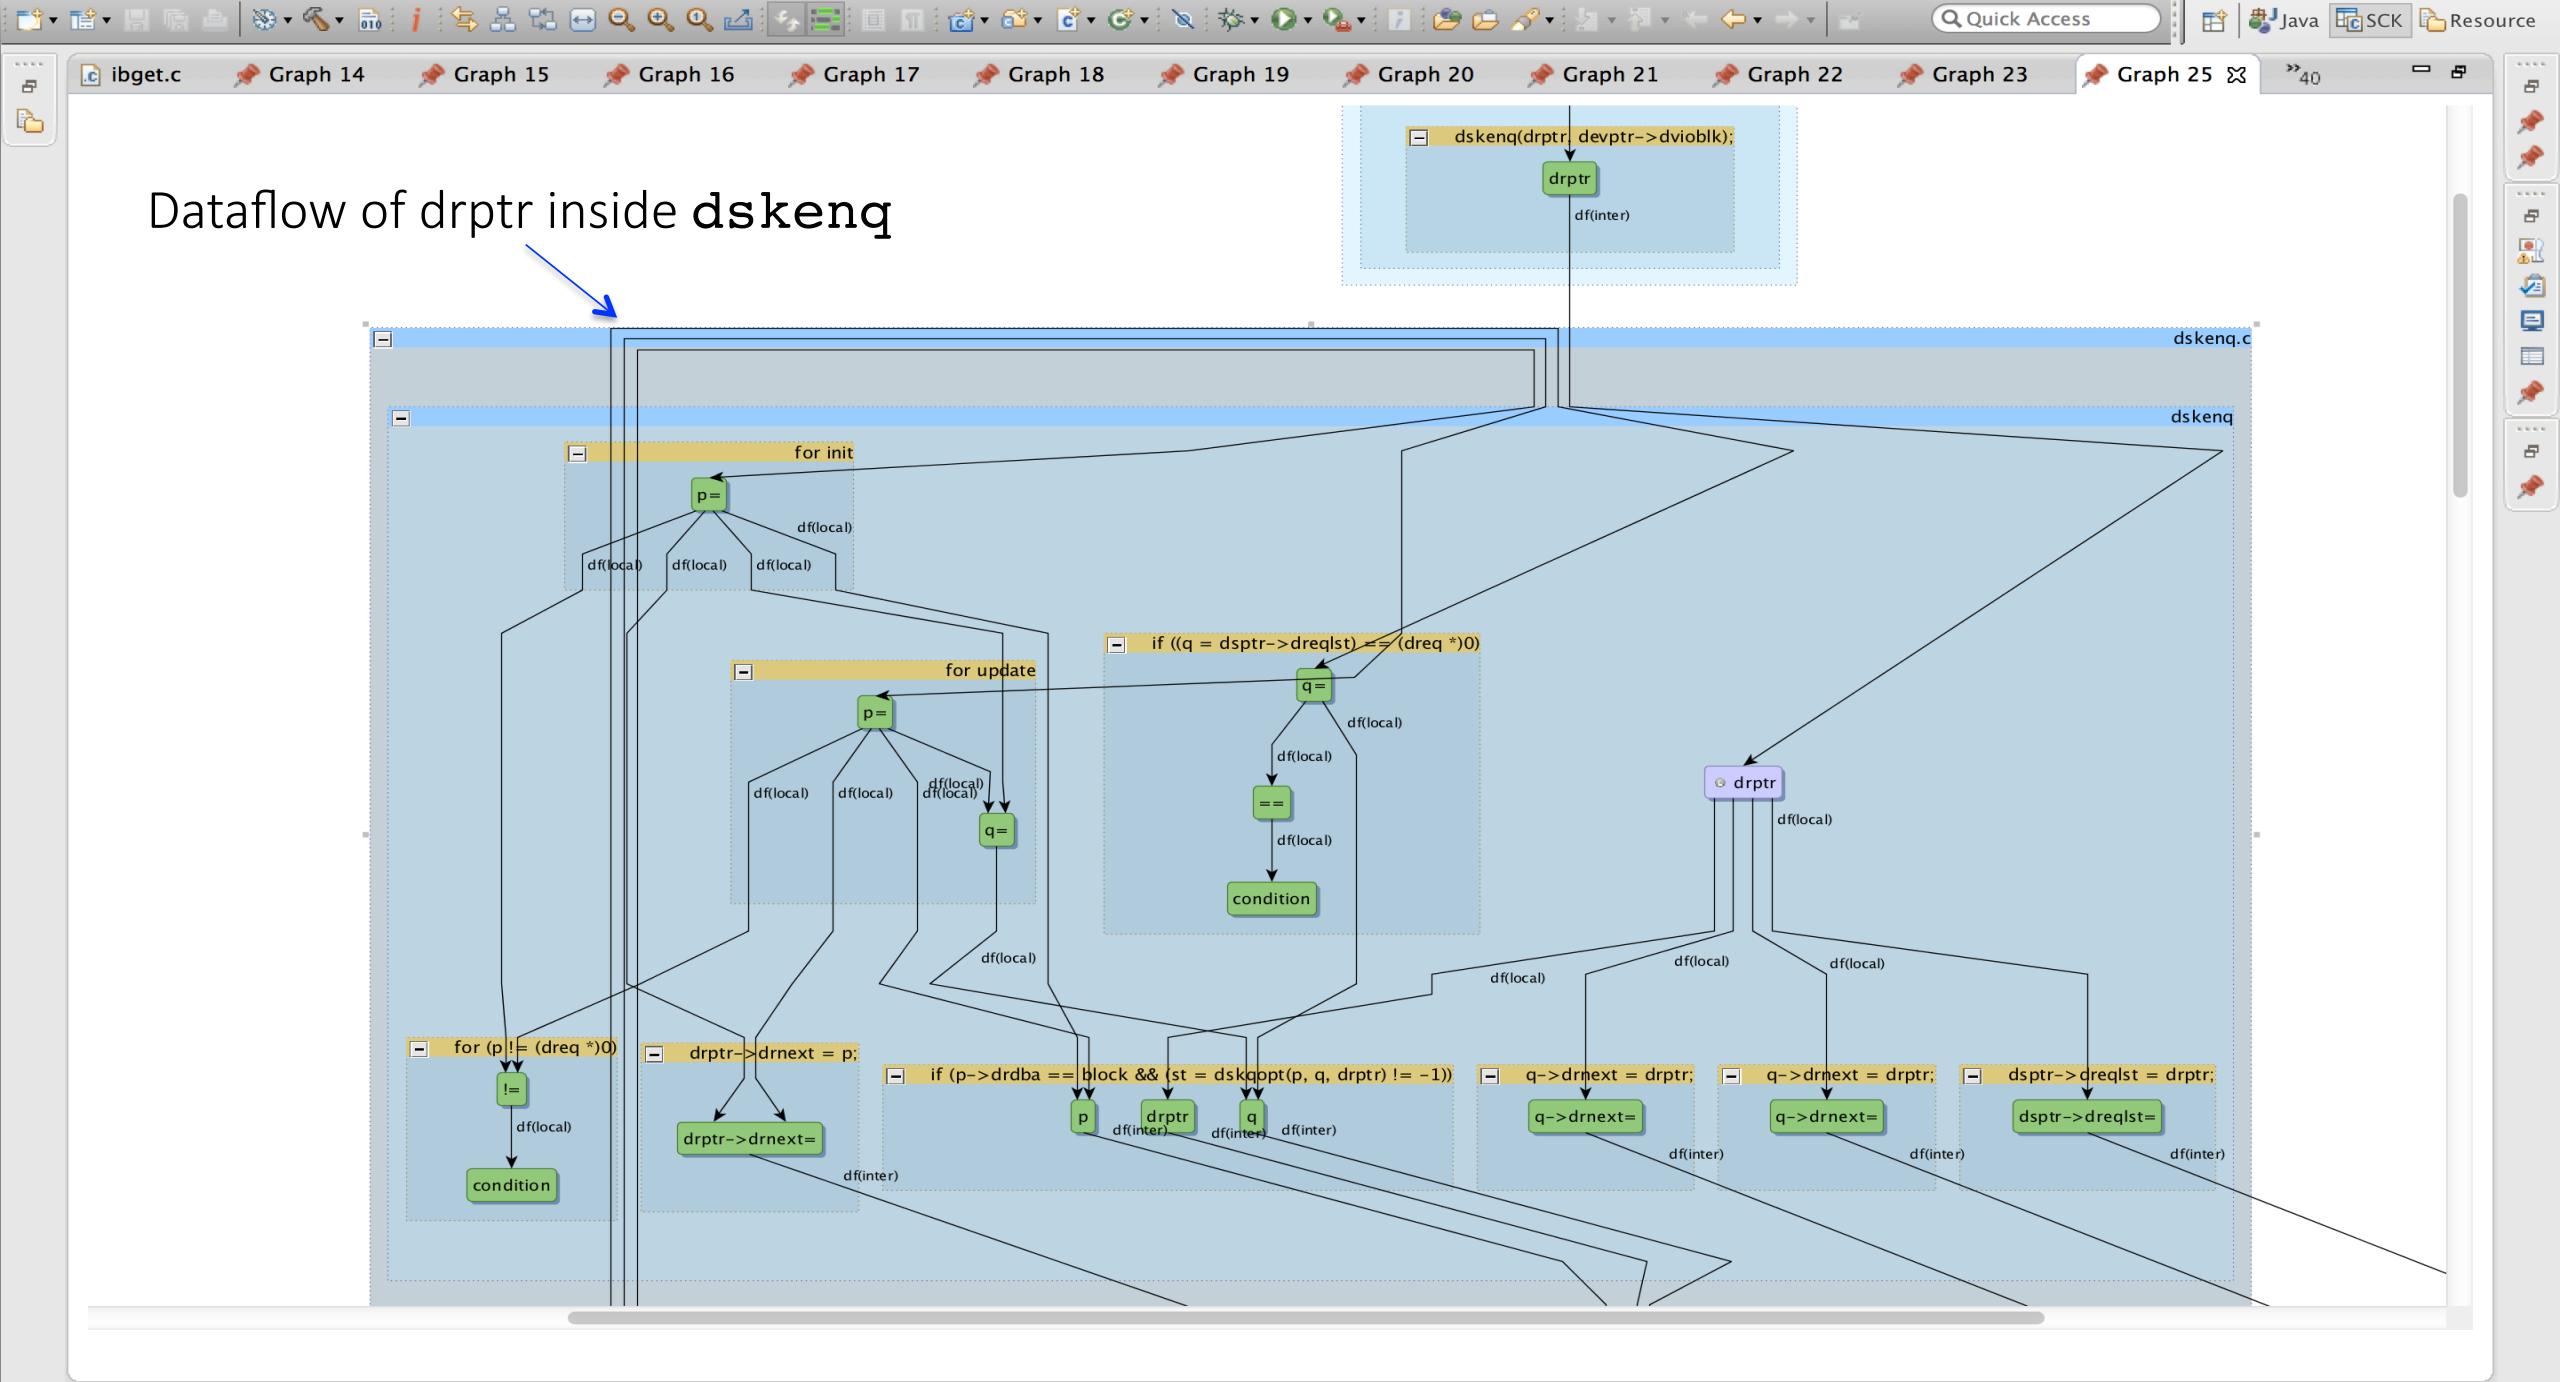
```
var dswritedfg = dfg("dswrite", "getbuf")
dswritedfg: com.ensoftcorp.atlas.core.query.Q = <Atlas query expression>
show(dswritedfg)
```

Evaluate: <type an expression or enter ":help" for more information>

Element Detail View

- dswrite
 - @ XCSG.ModelElement.id = [Xinu/Xinu/sys/dswrite(struct devsw*, char*, DBADDR)]
 - @ XCSG.ModelElement.name = dswrite
 - @ XCSG.ModelElement.sourceCorrespondence = [322+7+/Xinu/sys/dswrite.c]
 - ##index
 - external
 - isDef
 - XCSG.Function
 - XCSG.Language.C





Project Explorer

- bpdump.c
- chprio.c
- ckmode.c
- cksum.s
- clkinit.c
- clkint.c
- clkint.s
- close.c
- conf.c
- control.c
- create.c
- csv.s
- ctxsw.c
- ctxsw.s
- devdump.c
- dfalloc.c
- dfdsrch.c
- dgalloc.c
- dgclose.c
- dgcntl.c
- dgdump.c
- dginit.c
- dgmctrl.c
- dgmopen.c
- dgparse.c
- dgread.c
- dgwrite.c
- dot2ip.c
- dscntl.c
- dsinit.c
- dsinter.c
- dskbcpy.c
- dskenq.c
- dskqopt.c
- dskstrt.c
- dsksync.c
- dsopen.c
- dsread.c
- dsseek.c
- dswrite.c
- dudir.c
- dumkfs.c
- ethinit.c
- ethinter.c
- ethread.c
- reboot.c

Graph 21 Graph 22 Graph 23 Graph 25 Graph 26 >48

Dataflow within `dskenq` – obtained by projecting the DFG

The Dataflow graph for `dskenq` displays the flow of data between various memory locations. It includes nodes for `for init`, `for update`, and `condition` blocks. The graph shows how pointers like `drptr`, `q`, and `p` are updated and used across different operations. Annotations in the graph provide specific code snippets from the source code.

*var projection = projectdfg(dswritedfg, selected)
show(projection)*

Atlas Connection View Atlas Smart View

dswrite.c dsseek.c dsread.c

dskenq

Script: Reverse Call

Atlas Shell (AtlasToolbox)

```
show(dswritedfg)
show(projection)
```

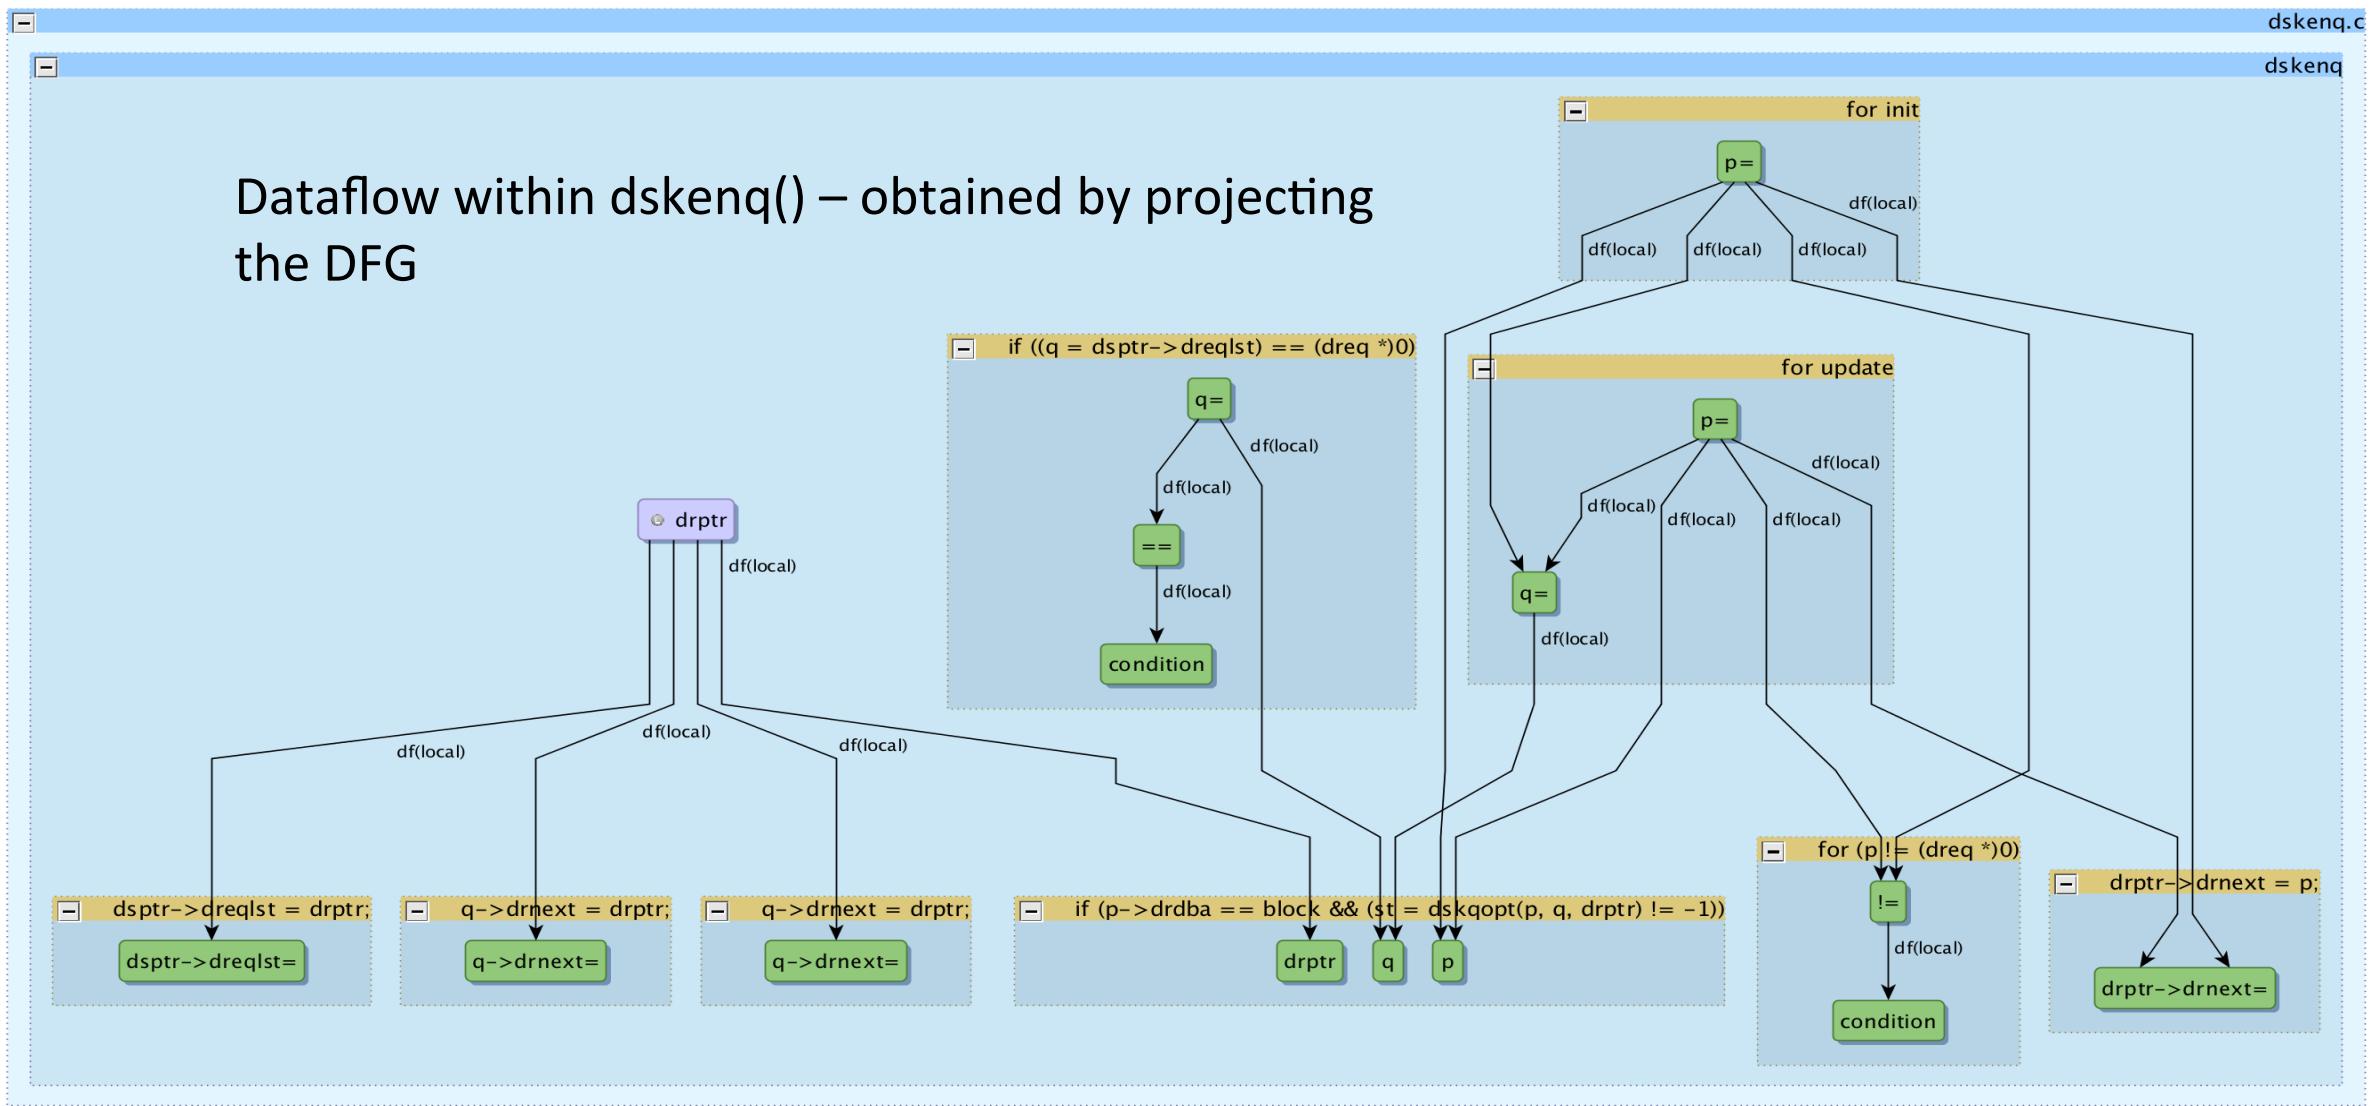
Evaluate: <type an expression or enter ":help" for more information>

Element Detail View

dskenq

- XCSG.ModelElement.id = [Xinu/Xinu/sys/dskenq.c#dskenq(struct dreq*, struct dsblk*)]
- XCSG.ModelElement.name = dskenq
- XCSG.ModelElement.sourceCorrespondence = [306+6+/Xinu/sys/dskenq.c]
- ##index
- external
- isDef
- XCSG.Function
- XCSG.Language.C

Dataflow within dskenq() – obtained by projecting the DFG



Project Explorer

- bpdump.c
- chprio.c
- ckmode.c
- cksum.s
- clkinit.c
- clkint.c
- clkint.s
- close.c
- conf.c
- control.c
- create.c
- CSV.S
- ctxsw.c
- ctxsw.S
- devdump.c
- dfalloc.c
- dfdsrc.h
- dgalloc.c
- dgclose.c
- dgctrl.c
- ddump.c
- dginit.c
- dgmctrl.c
- dgmopen.c
- dgparse.c
- dgread.c
- dgwrite.c
- dot2ip.c
- dscntl.c
- dsinit.c
- dsinter.c
- dskbcp.c
- dskenq.c
- dsqopt.c
- dsstrt.c
- dsksync.c
- dsopen.c
- dsread.c
- dsseek.c
- dswrite.c
- dudir.c
- dumkfs.c
- ethinit.c
- ethinter.c
- ethread.c
- ethstruct.c

Graph 21 Graph 22 Graph 23 Graph 25 Graph 26 »48

show(forwardSlice(projection, selected))

Atlas Connection View Atlas Smart View

Script: Reverse Call

Problems Tasks Console Properties Atlas Shell (AtlasToolbox)

```
show(dswriteln)
show(projection)
```

Evaluate: *show(forwardSlice(projection, selected))*

Element Detail View

drptr

- XCSG.ModelElement.id = [[Xinu/Xinu/sys/dskenq.c#dskenq(struct dreq*, struct dsblk*)], an]
- XCSG.ModelElement.name = drptr
- XCSG.ModelElement.sourceCorrespondence = [341+5+/Xinu/sys/dskenq.c]
- XCSG.Parameter.parameterIndex = 0
- ##index
- isDecl
- XCSG.Language.C
- XCSG.Parameter

Project Explorer

- bpdump.c
- chprio.c
- ckmode.c
- cksum.s
- clkinit.c
- clkint.c
- clkint.s
- close.c
- conf.c
- control.c
- create.c
- csv.s
- ctxsw.c
- ctxsw.s
- devdump.c
- dfalloc.c
- dfdsrch.c
- dgalloc.c
- dgclose.c
- dgcntl.c
- dgdump.c
- dginit.c
- dgmctl.c
- dgmopen.c
- dgparse.c
- dgread.c
- dgwrite.c
- dot2ip.c
- dscntl.c
- dsinit.c
- dsinter.c
- dskbcpy.c
- dskenq.c
- dskqopt.c
- dskstrt.c
- dsksync.c
- dsopen.c
- dsread.c
- dsseek.c
- dswrite.c
- dudir.c
- dumkfs.c
- ethinit.c
- ethinter.c
- ethread.c
- athbrt.c

Graph 22 Graph 23 Graph 25 Graph 26 Graph 27 >49

Forward Slice of a node from dataflow graph

dskenq.c
dskenq

Script: Reverse Call

Problems Tasks Console Properties Atlas Shell (AtlasToolbox)

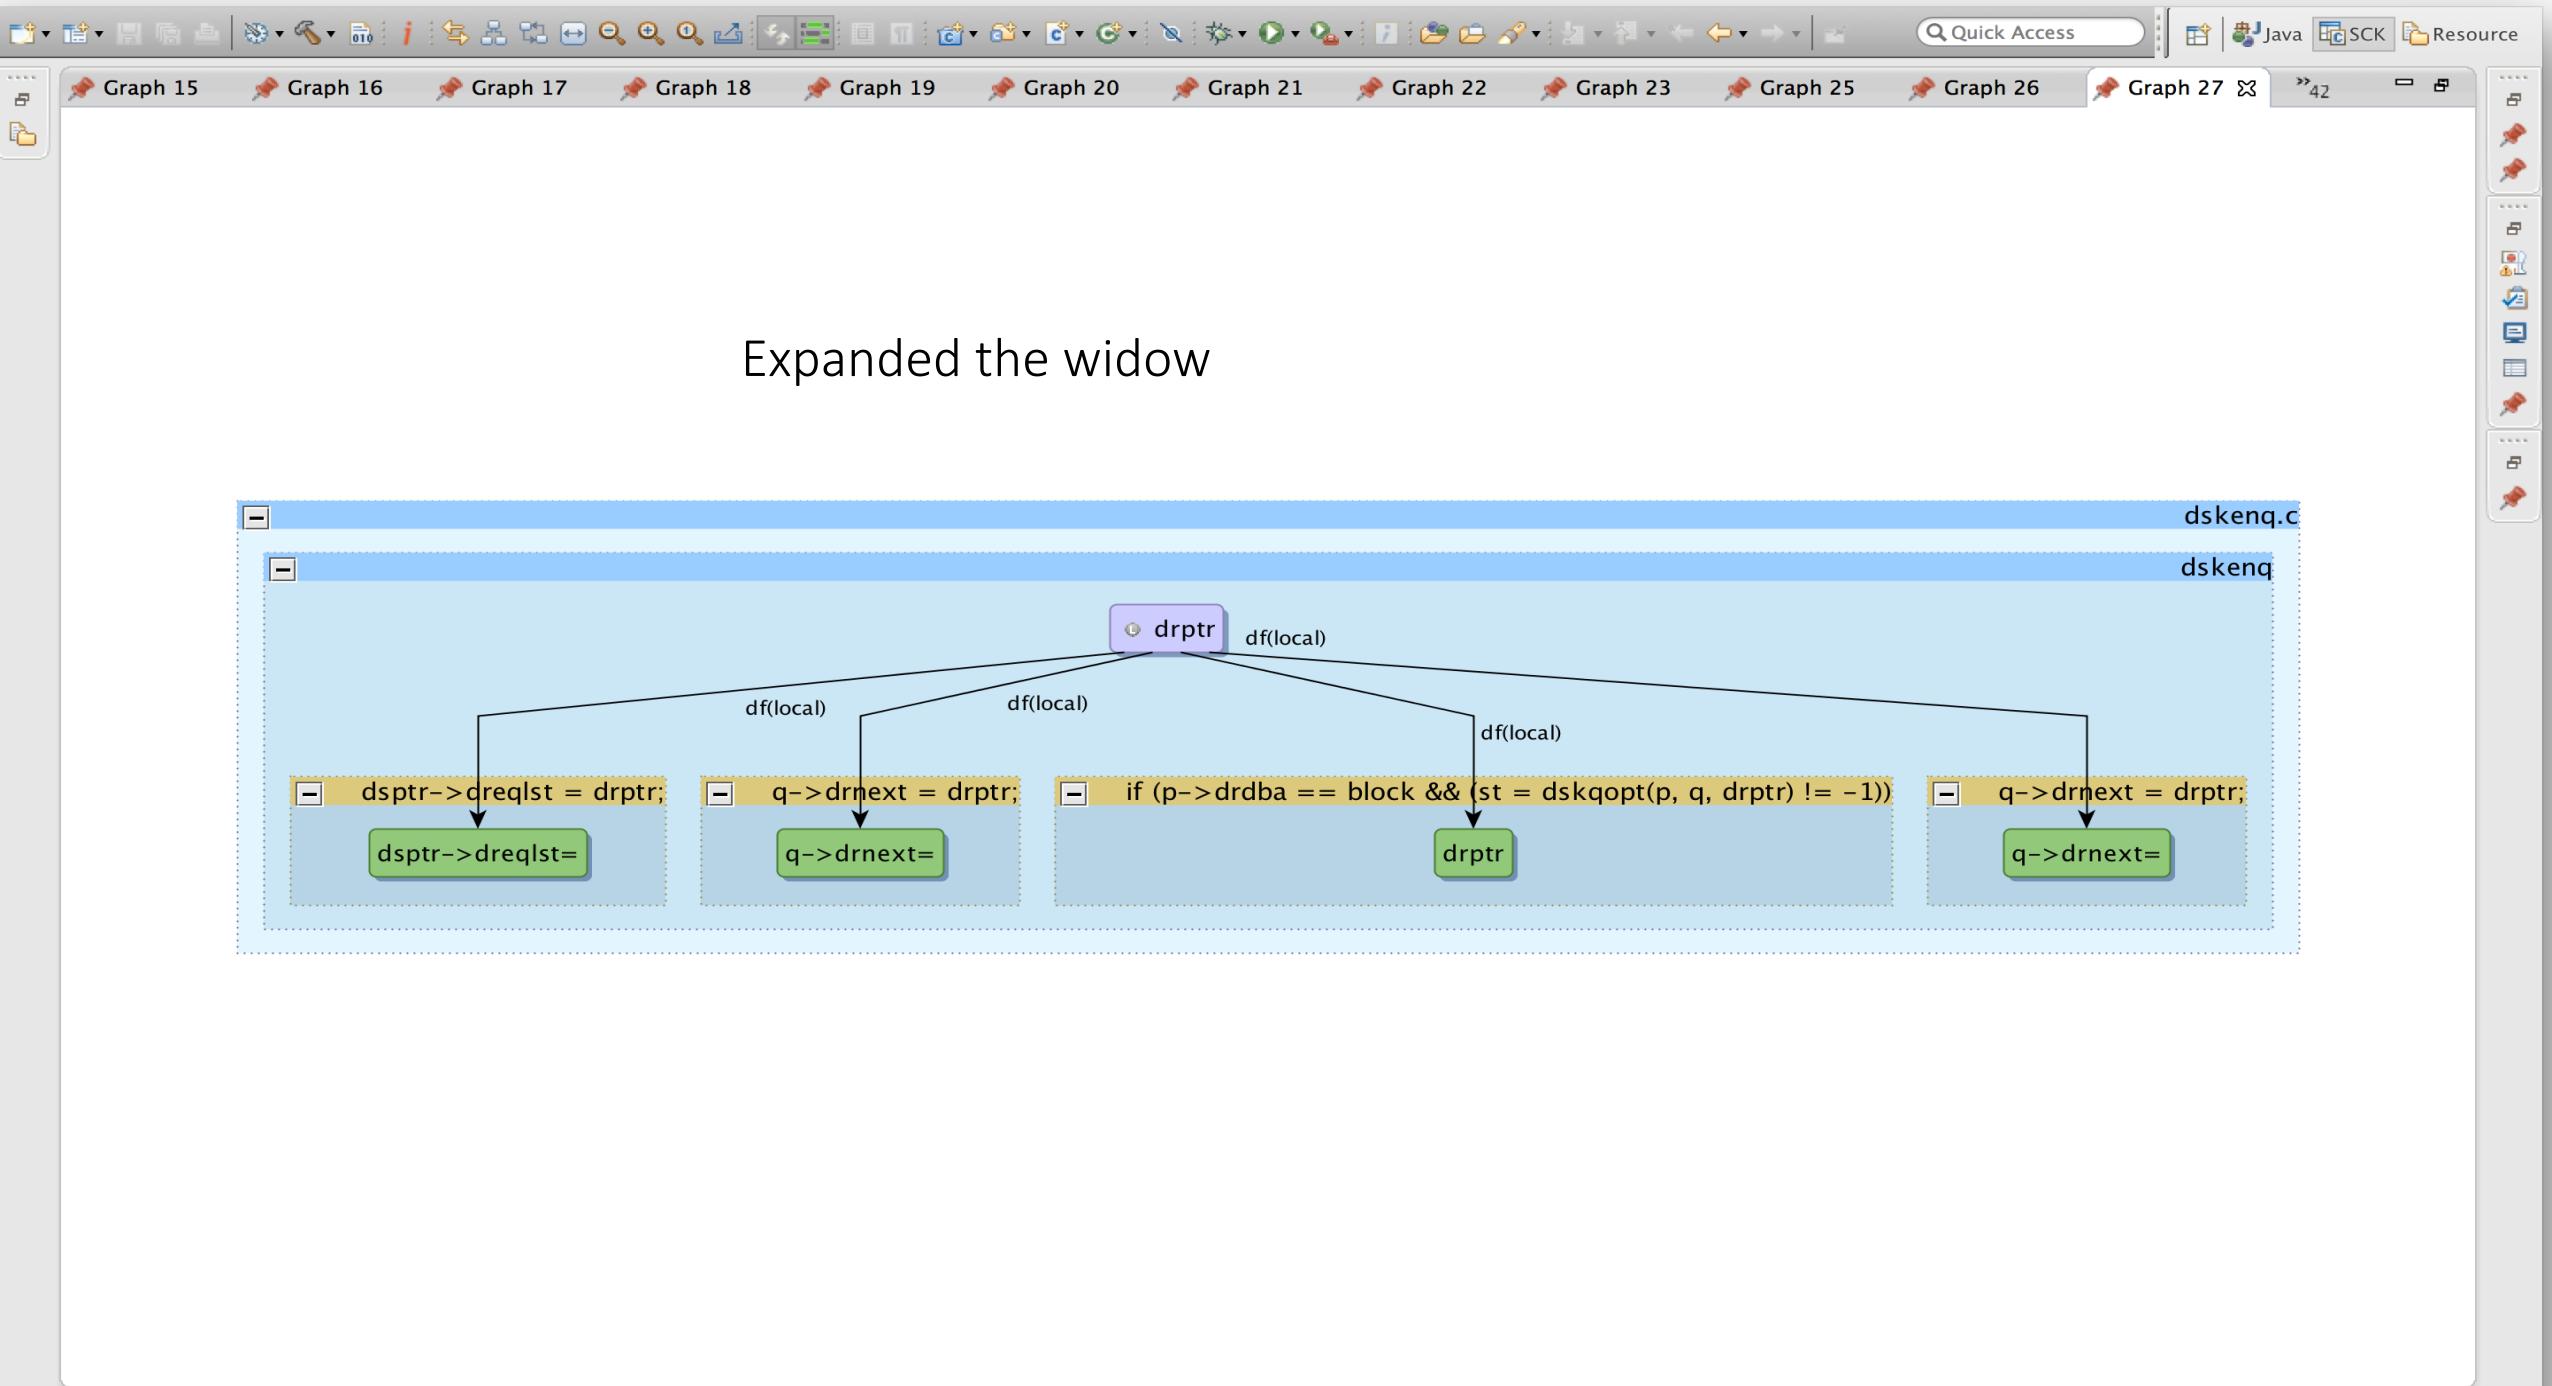
```
show(projection)
show(forwardSlice(projection,selected))
```

Evaluate: <type an expression or enter ":help" for more information> atlas

Element Detail View

drptr

- XCSG.ModelElement.id = [Xinu/Xinu/sys/dskenq.c#dskenq(struct dreq*, struct dsblk*)], an
- XCSG.ModelElement.name = drptr
- XCSG.ModelElement.sourceCorrespondence = [341+5+/Xinu/sys/dskenq.c]
- XCSG.Parameter.parameterIndex = 0
- ##index
- isDecl
- XCSG.Language.C
- XCSG.Parameter



Two Aspects of Hardness

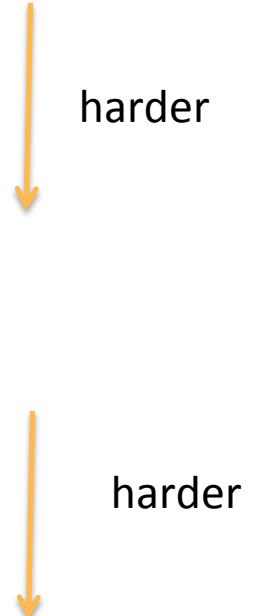
- Interprocedural Interactions get complicated because of the varied ways the functions communicate with each other.
- Within a function, it gets complicated because of many paths, and with different complications within paths.

Quest for Relevant Functions

- **Relevant:** functions necessary and sufficient to verify a problem instance (e.g. each memory allocation).
- **Butterfly Effect:** A small change at one point causes unforeseen effects at far away points.
- How do we find relevant functions?
 - Identify all the communication mechanisms (e.g. pointer to allocated memory gets passed as a parameter).
 - Formulate the constraints for identifying relevant functions (e.g. functions called by getbuf and freebuf should not be in the set of relevant functions).

Fundamentals of data and control flow

- Relevant data: Data (D) relevant to a problem instance (e.g. pointers to an allocated memory).
- Fundamental mechanisms of data flow:
 - f passes D as a parameter to a callee function g.
 - f passes D as a parameter or return to a caller function g.
 - f and g share D through a global variable.
- Fundamental mechanisms of control flow:
 - f calls g directly.
 - f calls g indirectly (e.g. using a function pointer).
 - f and g operate asynchronously (control transfer happens



Formulation of constraints

- **Data-Centric Constraints:** Exclude functions that do not access the data (D) relevant to the problem instance (e.g. exclude functions that do not access pointers to the allocated memory).
- **Control-Centric Constraints:** Exclude functions that cannot reach to functions relevant to the problem instance (e.g. exclude functions that do not have a call chain reaching to freebuf).
- **Design-Centric Constraints:** Exclude functions that imply an unnecessarily complicated design.

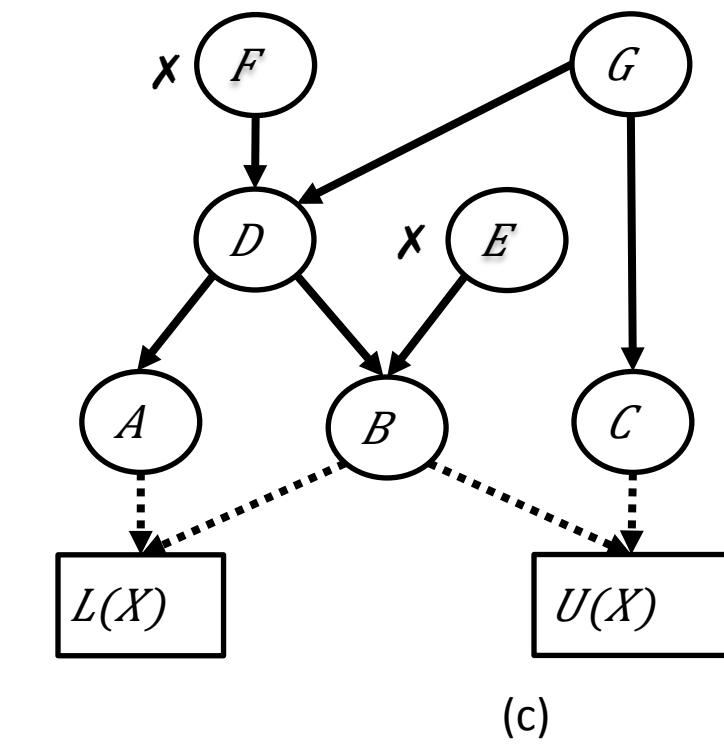
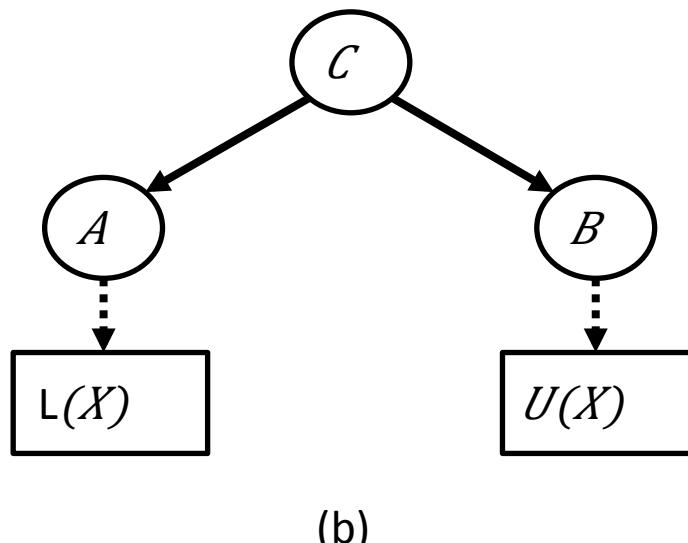
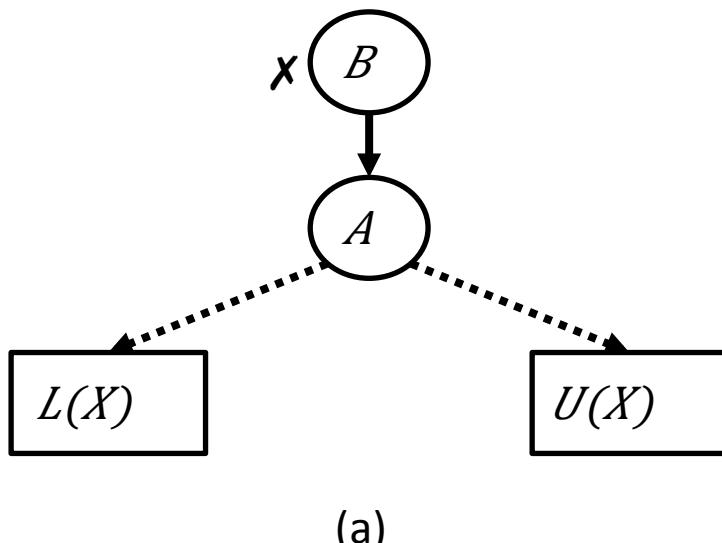
An example of design-centric constraint

A constraint for the MP problems – If a function f reaches getbuf and freebuf, the developer can implement the matching logic at f itself. The design is unnecessarily complicated if that logic is moved up to a function g that calls f .

→ calls

.....→ contains

What functions should not be there (X)?



(a)

(b)

(c)

Matching Pair Graph

- A visual graph model to discover functions relevant to each MP instance.
- Founded on the knowledge of the fundamental *propagation mechanisms* and the *propagation boundaries*.
- *Directed graph* with functions as the nodes and calls as the edges.
- *Roots* are asynchronous functions.
- *Leaves* are MP functions (e.g. `getbuf` and `freebuf`).

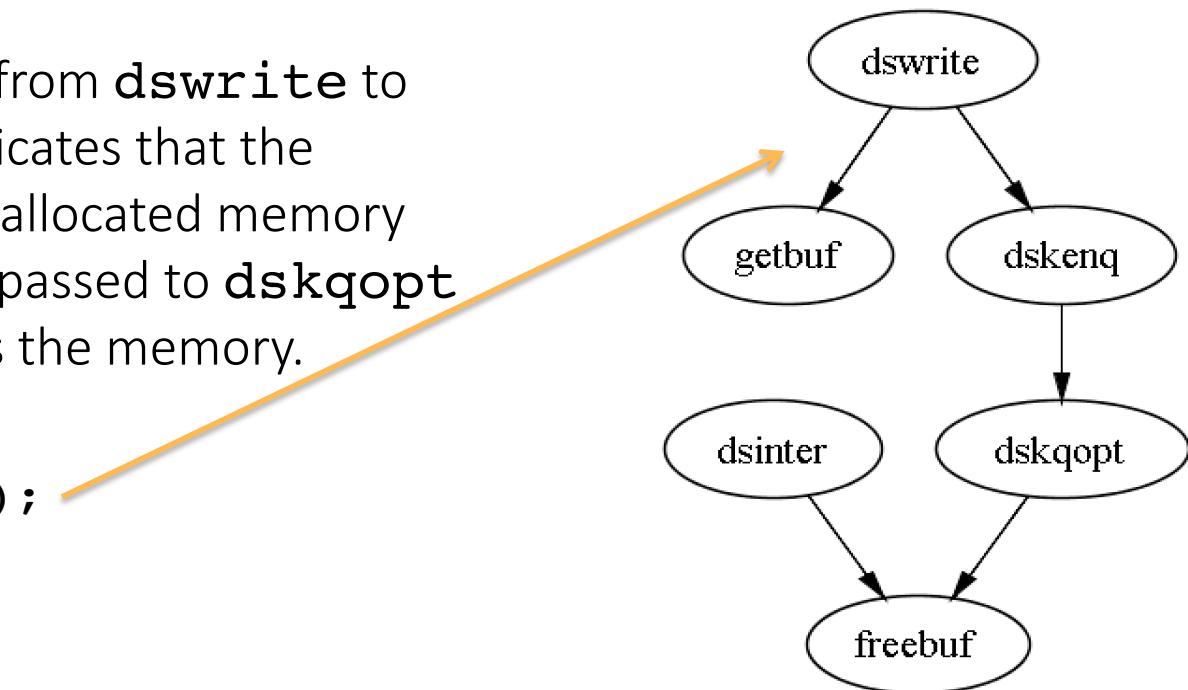
MPG can serve as verification-critical evidence

We illustrate how MPG provides important clues to verify an MP instance.

```
dswrite(devptr, buff, block)
    struct devsw *devptr;
    char *buff;
    DBADDR block;
{
    struct dreq *drptr;
    char ps;

    disable(ps);
    drptr = (struct dreq *) getbuf(dsrbkp);
    drptr->drbuff = buff;
    drptr->drdba = block;
    drptr->drpid = currpid;
    drptr->drop = DWRITE;
    dskenq(drptr, devptr->dvioblk);
    restore(ps);
    return(OK);
}
```

The call chain from **dswrite** to **dskqopt** indicates that the pointer to the allocated memory (**drptr**) gets passed to **dskqopt** which releases the memory.



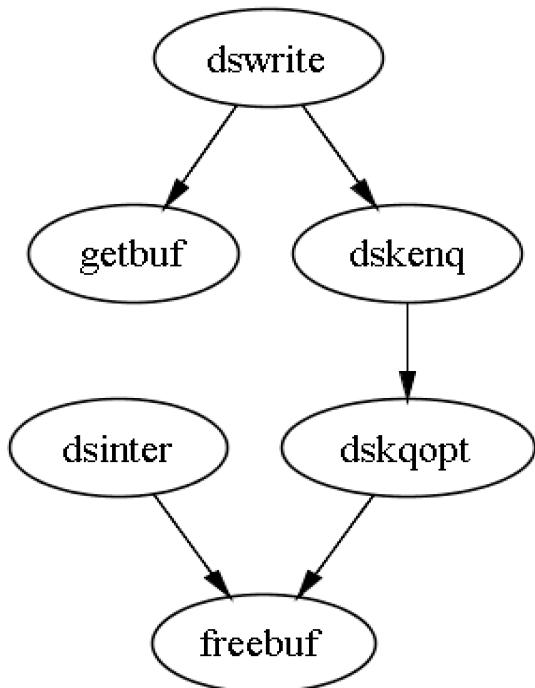
Interesting Question: What does the presence of **dsinter** indicate?

MPG as Verification - Critical Evidence

Question: What does the presence of **dsinter** indicate?

Note that:

1. The function **dsinter** deallocates memory.
2. **dsinter** does not allocate memory.
3. **dswrite** and **dsinter** operate asynchronously.



*On some paths the pointer to the allocated memory (**drptr**) is shared with **dsinter** and it deallocates the memory asynchronously.*

*Enables correct reasoning for a case involving **asynchronous deallocation** by an interrupt routine - a hard case for automation.*

A closer look dsinter – the interrupt-driven routine

```
INTPROC    dsinter(dsptr)
           struct      dsblk *dsprt;
{
    struct      dtc      *dtptr;
    struct      dreq      *drptr;

    dtptr = dsprt->dcsr;
    drptr = dsprt->drqlst;
    if (drptr == DRNULL) {
        panic("Disk interrupt when disk not busy");
        return;
    }
    X if (dtptr->dt_csr & DTERROR)
        drptr->drstat = SYSERR;
    else
        X drptr->drstat = OK;
    X if ( (dsprt->drqlst=drptr->drnnext) != DRNULL)
        dskstrt(dsprt);
    switch (drptr->drop) {

        X case DREAD:
        X case DSYNC:
            ready(drptr->drpid, RESCHYES);
            return;

        case DWRITE:
        X freebuf(drptr->drbuff);
        /* fall through */
        case DSEEK:
        X freebuf(drptr);
    }
}
```

How many relevant execution paths ?

Principle: A branch relevant only if it has interesting events

Three interesting events marked

Several branches eliminated from consideration, because no interesting events on those branches

Left with only 2 paths out of 17 potential paths – less work !

```

INTPROC      dsinter(dsprt)
    struct      dsblk *dsprt;
{
    struct      dtc   *dtptr;
    struct      dreq   *drptr;

    dtptr = dsprt->dcsr;
    drptr = dsprt->dreqlst;
    if (drptr == DRNULL) {
        panic("Disk interrupt when disk not busy");
        return;
    }
    if (dtptr->dt_csr & DTERROR)
        drptr->drstat = SYSERR;
    else
        drptr->drstat = OK;
    if ( (dsprt->dreqlst=drptr->drnext) != DRNULL)
        → dskstrt(dsprt);
    switch (drptr->drop) {

        case DREAD:
        case DSYNC:
            ready(drptr->drpid, RESCHYES);
            return;

        case DWRITE:
            freebuf(drptr->drbuff);
            /* fall through */
        case DSEEK:
            freebuf(drptr);
    }
}

```

An escape happened but
through parameter to dskstrt

Principle: Function calls not
interesting if the function is not in
MPG(X)

dskopt Seven execution paths shown as rows

| Input | Conditions: <code>dreq.drop ==</code> | | | | | | Event Description |
|--------------------------------------------|---------------------------------------|-------|----|-------------|------------|-------------|-------------------------------------------------------------------------------------------------------------------|
| Dreq <code>devsw.dvioblk.dreqlst</code> | DYNCH | DSEEK | C1 | DWRITE & C2 | DREAD & C3 | DWRITE & C3 | |
| | T | - | - | - | - | - | Error |
| | F | T | - | - | - | - | <code>freebuf, sig = dreq</code> |
| | F | F | T | - | - | - | <code>devsw.dvioblk.dreqlst ← dreq</code> <code>freebuf, sig = dreq</code> |
| | F | F | F | T | - | - | <code>devsw.dvioblk.dreqlst ← dreq</code> <code>freebuf, sig = char</code> <code>freebuf, sig = dreq</code> |
| | F | F | F | F | T | - | |
| | F | F | F | F | F | T | <code>devsw.dvioblk.dreqlst ← dreq</code> |
| | F | F | F | F | F | F | <code>devsw.dvioblk.dreqlst ← dreq</code> |

dsinter

| Input | Conditions | | | | | | | Event Description |
|-------|-----------------|----|----|----------------|-------|--------|-------|------------------------------------------------------------------------------------------------------------|
| dsblk | drptr == DRNULL | D1 | D2 | drptr->drop == | | | | |
| | | | | DREAD | DSYNC | DWRITE | DSEEK | |
| | F | - | - | F | F | T | | <code>dreq ← dsblk.dreqlst</code> <code>freebuf, sig = char*</code> <code>freebuf, sig = dreq</code> |
| | F | - | - | F | F | F | T | <code>dreq ← dsblk.dreqlst</code> <code>freebuf, sig = dreq</code> |

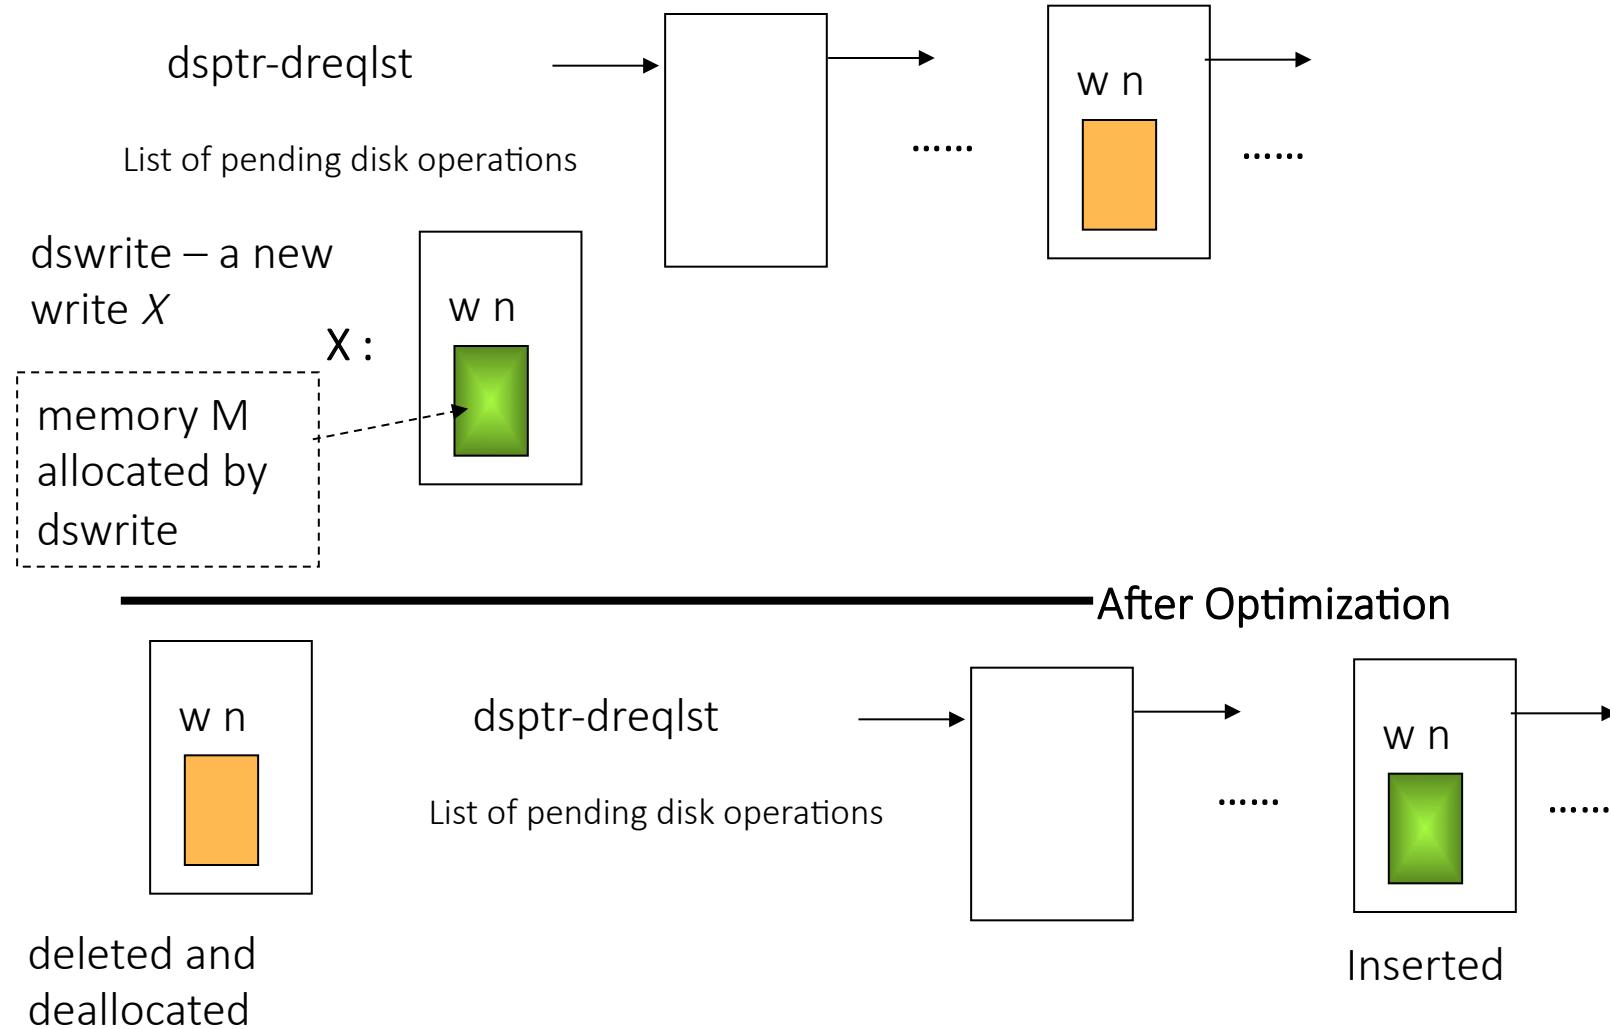
A logical explanation of the disk driver

- This part of the XINU code is very well written – it is complicated code because the underlying logic requires it to be.
- We will go through a relevant part of the underlying logic to appreciate the code.
- Note that the code had to be written very carefully – several lines of code are critical, miss one of them and it will create a memory leak.
- For example, miss the line `q->drnext = drptr;` in path 2-F and it will be a serious memory leak that will crash the system very quickly.

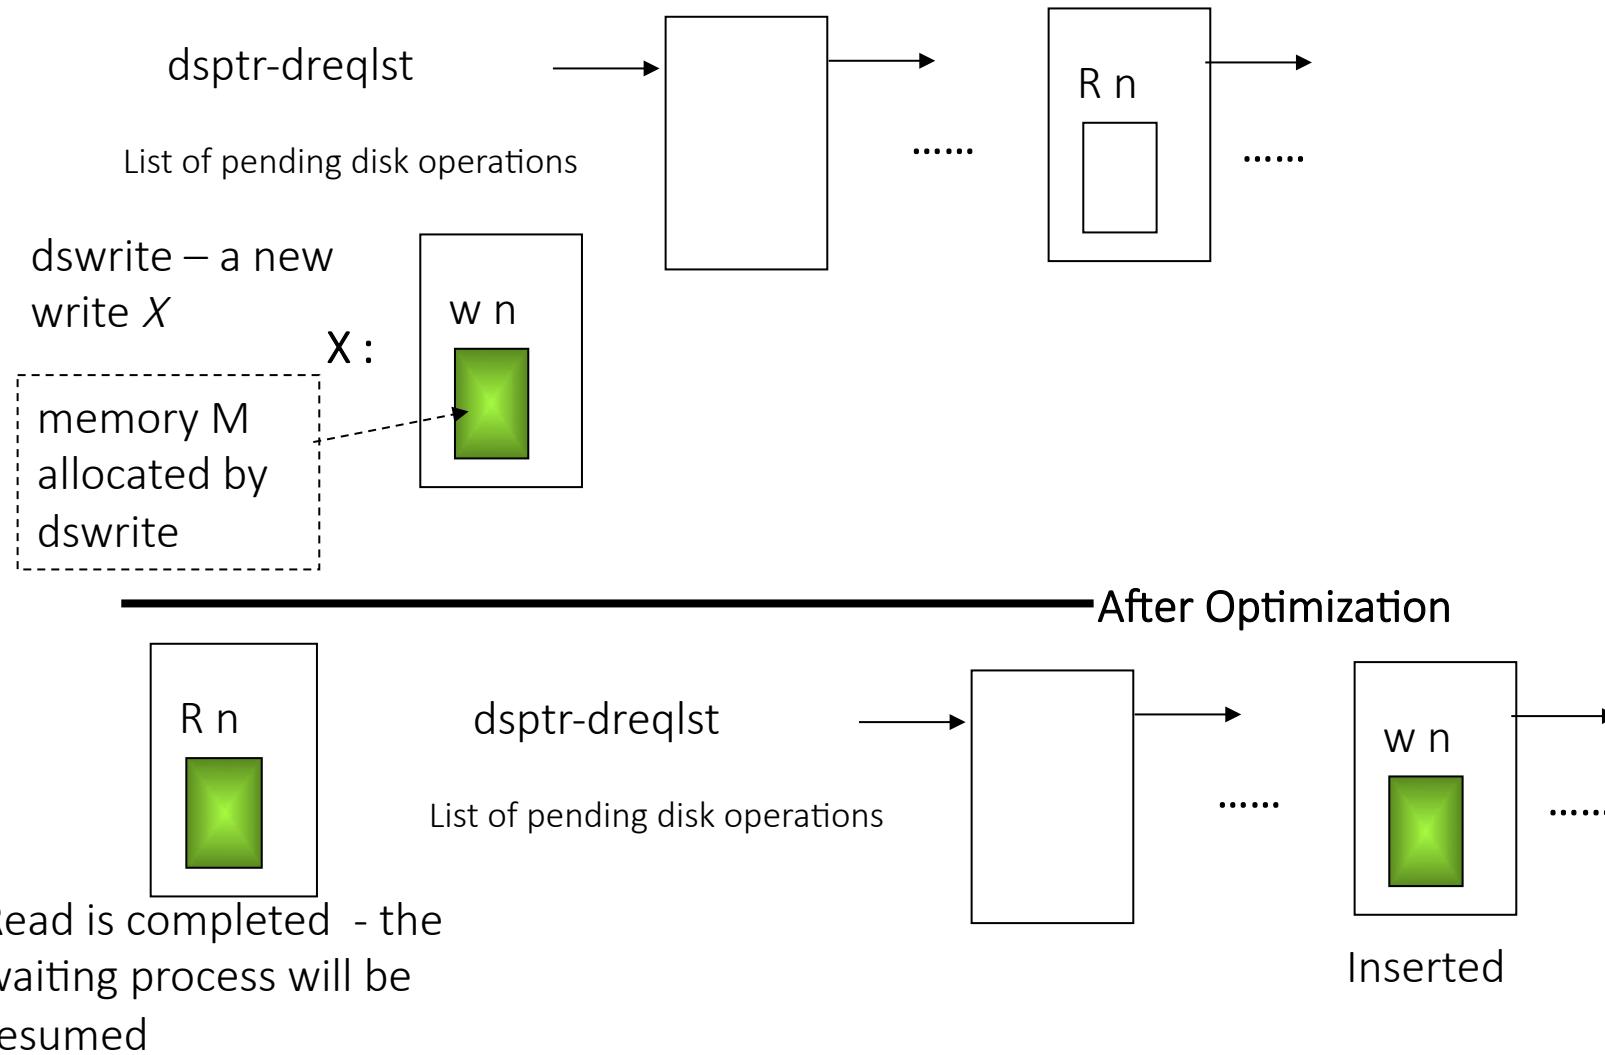
The underlying optimization

- A *write X* (or another disk operation) is issued for disk block n
- Disk operations, since they are slow, get queued in a linked list L of pending operations to be performed.
- Optimize if there is already a pending operation for the same block:
 - Path D (*/* dup write */*): A previous pending *write* to block n is deleted from L and the X is inserted in L .
 - Path F (*/* sat. old read */*): A previous pending *read* to block n is completed (and deleted from L) in memory using X the X is inserted in L .

Duplicate write optimization



Duplicate write



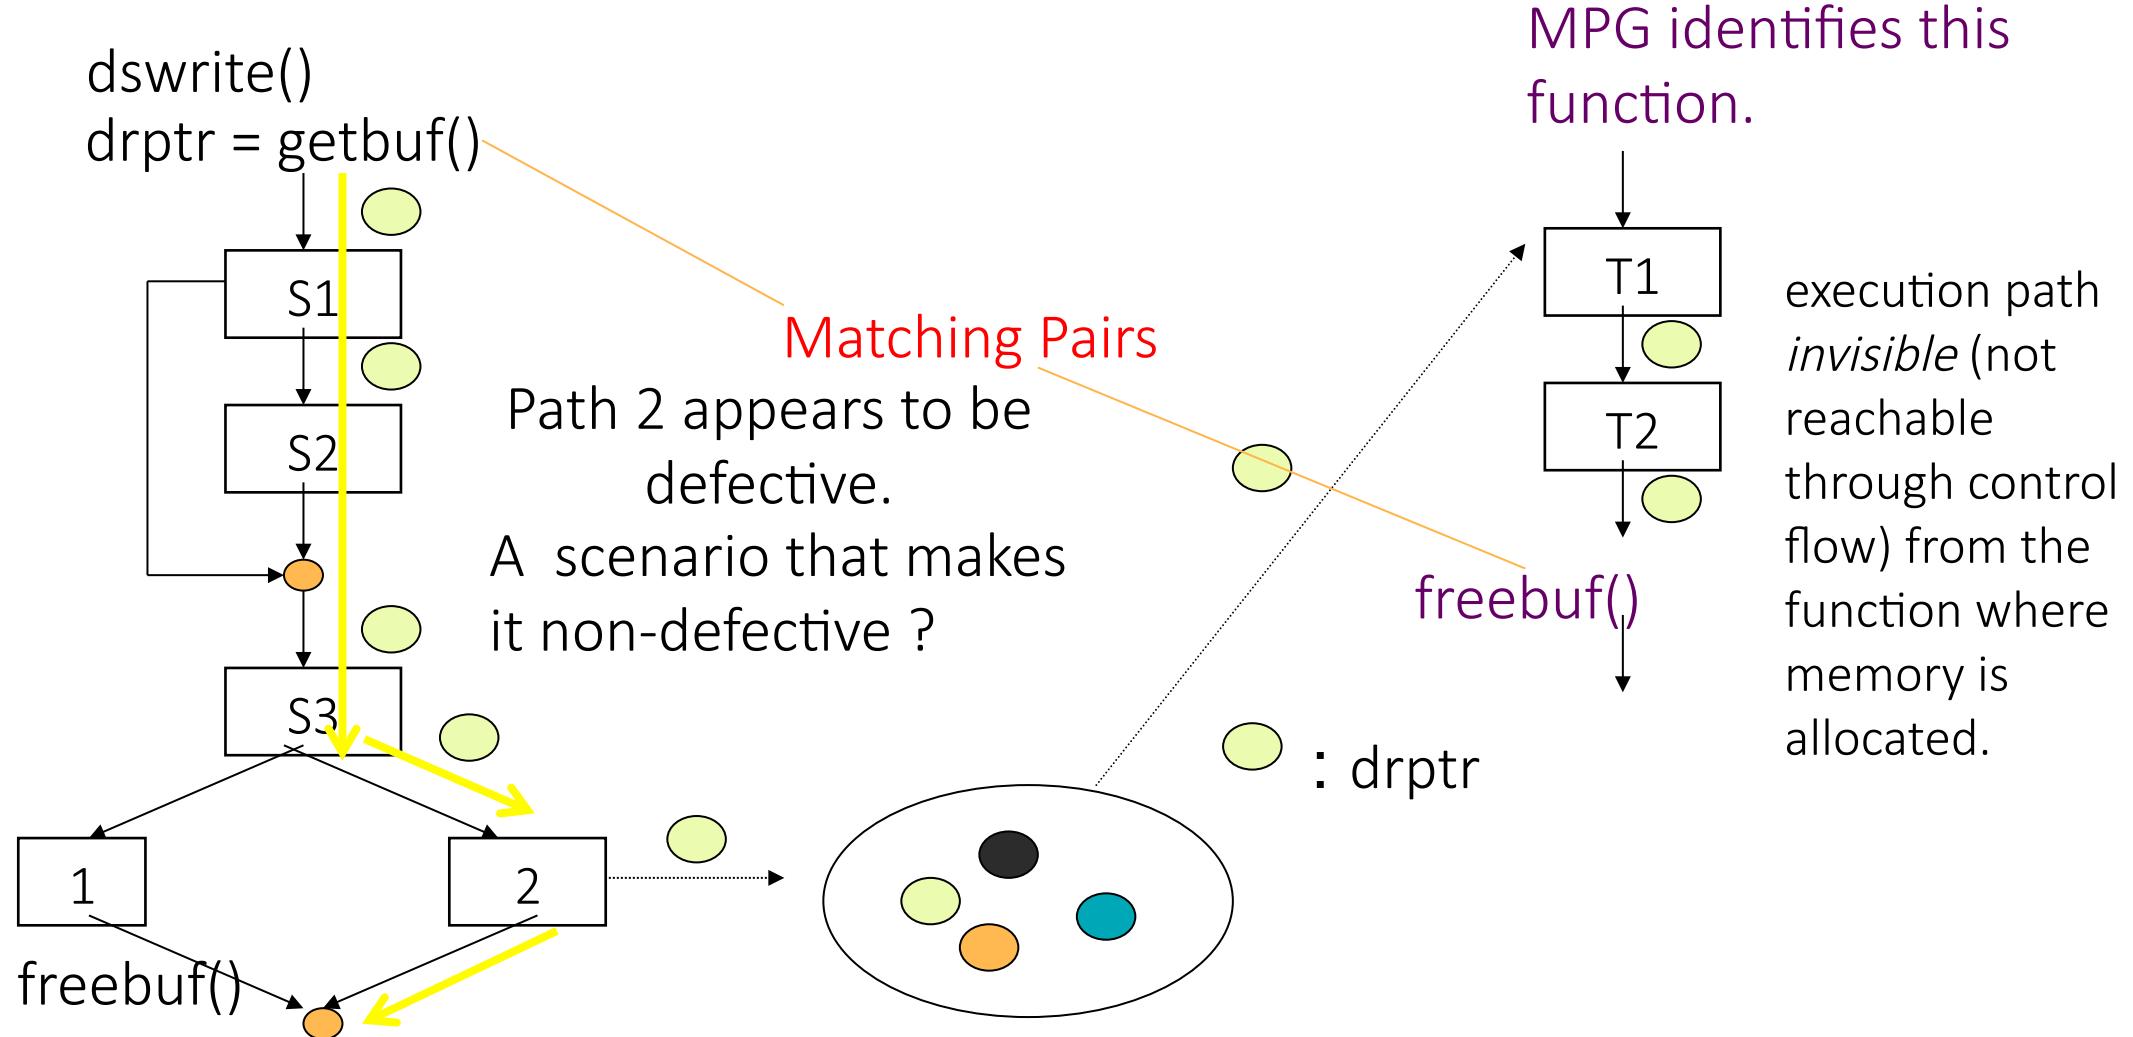
About the paths in dskeng

- Path 1: List L of pending disk operations is empty
- Path 2: L already has a pending operation for the same block
- Path 3: Insertion in the *middle* of the list – thus L is kept sorted to minimize time spent on moving the disk head
- Path 4: Insertion at the *end* of the list

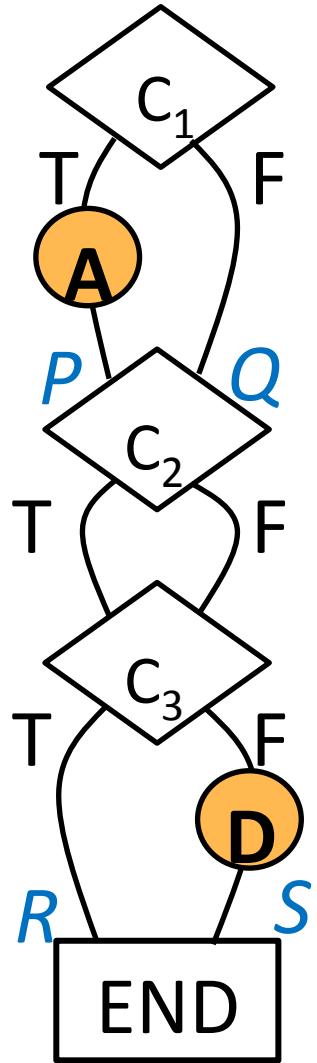
The producer-consumer pattern

- The device driver - follows a widely used coding pattern, called the *producer-consumer problem*.
 - Producer P – a thread that generates requests, *inserts* them in a list L . P *allocates* memory for the request.
 - Consumer C – another thread that consumes requests, deletes them from the list L . P *deallocates* memory for the request.
- A classic example of the *level-4 MP* problem
- Also, a *synchronization problem* because of the need for mutually exclusive access to the shared list L .
 - trivially solved in XINU by disabling interrupts

Hardness due to asynchronous paths



Irrelevant branch nodes

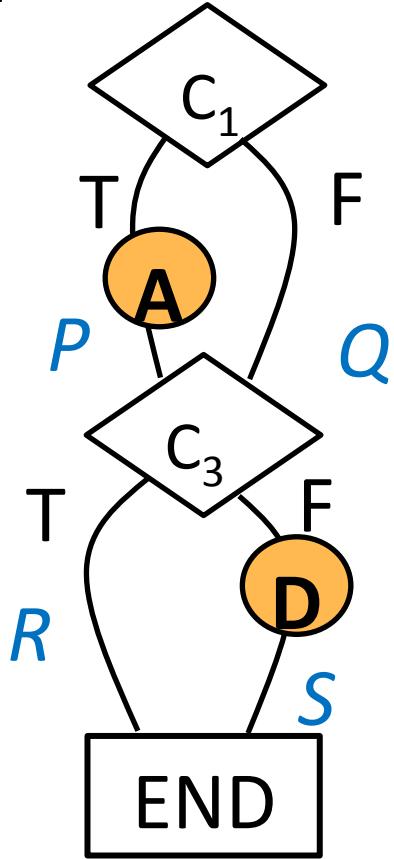


C_2 Irrelevant to
path-sensitive analysis



Remove the Irrelevant branch nodes to
avoid unnecessary path explosion & also
simplify the path feasibility check.

paths reduced from 8 to 4

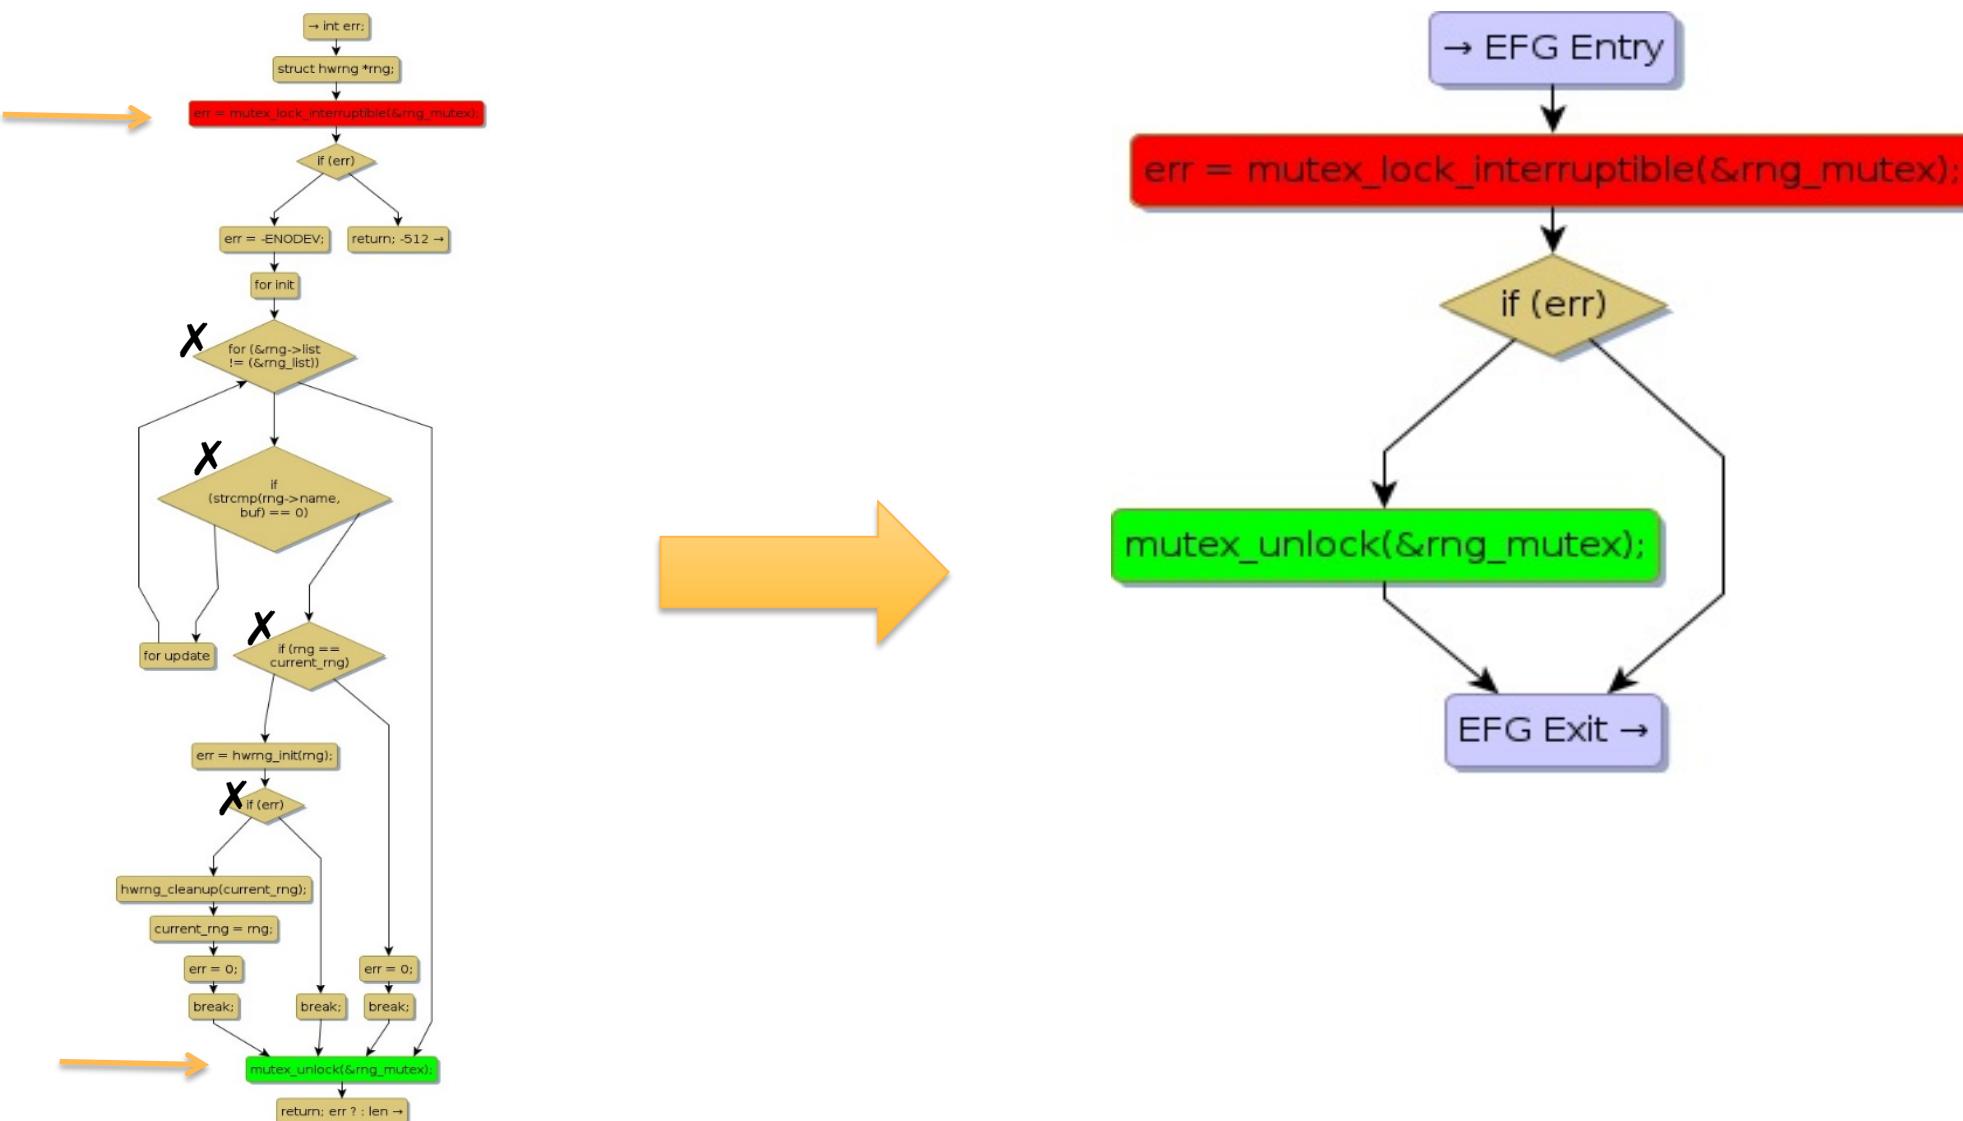


conditions for feasibility
check reduced from 3 to 2

Efficient path-sensitive analysis

- A large number of paths could be partitioned into a small number of groups.
- All Paths in a group are equivalent – have the same execution behavior w.r.t. the property to be verified.
- Efficient computation by examining only one path from each group.
- **Challenge:** How can the groups be formed without examining each path at least once?

A Linux example to illustrate efficient path-sensitive analysis



PCG for efficient path-sensitive analysis

- Projected Control Graph (PCGG) is derived from the Control Flow Graph (CFG).
- Relevant events: CFG nodes relevant to the property (e.g. `getbuf`, `freebuf`, nodes of type `p = drptr` (pointer to the allocated memory)).
- Behavior: An event trace along a CFG path
- PCG has one path per group of CFG paths with identical behaviors.
- PCG construction: $O(K)$ where $K = N + E$, N : # nodes in CFG, E : # edges in CFG.

Hardness captured by MPG and PCG

- MPG captures interprocedural complexity:
 - the functions to be analyzed.
 - the data flow between the relevant functions.
 - the control flow between the relevant functions.
- EFG captures intra-procedural Complexity:
 - the paths with distinct behaviors.
 - the relevant conditions for performing path feasibility check.
- MPG & EFG together serve as:
 - The evidence to verify each MP instance.
 - They reveal the hardness challenges of each MP instance and thus can be the basis for defining hardness metrics.

Using Atlas to create matching pair graph

- MPG:
 - Functions and their call relationships *relevant* for verifying an MP instance.
 - *Relevant*: functions that contain the matching events and also the functions propagating the pointers to objects for matching the events.
 - Based on fundamental principles of how functions communicate.
- The lab covers simple experiments to evolve ideas

Project Explorer

```

dswrite.c Graph 2 Graph 3 Graph 4 Graph 5 Graph 6 >56
/* dswrite.c - dswrite */

#include <conf.h>
#include <kernel.h>
#include <proc.h>
#include <disk.h>

* dswrite -- write a block (system buffer) onto a disk device.
dswrite(devptr, buff, block)
    struct devsw *devptr;
    char *buff;
    DBADDR block;
{
    struct dreq *drptr;
    char ps;

    disable(ps);
    drptr = (struct dreq *) getbuf(dskrbp);
    drptr->drbuff = buff;
    drptr->drdba = block;
    drptr->drpid = currpid;
    drptr->drop = DWRITE;
    dskenq(drptr, devptr->dvioblk);
    restore(ps);
    return(OK);
}

```

Atlas Connection View

Forward call graph is neither sufficient nor necessary.

Let us try some ideas for the MPG (the graph of relevant functions with call relationships for verifying the allocation instance in **dswrite**).

Atlas Shell (Project: AtlasToolbox)

```

var z = graph(rootx,y)
z: com.ensoftcorp.atlas.core.query.Q = <Atlas query expression>
show(z)

Evaluate: <type an expression or enter ":help" for more information>

```

Element Detail View

- dswrite**
 - XCSG.ModelElement.id = [Xinu/Xinu/sys/dswrite(struct devsw*, char*, DBADDR)]
 - XCSG.ModelElement.name = dswrite
 - XCSG.ModelElement.sourceCorrespondence = [322+7+/Xinu/sys/dswrite.c]
 - ##index
 - external
 - isDef
 - XCSG.Function
 - XCSG.Language.C

Project Explorer

- csv.s
- ctxsw.c
- ctxsw.s
- devdump.c
- dfalloc.c
- dfdsrch.c
- dgalloc.c
- dgclose.c
- dgcntl.c
- dgdump.c
- dginit.c
- dgmctl.c
- dgmopen.c
- dgparse.c
- dgread.c
- dgwrite.c
- dot2ip.c
- dscntl.c
- dsinit.c
- dsinter.c
- dskbcpy.c
- dskenq.c
- dskqopt.c
- dskstrt.c
- dsksync.c
- dsopen.c
- dsread.c
- dsseek.c
- dswrite.c
- dudir.c
- dumkfs.c
- ethinit.c
- ethinter.c
- ethread.c
- ethrst.c
- ethstrt.c
- ethwrite.c
- ethwstrt.c
- freebuf.c
- freemem.c
- getaddr.c
- getbuf.c
- getc.c
- getitem.c
- getmem.c
- getname.c

Graph 32 Graph 33 Graph 34 Graph 35 Graph 36

```

graph TD
    dswrite --> dskqopt
    dsread --> dskqopt
    dsseek --> dskqopt
    dsinit --> dskqopt
    dskqopt --> dskstrt
    dskqopt --> dskqopt
    dskstrt --> dskenq
    dsinter --> dskenq
    dskenq --> dskqopt
  
```

Atlas Connection View

Find all functions that reference the structure (dreq) for which memory is allocated

*var x = ref(types("dreq"))
show(x)*

Script: Forward Call

Element Detail View

- dswrite
 - @ XCSG.ModelElement.id = [Xinu/Xinu/sys/dswrite.c#dswrite(struct devsw*, char*, DBADDR)]
 - @ XCSG.ModelElement.name = dswrite
 - @ XCSG.ModelElement.sourceCorrespondence = [322+7+/Xinu/sys/dswrite.c]
 - ##index
 - external
 - isDef
 - XCSG.Function
 - XCSG.Language.C

Problems Tasks Console Properties Atlas Shell (AtlasToolbox)

```

var x = ref(types("dreq"))
x: com.ensoftcorp.atlas.core.query.Q = <Atlas query expression>
show(x)
  
```

Evaluate: <type an expression or enter "help" for more information>

Project Explorer

- csv.s
- ctxsw.c
- ctxsw.s
- devdump.c
- dfalloc.c
- dfdsrch.c
- dalloc.c
- dgclose.c
- dgcntl.c
- ggdump.c
- dginit.c
- dgmctl.c
- dgmopen.c
- dgparse.c
- dgread.c
- dgwrite.c
- dot2ip.c
- dscntl.c
- dsinit.c
- dsinter.c
- dskbcpy.c
- dskenq.c
- dskqopt.c
- dskstrt.c
- dsksync.c
- dsopen.c
- dsread.c
- dsseek.c
- dswrite.c
- dudir.c
- dumkfs.c
- ethinit.c
- ethinter.c
- ethhread.c
- ethrstrt.c
- ethstrt.c
- ethwrite.c
- ethwstrt.c
- freebuf.c
- freemem.c
- getaddr.c
- getbuf.c
- getc.c
- getitem.c
- getmem.c
- getname.c

Graph 33 Graph 34 Graph 35 Graph 36 Graph 37

*var rootx = x.roots()
show(rootx)*

Atlas Connection View Atlas Smart View

Script: Forward Call

Atlas Shell (Project: AtlasToolbox)

```
var rootx = x.roots()
rootx: com.ensoftcorp.atlas.core.query.Q = <Atlas query expression>
show(rootx)
```

Evaluate: <type an expression or enter ':help' for more information>

Element Detail View

- dswrite
 - XCSG.ModelElement.id = [Xinu/Xinu/sys/dswrite.c#dswrite(struct devsw*, char*, DBADDR)]
 - XCSG.ModelElement.name = dswrite
 - XCSG.ModelElement.sourceCorrespondence = [322+7+/Xinu/sys/dswrite.c]
 - ##index
 - external
 - isDef
 - XCSG.Function
 - XCSG.Language.C

Project Explorer

- Graph 34
- Graph 35
- Graph 36
- Graph 37
- Graph 38
- 60
- Atlas Connection View
- Atlas Smart View

freebuf.c

getbuf

var leaves = functions("getbuf","freebuf")
show(leaves)

getbuf.c

getbuf

freebuf.c

freebuf

Script: Forward Call

Problems Tasks Console Properties

Atlas Shell (AtlasToolbox)

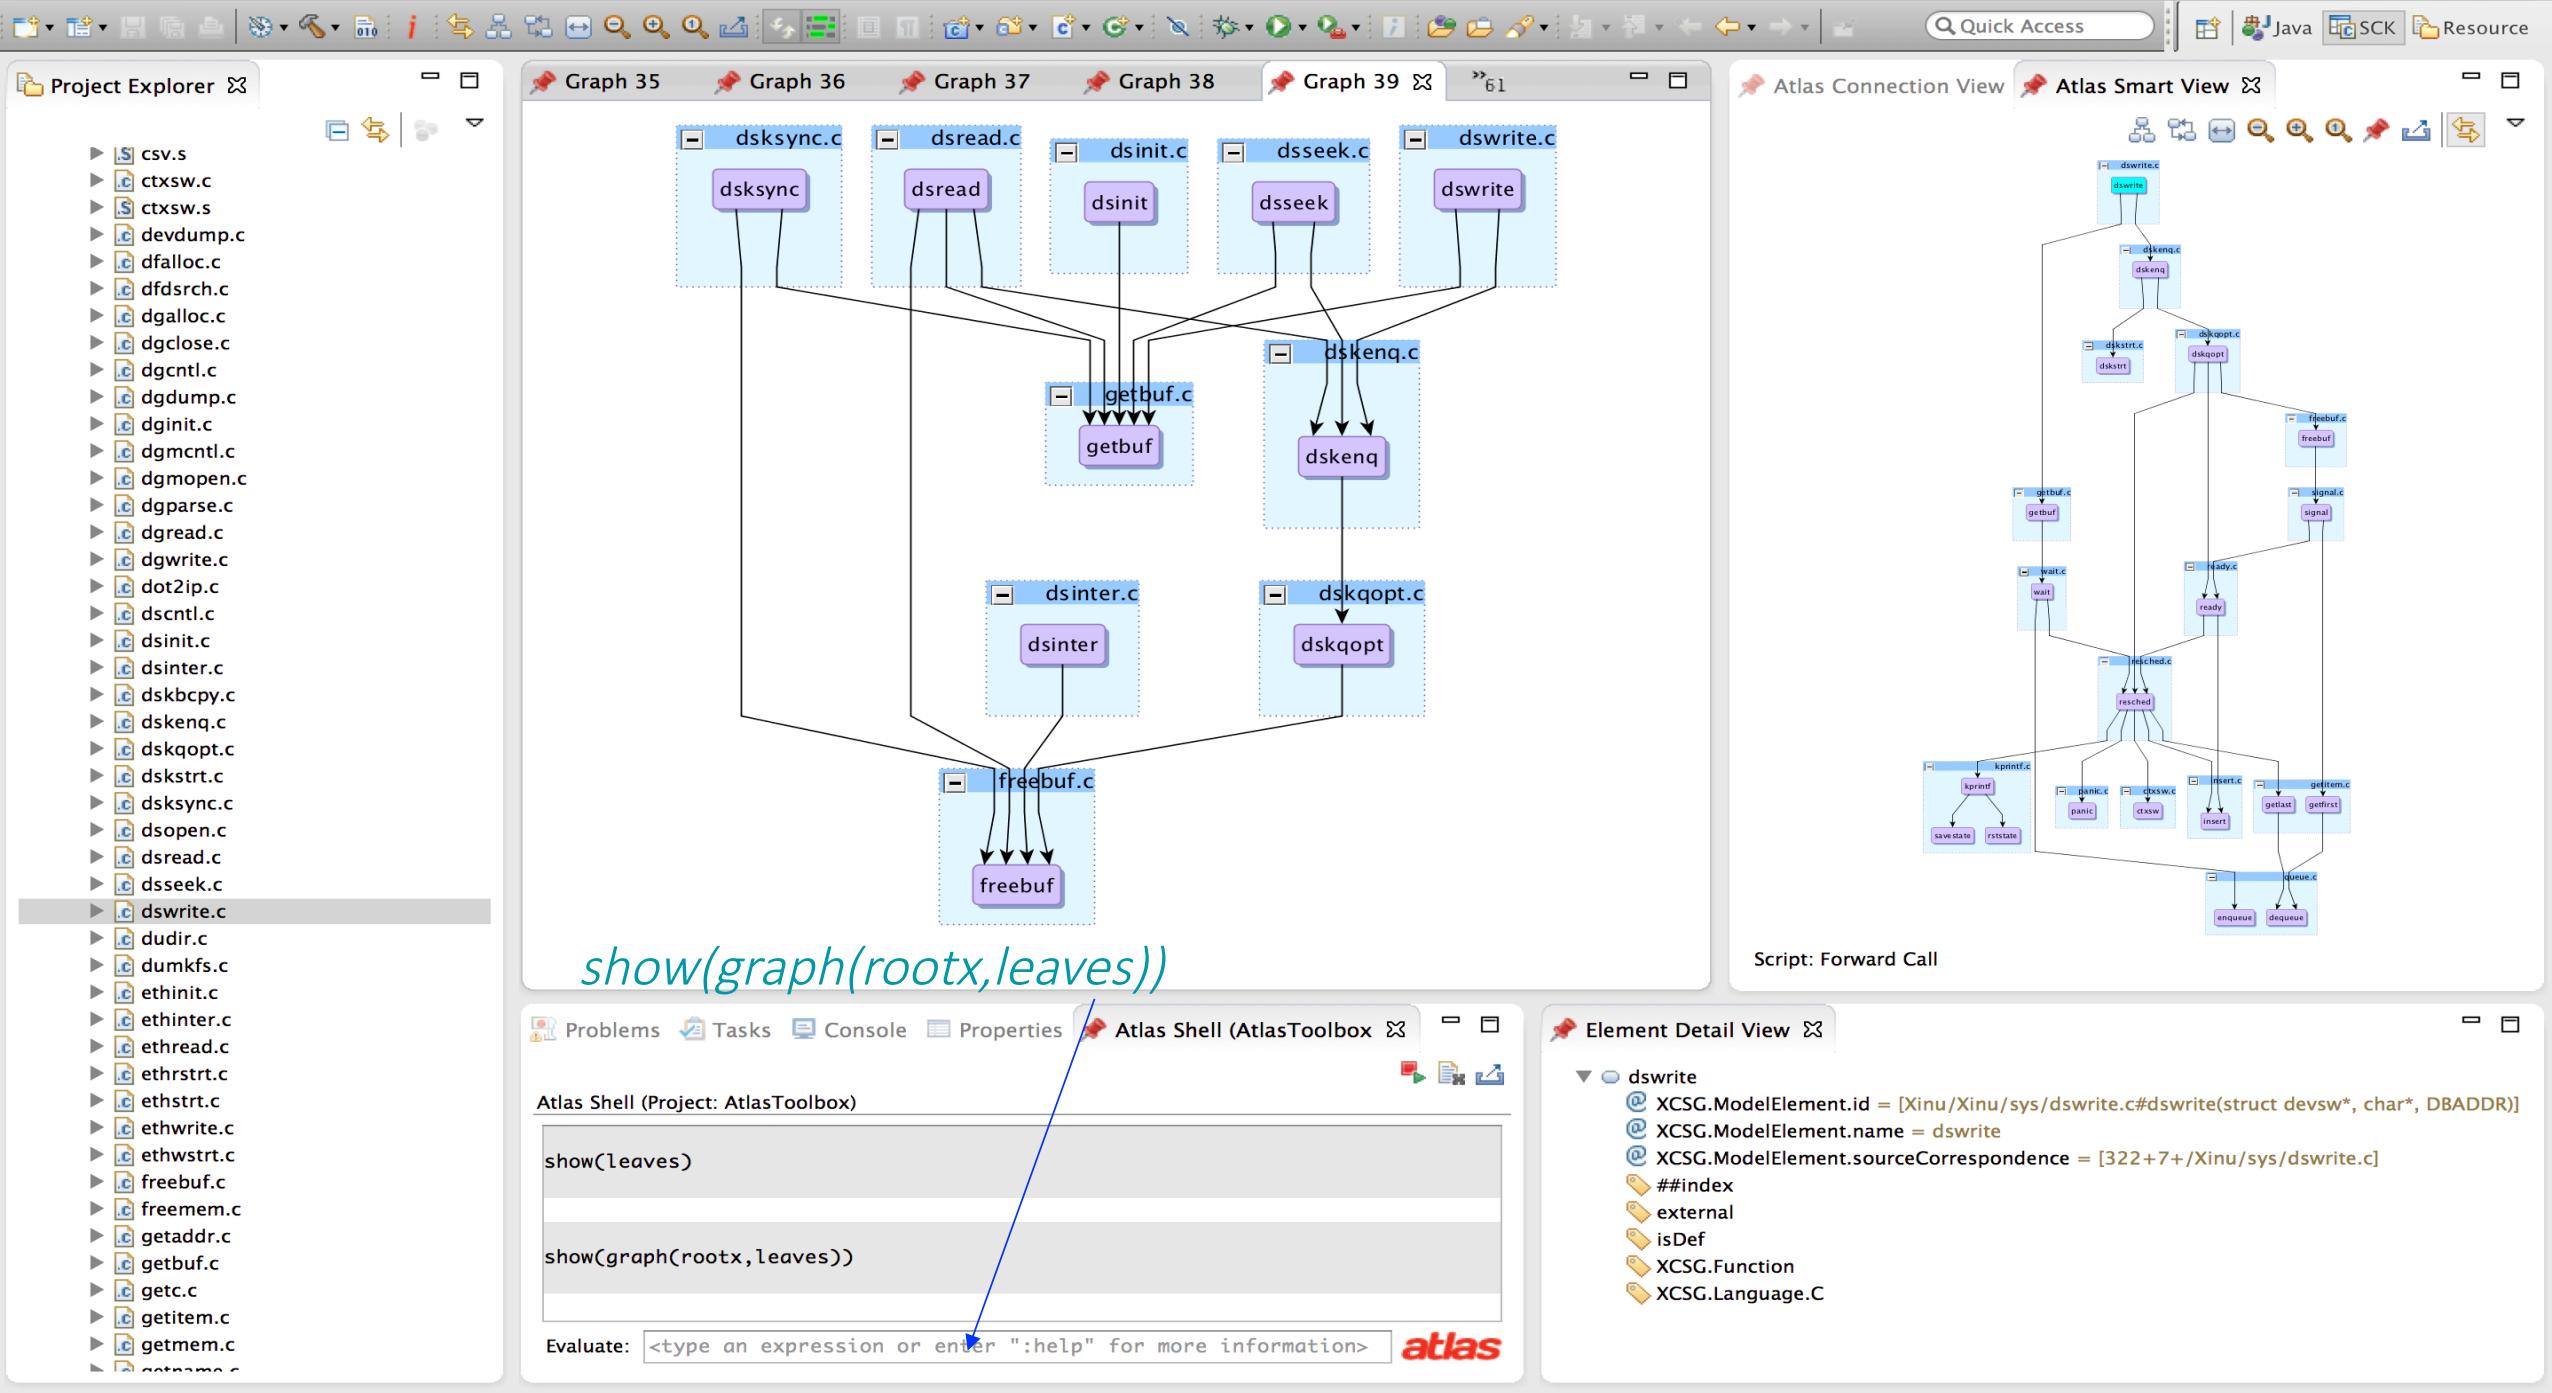
```
var leaves = functions("getbuf","freebuf")
leaves: com.ensoftcorp.atlas.core.query.Q = <Atlas query expression>
show(leaves)
```

Evaluate: <type an expression or enter ":help" for more information>

Element Detail View

dswrite

- XCSG.ModelElement.id = [Xinu/Xinu/sys/dswrite(struct devsw*, char*, DBADDR)]
- XCSG.ModelElement.name = dswrite
- XCSG.ModelElement.sourceCorrespondence = [322+7+/Xinu/sys/dswrite.c]
- ##index
- external
- isDef
- XCSG.Function
- XCSG.Language.C



Project Explorer

- csv.s
- ctxsw.c
- ctxsw.s
- devdump.c
- dfalloc.c
- dfdsrch.c
- dalloc.c
- dgclose.c
- dgcntl.c
- ddump.c
- dinit.c
- dmcntl.c
- dgmopen.c
- dgparse.c
- dgread.c
- dgwrite.c
- dot2ip.c
- dsctl.c
- dsinit.c
- dsinter.c
- dskbcpy.c
- dskenq.c
- dsqopt.c
- dsstrt.c
- dsksync.c
- dsopen.c
- dsread.c
- dsseek.c
- dswrite.c**
- dudir.c
- dumkfs.c
- ethinit.c
- ethinter.c
- ethread.c
- ethrstc.c
- ethstrc.c
- ethwrite.c
- ethwstrt.c
- freebuf.c
- freemem.c
- getaddr.c
- getbuf.c
- getc.c
- getitem.c
- getmem.c
- getname.c

Graph 37 Graph 38 Graph 39 Graph 40 Graph 41 63

```

graph TD
    dswrite[dswrite] --> dskenq[dskenq]
    dswrite --> getbuf[getbuf]
    dskenq --> dskqopt[dsqopt]
    getbuf --> freebuf[freebuf]
    dskqopt --> freebuf
    dsinter[dsinter] --> freebuf
  
```

Atlas Connection View

MPG (relevant functions with call relationships) to verify the **dswrite** allocation instance for memory leak

Script: Forward Call

Problems Tasks Console Properties Atlas Shell (AtlasToolbox)

```

show(a)
show(graph(a,leaves))
  
```

Evaluate: <type an expression or enter ":help" for more information>

Element Detail View

dswrite

- XCSG.ModelElement.id = [Xinu/Xinu/sys/dswrite.c#dswrite(struct devsw*, char*, DBADDR)]
- XCSG.ModelElement.name = dswrite
- XCSG.ModelElement.sourceCorrespondence = [322+7+/Xinu/sys/dswrite.c]
- ##index
- external
- isDef
- XCSG.Function
- XCSG.Language.C

Project: Verify the Linux kernel for a safety property

- Goal

- Bring out the challenges of real-world software.
 - Demonstrate how graph models can address the challenges.

- Planned Activities

- Perform experiments to understand the challenges of multiplicative growth of execution paths and path feasibility.
 - Identify the fundamentals necessary to address the challenges.
 - Apply the principles to design graph models for effective detection and verification.
 - Practice critical thinking through experiments to design the graph models.
 - Discuss research projects that participants can take up as a follow-on.

Project details

- Purpose of the project: verify that the lock is followed unlock on all feasible paths for every locking instance in the Linux kernel.
- The details of the project are provided in a separate document.
- A link is given in the project document to a website that provides the graph models for all the locking instances in a recent version of the Linux kernel.

Overarching ideas for mission-critical verification

Intelligence Amplifying Automation (Frederick Brooks): “If indeed our objective is to build computer systems that solve very challenging problems, my thesis is that IA > AI, that is, that intelligence amplifying systems can, at any given level of available systems technology, beat AI systems. That is, a machine and a mind can beat a mind-imitating machine working by itself.”

Automation Amenable to Human Scrutiny (De Millo, Perllis, and Lipton): “In mathematics, the aim is to increase the human confidence in the correctness of a theorem. [A] proof [does not] settle the matter, contrary to what its name suggests, a proof is only one step in the direction of confidence. Verification is nothing but a model of believability. It cannot be a model where proofs are accepted in blind faith. A proof should be amenable to human scrutiny. *A good proof is one that makes us wiser.*”

The role for automation in mission-critical software verification

- *Prove properties* of programs to find and verify vulnerabilities, where
 - find = identify the relevant program artifacts + hypothesize the vulnerability
 - verify = ascertain the hypothesized vulnerability
- *Produce evidence* to support human scrutiny of verification by the automation.
- *Produce evidence* to amplify human intelligence to complete the verification tasks where automation falls short.
- Examples: Automation such that,
 - it can prove loop termination and produce the termination pattern.
 - if it cannot prove the termination, but it still helps the human to reason about the termination by producing relevant loop artifacts and their characteristics.

Our pragmatic approach to advance automation

- *Advances* are designed to:
 - Increase the share of properties that can be proved automatically
 - Enhance the evidence quality for human scrutiny and for amplifying human intelligence
- *Experimentation Capabilities* to advance theory and automation:
 - Quick prototyping of new property verifiers
 - Exploration to discover new concepts to formulate properties
 - Developing new ideas for high-quality evidence

On experimentation: Gauss declared that his way of arriving at mathematical truths was through systematic experimentation — his notebooks attest to it. As mathematicians do, we need to perform experiments; experiments that motivate future development of theory and automation technology.

Our practical framework

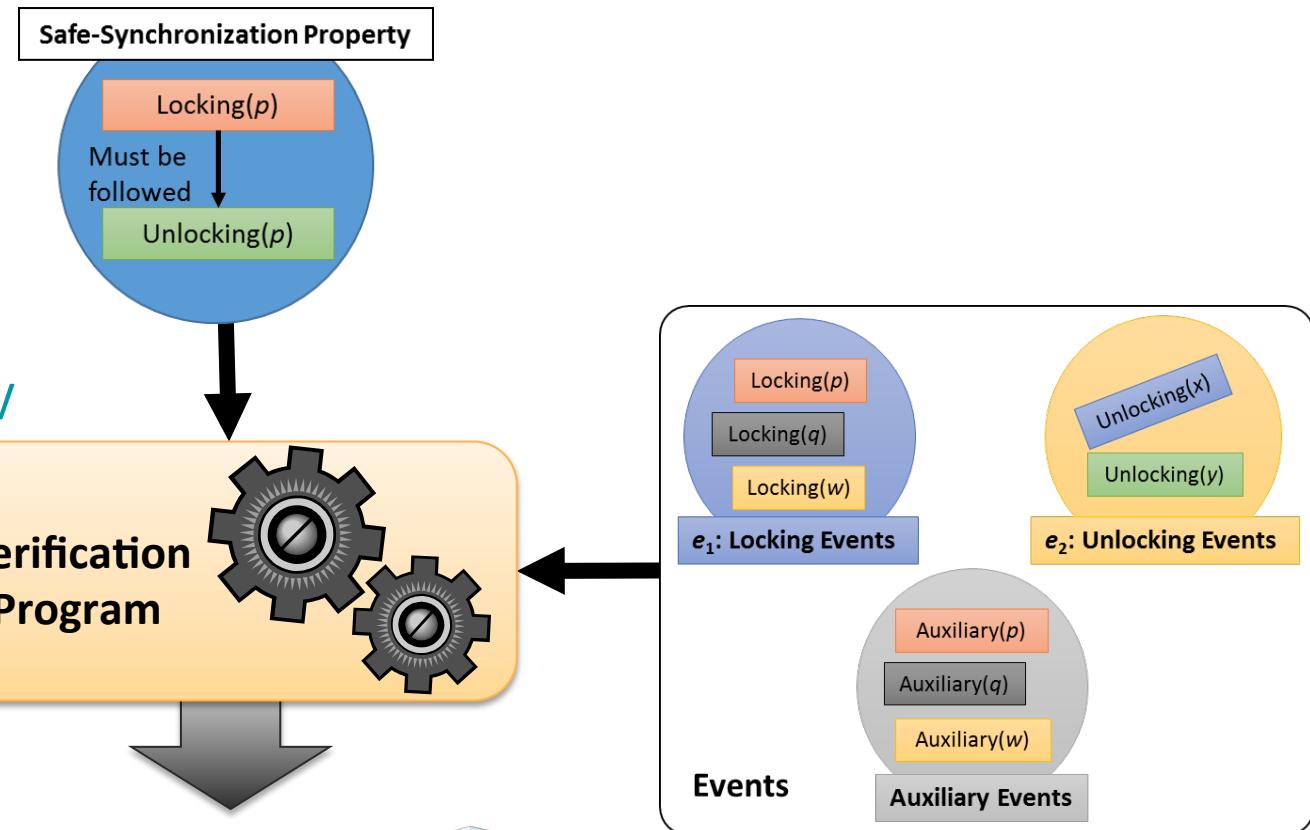
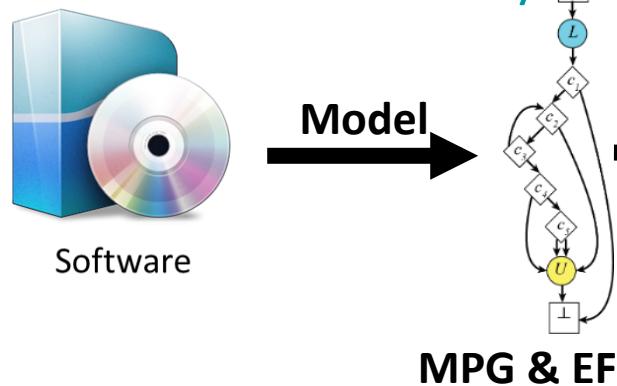
- *Graph Paradigm:*
 - *eXtensible Common Software Graph* (XCSG) schema to capture language semantics
 - currently works across Java, Java byte code, C, C++
 - *Graph database* of pre-computed program artifacts and relationships
 - *Query language* to traverse and transform program graphs
 - *Analysis composition* through the medium of graphs
 - *Interactive visualization, source correspondence* and a query *interpreter*
 - *Plug-in architecture* to build powerful toolboxes that leverage these capabilities
- *Benefits:*
 - Enables quick prototyping of property verifiers
 - Provides a rich infrastructure for experimenting
 - Helps to develop language-independent graph models to enhance scalability and accuracy of analyses
 - Facilitates the leveraging of graph theory to develop new concepts

An overview of Linux verification with our tool

Linux (3.17, 3.18, 3.19), 37 MLOC, Results:

- ✓ : 66,151 instances verified as no bugs
- ✗ : 7 instances reported as bugs
- ? : 451 inconclusive instances

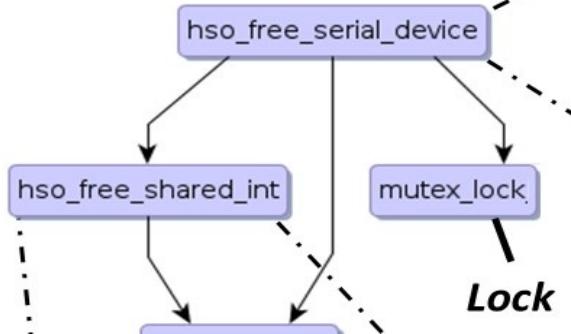
99.3% instances correctly verified by EECV



Three Result Categories with
MPG and EFG as evidence

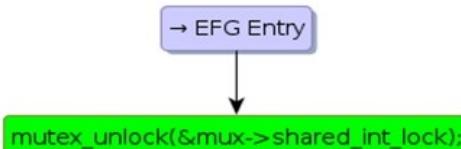
An illustration of an automatically verified instance

(a) Matching Pair Graph



Lock

Unlock



(b) PCG

(hso_free_shared_int)

→ EFG Entry

if (!serial)

if (serial->shared_int)

mutex_lock(&serial->shared_int->shared_int_lock); — Lock

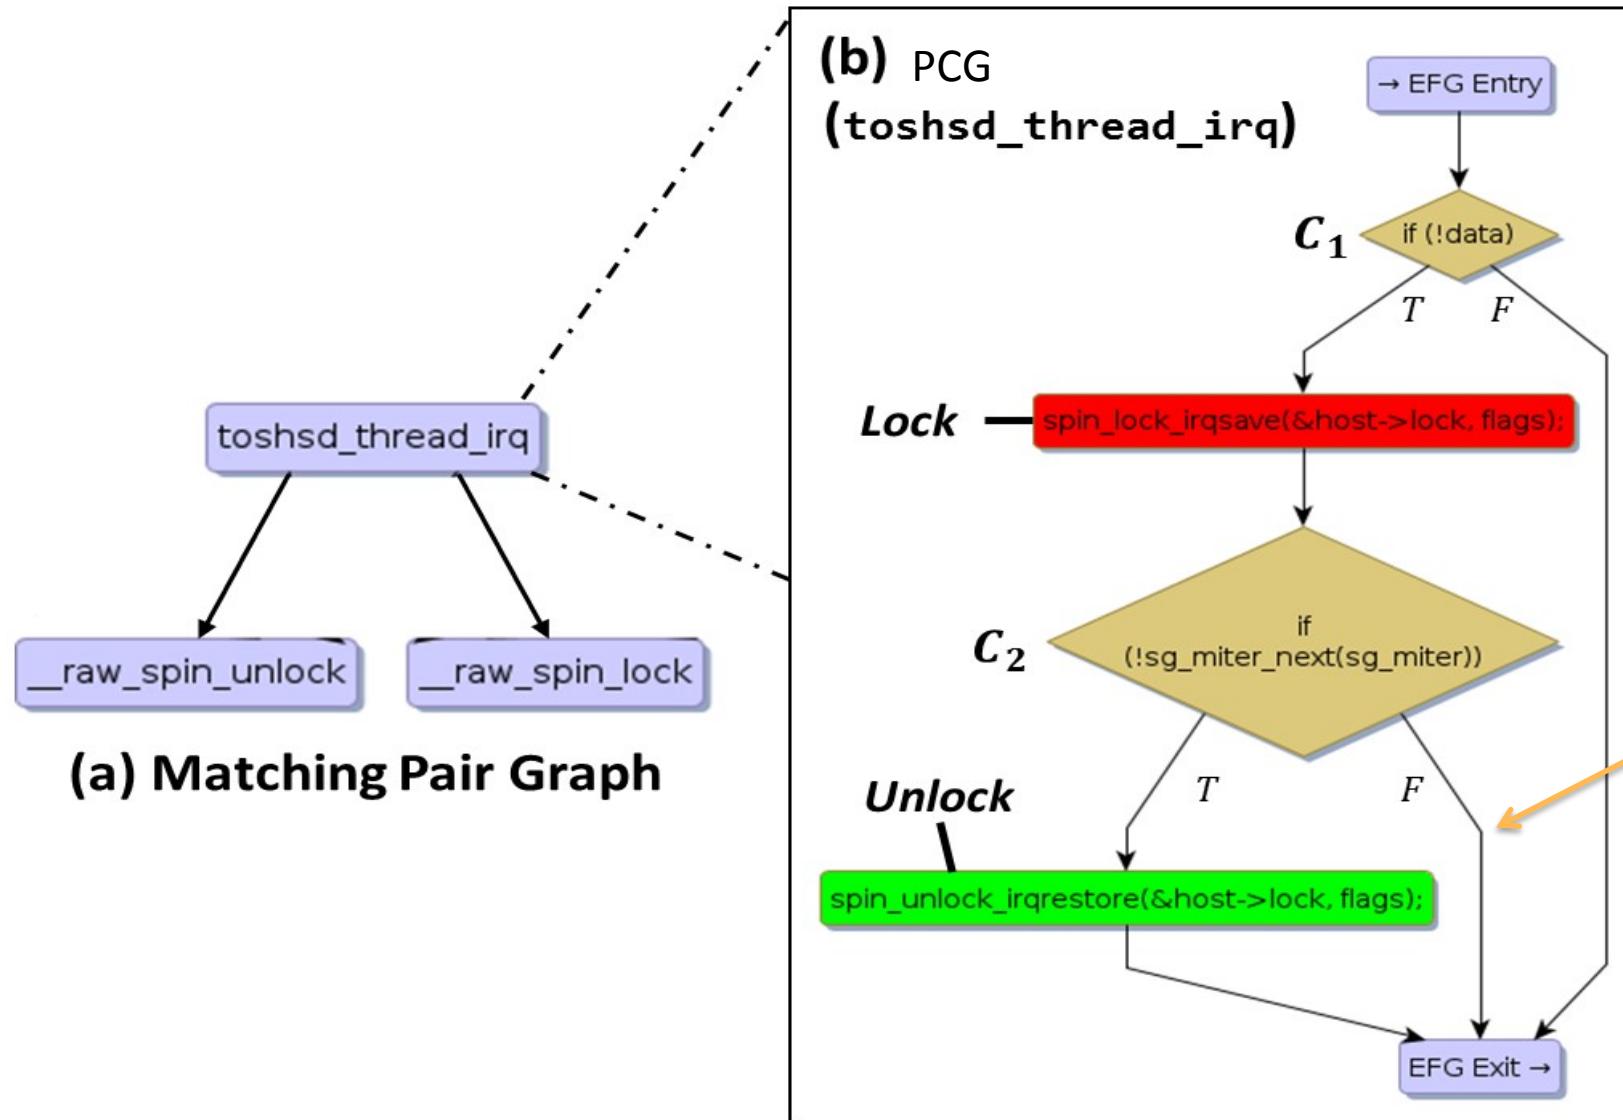
if (--serial->shared_int->ref_count
== 0)

Unlock

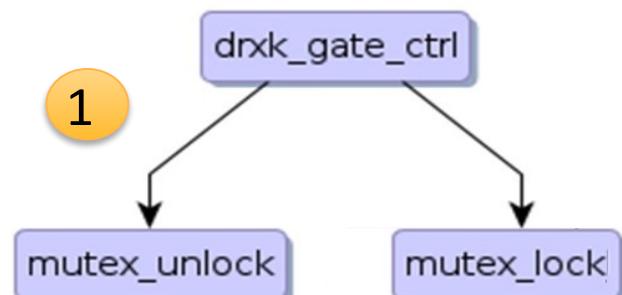
(c) PCG

(hso_free_serial_device)

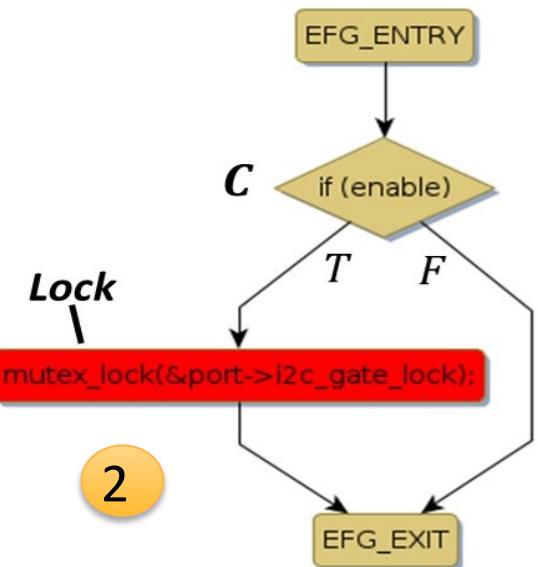
An illustration of automatically discovered bug



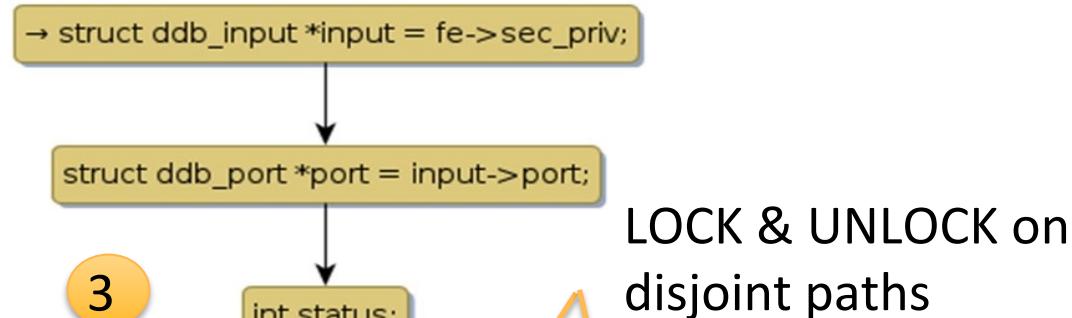
Evidence produced for the bug



(a) MPG



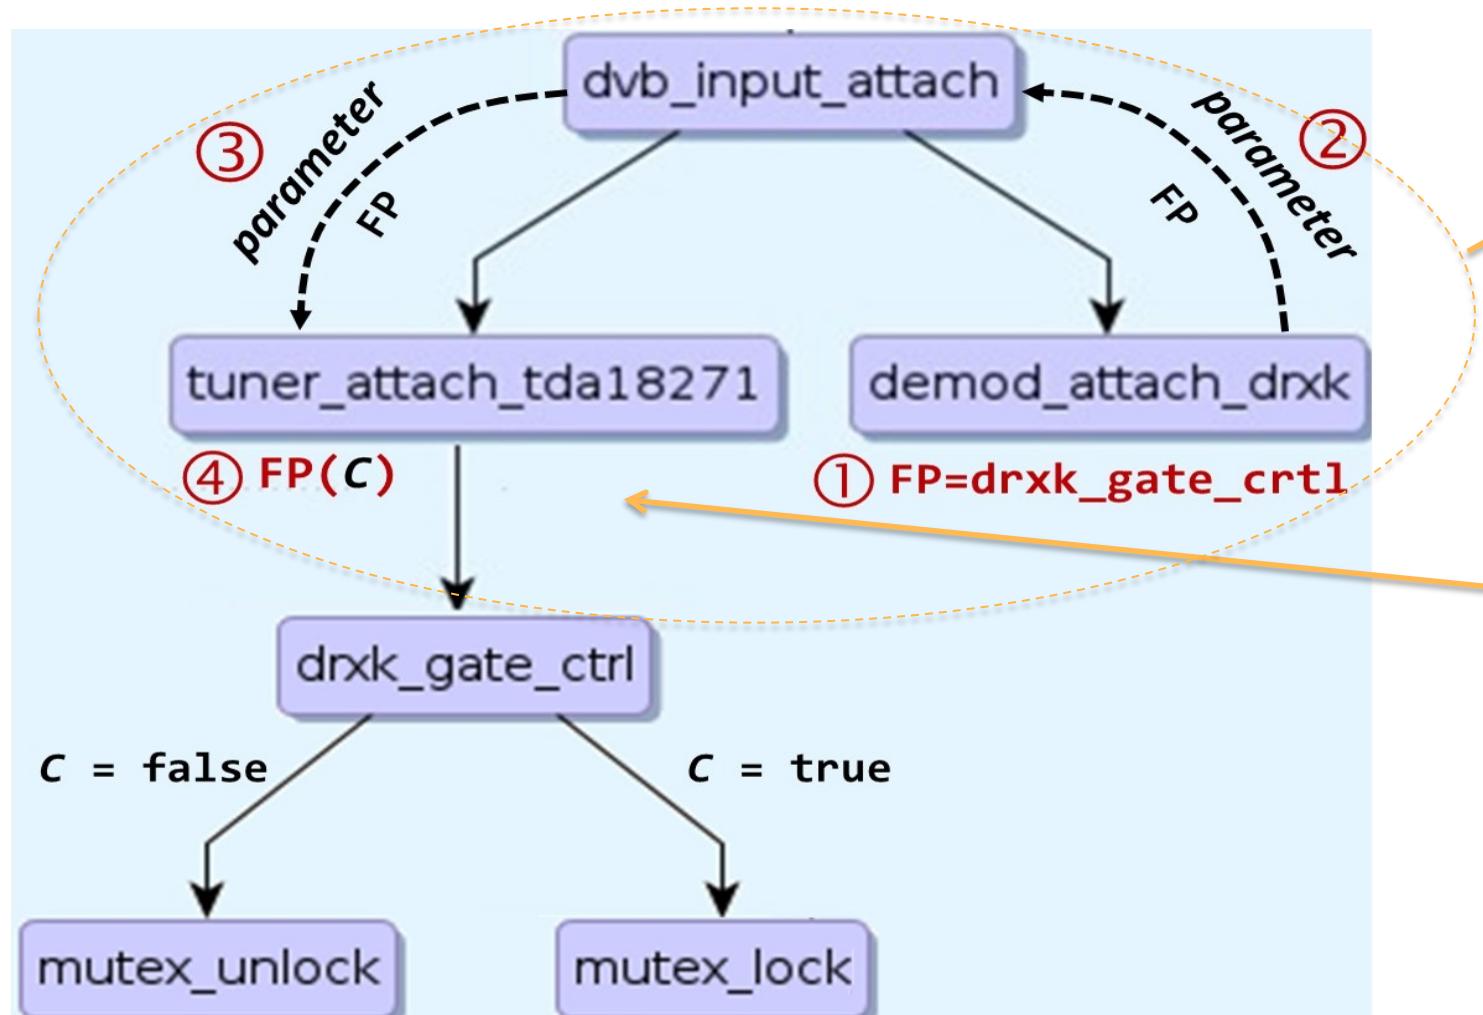
(b) PCG drxk_gate_ctrl



LOCK & UNLOCK on disjoint paths

(c) CFG for drxk_gate_ctrl

Automatic evidence + manual effort for highly complex bug



Manual Effort: A set of queries to find where the function pointer (FP) is set to the function being investigated, and to find the functions that make calls using FP

The function called twice using FP, first call acts like LOCK and second call acts like UNLOCK

Bug: The usage is correct on all but one path where the second call does not happen

Linux Lock/Unlock Stats

- Versions 3.17-rc1, 3.18-rc1, 3.19-rc1
 - 37 MLOC & 66,609 instances
 - 62,663 Intra-Procedural (MPG size =1), average MPG size 1.3, the maximum MPG size 40 for one MPG
 - 55,251 functions in the union of all MPGs
 - 4 nodes per EFG, the maximum EFG size 63 for one EFG

CFG to EFG Reduction – Top Ten

Linux version 3-19-rc-1

| Function Name | Nodes | | Edges | | Branch Nodes | |
|---------------------------|-------|-----|-------|-----|--------------|-----|
| | CFG | EFG | CFG | EFG | CFG | EFG |
| client_common_fill_super | 1,101 | 15 | 1,179 | 28 | 249 | 13 |
| kiblnd_create_conn | 731 | 18 | 925 | 34 | 197 | 15 |
| CopyBufferToControlPacket | 392 | 20 | 559 | 39 | 180 | 18 |
| kiblnd_cm_callback | 662 | 38 | 831 | 56 | 170 | 15 |
| kiblnd_passive_connect | 622 | 22 | 784 | 44 | 164 | 20 |
| dst_ca_ioctl | 349 | 2 | 518 | 1 | 163 | 0 |
| qib_make_ud_req | 621 | 10 | 821 | 15 | 156 | 5 |
| cfs_cpt_table_al | 522 | 7 | 672 | 13 | 153 | 6 |
| private_ioctl | 569 | 16 | 732 | 24 | 148 | 8 |
| vCommandTimer | 490 | 47 | 623 | 75 | 143 | 28 |

Our automatic verification results

- Linux (3.17, 3.18, 3.19), 37 MLOC, LOCK-UNLOCK MPV
- SAP Verification Results:
 - Automatically verified **66,151** instances with no violation (Category-1)
 - Automatically verified **7** instances with no violation (Category-2)
 - **451** Instances of inconclusive automatic verification (Category-3)
- **99.3%** instances automatically verified with no false positives or negatives

Inconclusive Verification – Sources of Difficulty

99.3% automatically verified with remaining 451 Instances of inconclusive automatic verification.

| Kernel | Type | C3 | Intra-procedural Feasibility | Inter-procedural Feasibility | Function Pointers | Signature Problems |
|--------------------|-------|------------|------------------------------|------------------------------|-------------------|--------------------|
| 3.17-rc1 | spin | 82 | 14 | 16 | 22 | 30 |
| | mutex | 73 | 10 | 31 | 28 | 4 |
| 3.18-rc1 | spin | 74 | 13 | 20 | 23 | 18 |
| | mutex | 92 | 13 | 35 | 39 | 5 |
| 3.19-rc1 | spin | 77 | 12 | 16 | 28 | 21 |
| | mutex | 53 | 6 | 22 | 21 | 4 |
| All Kernels | | 451 | 68 | 140 | 161 | 82 |

Targets for the future tool improvements:

1. Automated path feasibility checking.
2. Automated resolution of function pointers to complete call graphs.
3. Improve data flow techniques to handle identified patterns of passing pointers.